

Aggregation is a claim, not a presentation: the geometry of the acceptance gap

Highlights

Aggregation is a claim, not a presentation: the geometry of the acceptance gap

Amin Abaee

- Compensatory and noncompensatory (weak- and strong-sustainability) criteria are compared as robust transition-safety problems on a common datum.
- Aggregate certification can fail path-wise floors on a region with nonempty interior.
- The gap is identified with the failure of quantifier interchange over plans and weights, and characterized geometrically for any finite plan menu.
- Exact per-weight licensing thresholds and a closed-form reserve threshold are derived.
- Convexifying the plan menu closes the gap exactly; discrete time-sharing does not.

Aggregation is a claim, not a presentation: the geometry of the acceptance gap

Amin Abaee^a

^a*Independent Researcher, Tehran, Iran, [ORCID: 0000-0002-0019-1842](https://orcid.org/0000-0002-0019-1842), amin_abaee@ut.ac.ir*

Abstract

In multi-criteria assessment, compensatory and noncompensatory aggregation give two readings of the same constraint set: the scalarized reading, which permits substitution through a weighted aggregate (weak sustainability), and the coordinate-wise reading, which imposes separately-binding floors (strong sustainability). This paper compares the two readings as robust transition-safety problems on a common finite-horizon datum with path-wise constraints. At any fixed path the aggregate is lossless; the two readings nevertheless separate, because common-plan acceptance and per-weight acceptance do not commute.

On an explicit rational datum the acceptance gap is a region with nonempty interior: every nonnegative weighting licenses some transition, no transition is licensed by all, and no transition satisfies the floors path-wise. The gap is the aggregate-certified part of the states whose reserve cannot finance a staged transition; mixing plans as fractional blends closes every aggregate-certified state exactly, while alternating between plans over time does not. A general geometric result characterizes the gap for any finite menu of margin-reducible plans: it is the part of the convex hull of the plans' required margins that no single plan dominates.

Practical message: where separately-binding floors matter, reporting them individually is a detection requirement rather than a presentation preference — the aggregate alone cannot tell the rescuable from the impossible without separately checking the floors or verifying a common implementable plan.

A resource-transition benchmark realizes the datum as a fishery transition — pulse, sustained-yield, and reserve-financed staged plans under a heatwave disturbance — with every value verified in exact rational arithmetic and units anchored to a regulated biomass limit; a terminology guide maps the results onto decision-analysis, control-theoretic, and environmental vocabulary.

Keywords: multi-criteria decision analysis, compensatory and noncompensatory aggregation, composite indicators, scalarization, robustness, sustainability assessment, transition safety

1. Introduction

1.1. Two failure modes

Sustainability assessment sits on a fault line that has structured the field since its foundation. The weak-sustainability tradition, formalized in the genuine-savings and inclusive-wealth frameworks (World Bank, 2011; Boos, 2015; Neumayer, 2013), treats distinct capital forms as substitutable through prices. Under this view, a decline in one stock is consistent with sustainability if the aggregate value of the capital base does not decline. The strong-sustainability tradition holds,

in contrast, that certain stocks and services are separately binding — critical natural capital whose loss cannot be compensated at any price (Daly, 1990; Ekins et al., 2003).

Its foundation statement in ecological economics is the thesis of weak comparability: values relevant to environmental decisions may not be commensurable in a single metric at all (Martinez-Alier, Munda, and O'Neill, 1998). The debate is mature in the policy and multi-criteria decision literatures (Cinelli, Coles, and Kirwan, 2014; Schär, Pohl, and Geldermann, 2025; Hanley et al., 1999; Usubiaga-Liaño, 2025), and the theory of sustainability indicators is correspondingly well developed (Martinet, 2011; Cairns and Martinet, 2014), including the maximin formulation in which strong sustainability emerges from the viability of the constraints rather than from an optimality criterion (Solow, 1974; Doyen and Gajardo, 2020).

Read in this paper's terms, the two traditions are distinguished by whether compensatory substitution keeps pace with depletion. The material-cycle reading of that distinction — waste as a relational status, the closure of the material cycle at the rate of use, substitution against deep-time renewal — is developed in the companion ledger study (Abaee, 2026, *Typed Flux Ledgers and Depletion Arithmetic*) and is not re-argued here; no theorem of this paper concerns material cycles. The doctrines this paper does compare are the operators formalized in Section 3.1 and read doctrinally in Section 5.1.

The case for the abstraction. For a reader whose primary interest is a specific managed system — a fishery, a forest, an aquifer — the value of carrying the comparison at the level of operators, weights, and tubes is that the resulting statements transfer between systems without re-derivation. Whatever the stocks and whatever the plan menu, the statements have the same form: an aggregate criterion can license, at every weighting, a transition that no weighting licenses safely path-wise; the licensed set decouples into a reserve-financed part and an impossible part with an explicit resource threshold; and blending plans, rather than alternating them, is what erases the artifact. Section 6.3 makes the translation concrete on a standard biomass-yield model, and every number displayed in the main text is re-derivable by exact rational arithmetic from the deposited verification code.

The assessment question here is the certification form of that distinction. The blend-collapse result (Theorem 9) closes the acceptance gap exactly on the compensatory region where convexified substitution — fractional allocation of control flows — keeps pace, and the impossibility region of Theorem 5(4) marks where substitution, even under discrete time-sharing, does not. Substitution is admissible only through an identified physical pathway, and no compensation is assumed merely because a weighted aggregate permits it. The aggregate question — whether total capital can be maintained — is ambiguous between the two quantifier orders made precise in Remark 1 and Section 4.5.

This paper addresses a question that is less explicit in much of the composite-indicator literature: the *dynamic* question. During a sustainability transition, when can a compensatory aggregate certify a trajectory that a noncompensatory assessment rejects? Two failure modes motivate the question.

The first is *commensurability drift*. Assessments aggregate many different things into single indices: stocks, services, liabilities, floors; the compensation principles are rarely stated as explicit mathematics. Composite indices are routinely criticized on exactly this ground: the weights they embed are value judgements presented as measurement (Hickel, 2020; Martinez-Alier, Munda, and O'Neill, 1998; Fischer et al., 2022). The critique typically targets the *choice* of weights; the methodology of that choice is itself a developed field (Nardo et al., 2008), and the weights-versus-importance distinction it has produced is part of good composite-indicator practice (Becker et

al., 2017). The compensability question itself is foundational in multi-criteria decision analysis: multi-attribute value theory builds substitution rates between criteria into its tradeoff structure (Keeney and Raiffa, 1976); the outranking tradition constrains compensation through concordance and veto (Roy, 1996); and the decision-aiding treatment of sustainable development centres on incommensurable criteria (Munda, 2005). What is established below is orthogonal to all of these: even if weights perfectly reflect importance and substitution rates are elicited faithfully, the quantifier order can still produce the gap. A complementary methodological strand avoids weights altogether, constructing composite indices from Pareto fronts (Gao et al., 2023). The stronger possibility, established here, is that no choice of weights repairs the difficulty under the per-weight, weight-adaptive certification protocol. The failure is structural — it lies in the logic of aggregation itself.

The second is the circulation of conditional results as unconditional ones. Each result below therefore states its assumptions and status explicitly.

The masking formalized here is the following: an aggregate of two typed floors can stay nonnegative along its worst-case tube while each floor in turn takes a dip that no common plan accepts. It is distinct from the *productivity illusion* — adequate delivery from a reduced productive base — treated in Abaee (2026, *Typed Flux Ledgers and Depletion Arithmetic*). On the witness constructed below the aggregate is taken over the two service floors s_1, s_2 themselves; what it does not detect is the individual floor mid-interval — the compensatory form of the illusion in aggregation, the acceptance gap of Theorem 5(4).

1.2. The central result

The central result is a quantifier noncommutation. Fix a typed transition datum: a state space carrying typed floors $s_i \geq 0$ (normalized at zero), an action set, a disturbance set, and exact-tube semantics under which a transition is safe only if every state visited along it — not merely its endpoint — satisfies the constraints. For each nonnegative scalarization weight w in the full cone $W_+ = \mathbb{R}_+^n \setminus \{0\}$, let $E_w(z)$ be the set of actions admissible at state z under the aggregate floor $w \cdot s \geq 0$, and let $E_{\text{typ}}(z)$ be the set of actions admissible under the typed floors $s_i \geq 0$ taken separately. Two acceptance criteria then differ by quantifier order:

- **Common-plan acceptance (noncompensatory, coordinate-wise).** There exists one action lying in every $E_w(z)$:

$$\exists a \forall w : a \in E_w(z).$$

- **Per-weight acceptance (compensatory, scalarized).** For each weight there exists some admissible action, possibly depending on the weight:

$$\forall w \exists a_w : a_w \in E_w(z).$$

To isolate this mechanism, the assessment protocols an indicator can serve must be distinguished from the indicator itself. Table 1 collects the four protocols in play; the separation studied in this paper is specifically Protocol 2 versus Protocol 4 — a statement about the certification protocol applied to a composite index, not an objection to composite indices as such (Protocol 1, at a fixed declared weight, is ordinary fixed-weight assessment, and Protocol 3 collapses into Protocol 4 by the full-cone identity of Proposition 3(ii)). In the vocabulary of decision making under uncertainty, the two protocols differ in when the plan is fixed relative to the assessment weight: Protocol 4 commits to one plan before the weight is known, while Protocol 2 lets the plan adapt once the weight is

observed — the here-and-now versus wait-and-see distinction of adjustable robust optimization (Ben-Tal et al., 2004), with the weight in the role of the uncertain parameter.

Table 1: Four assessment protocols for a composite index on a transition datum. The separation studied in this paper is Protocol 2 versus Protocol 4.

Protocol	Quantifier form	Reading
1. Fixed-weight assessment	$a \in E_w(z)$ at a declared w	one index, one selected plan: ordinary fixed-weight certification
2. Weight-adaptive (per-weight) acceptance	$\forall w \exists a_w \in E_w(z)$	every weighting can be paired with some plan, possibly weighting-dependent: the compensatory reading; $\mathcal{V}_{\text{weak}} = \bigcap_w \mathcal{V}_w$
3. Common-plan robustness	$\exists a \forall w : a \in E_w(z)$	one plan serves every weighting; equals the typed operator on the full cone (Proposition 3(ii))
4. Coordinate-wise (typed) acceptance	$\exists a : a \in E_{\text{typ}}(z)$	separately binding floors enforced path-wise: the noncompensatory reading

Where the failure is not. Static scalarization is not the source of the counterexample: at any fixed trajectory the full-cone aggregate detects every component violation (Remark 2), and on a fixed action the typed operator is exactly the common-plan set over the full cone (Proposition 3(ii)). The failure studied here appears one level up — between protocols — because different weights may select different trajectories: the quantifier order, not the scalarization, is the seat of the separation (Sections 4.5–4.6).

A firm or public authority must carry out a transition while satisfying two separate minimum-service constraints. Rapid restructuring exposes one constraint to a temporary drawdown; a gradual plan avoids drawdowns but requires a liquidity reserve. An evaluator who aggregates the two constraints into one weighted index may approve the transition for every relative price, while an evaluator who requires each constraint to hold along the entire path approves none of the available plans. The two quantifier orders make this divergence precise.

The first implies the second; the converse fails. The failure is not a defect of any particular weight. It is the noncommutativity of the existential quantifier over actions with the universal quantifier over weights. The noncommutation is the assessment-theoretic instance of the strict minimax pattern of game theory: for the binary payoff “action a is admissible at z under w ”, the two quantifier orders are the two orders of the minimax interchange, and the witness of Section 4.5 exhibits the interchange failing strictly. Remark 1 states the always-valid inclusion; Proposition 3(ii) identifies the noncompensatory operator with the common-plan set over the full cone; and Theorem 5 exhibits an explicit rational witness on which the acceptance gap $\text{FP}_{\text{agg}} = \mathcal{V}_{\text{weak}} \setminus \mathcal{V}_{\text{typ}}$ is a region with nonempty interior — the *impossibility region* (Theorem 5(4)). The witness further partitions the aggregate-versus-direct-floor discrepancy region $\mathcal{Q}$ by a resource threshold into the impossibility region and the *rescue set*, whose states are already typed-transformable through the resource-controlled action (Theorem 5(7)). The rescue operation of Section 5.7 — resource augmentation — acts on the impossibility region, not on the rescue set. A general geometric result in Section 4.4 characterizes the acceptance gap of any finite plan menu: it is the part of the convex hull of the plans’ required margins that no single plan dominates.

1.3. Contribution and scope

Contributions. (i) The action-set identity $E_{\text{typ}}(z) = \bigcap_{w \in W_+} E_w(z)$, with the full-cone choice (the cone without the origin) isolating the separation as purely dynamic (Proposition 3(ii)). (ii) A general quantifier-separation remark (Remark 1). (iii) Monotonicity of the accepted-state hierarchy in the weight family (Proposition 4). (iv) An explicit rational witness with an open region of strict separation, the discrepancy-region split $\mathcal{Q} = R \cup I$, the identity $\text{FP}_{\text{agg}} = I$, the exact per-weight licensing thresholds (Theorem 5), and the closed-form rescue threshold $\kappa^*(z) = (1 - x) \cdot \mathbf{1}[z \notin \mathcal{V}_{\text{typ}}]$ with the error-bound remark (Proposition 11, Section 5.7). (v) Persistence of the separation under hold-prefix extension of the horizon (Remark 7). (vi) An explicit two-stage erasure datum on which a later interval removes the gap (Theorem 8), delimiting the propagation claim. (vii) The blend-collapse theorem: menu convexification closes the acceptance gap exactly at the compensatory region (Theorem 9), while the converse delimitation (Proposition 10) shows that discrete time-sharing — whose visited set is the union of the action tubes — closes it only where both dips are independently subsumed, so the gap survives on the impossibility region. (viii) The finite-menu geometric characterization (Theorem 6, Section 4.4): the acceptance gap of any finite plan menu is the part of the convex hull of the required margins that no single plan dominates, subsuming Theorem 5’s gap and Theorem 9’s collapse as instances. (ix) A resource-transition benchmark (Section 6.3): the witness datum realized as a Schaefer-type fishery transition — pulse-plus-closure, sustained-yield, and reserve-financed staged plans as management schedules — with all quoted values machine-verified and the certified tubes shown to be conservative for the nonlinear realization. (x) A terminology translation guide (Section 5.3) mapping the paper’s central objects onto the working vocabularies of MCDM and operations research, control and viability theory, and environmental modelling and governance. (xi) A labelled substitutability-elasticity extension on the witness datum (Section 5.8): the linear aggregate and the typed operator as the $\sigma = \infty$ and $\sigma = 0$ members of one CES family, an exact critical-floor ladder with algebraic floors ($\sqrt{2}$, φ , $\sqrt{3}$), and two machine-anchored theorems (S1 nesting; S2 the two-sided Leontief statement). (xii) A labelled common-shock extension of the witness datum (Section 5.9): the middle-regime disturbance verdict of Table 5 computed exactly (the acceptance gap survives the common shock, reduced by the exact 1/2-cut, licensing thresholds unshifted on the surviving region), the class-conditional rescue trichotomy ($\kappa^* = 0$, $(1 - x)_+$, or ∞ , with the first unrescuable-by-reserve failure states on this datum), the class-conditional caveat on the blend collapse (Theorem 9’s window, not blend-closed on the fragile band), and two exact reading-conditional theorems for the coupled regime (the additive gap relocating rather than vanishing; the replace reading collapsing the typed/weak distinction itself) — machine-verified as a fourth check list. (xiii) The off-diagonal closure of the same programme (Section 5.8): the exact off-diagonal cover criterion (Theorem G) — the log-transcendental comparison decided by certified rational enclosures, the acceptance region bounded by the strictly decreasing involution ρ with its exact $\sqrt{2}$ pivot — the completed two-dimensional σ^* -landscape (all 13 rungs decided at all 64 grid states, zero cells undecided), and the fixed-sum tolerance structure whose interior window shows moderate concentration of a deficit more certifiable than balance — the benchmark’s asymmetric-allocation decision, where the same total margin draws opposite verdicts and opposite management responses — together with the continuum slice closure (the ϕ -reduction deciding the slice-minimum: the flip total $t^* = a^* + b^*$ certified to ten places by two independent code paths, the accepting windows proved to be intervals, the total-3 boundary proved in the interior, the harmonic circle, the σ^* -depth sandwich, and the flip-total ladder) — machine-verified as fifth and sixth check lists.

Scope. No universal ranking of weak- and strong-sustainability doctrines; no proof that any particular set of environmental boundaries ought to be treated noncompensatorily; no empirical result about any specific resource system. The governance, intergenerational, and composition extensions are stated at partial status in the Supplementary Material.

2. A Typed Framework for Transition Assessment

This section fixes the definitions the theorems require.

2.1. Types and physical state

Physical state is typed. A state variable denotes a *moiety* — a named conserved substance — with a unit, and typed fluxes connect typed stocks. Conservation claims are per-moiety; the framework does not authorize adding biomass, money, and biodiversity into one conserved scalar. Services, thresholds, information states, and institutional variables are separate types. No statement combines quantities of different type except through an explicit bridge. Typing serves two purposes: it makes the domain of every conservation law explicit, and it records the provenance of every floor — physical, contractual, or normative — as part of the floor’s type.

2.2. The assessment datum

The results use a typed exact-tube transition datum: a state space Z carrying typed floors $s_i \geq 0$ (normalized at zero), an action set, a disturbance class D , a typed safe set S , a destination set G , and tube and successor maps $\text{Tube}(a, d)$, $\text{Succ}(a, d)$ as defined in Section 3.1. A transition is safe only if every state visited along it, not merely its endpoint, satisfies the constraints. The witness datum of Section 4.5 is an instance of this datum.

2.3. Quantifier discipline

Every assessment operator evaluates the disturbance quantifier innermost: admissibility of an action requires safety for every $d \in D$. The separation analysed in Section 3 concerns a second quantifier, over scalarization weights, placed outside it: the noncommutation studied is $\exists a \forall w$ versus $\forall w \exists a_w$.

2.4. Notation

The following symbols are used throughout the paper.

- $z \in Z$: a state; a : an action; $d \in D$: a disturbance; $s = (s_1, \dots, s_n)$: typed floors (normalized so that $s_i \geq 0$ is the binding constraint); x : the reserve stock (bridge stock) on the witness.
- $W_+ = \mathbb{R}_+^n \setminus \{0\}$: the full nonnegative scalarization cone excluding the origin (Section 3.1); $w \in W_+$: a nonnegative scalarization weight; a subfamily is written W , as in Proposition 4; $r = w_2/w_1$: the weight ratio used on the witness (the resource increment of Section 5.6 is κ).
- $S = S^{\text{phys}} \cap \{s_i \geq 0, i = 1..n\}$ and $G = G^{\text{phys}} \cap \{s \geq 0\}$: typed safe and destination sets; S^w, G^w : their scalarized counterparts at weight w ; $S^{\text{phys}}, G^{\text{phys}}$: their physical counterparts; on the witness $S = S_0$, the transition-safe set of Section 4.5.
- $\text{Tube}(a, d)$: the full set of states visited during the interval under action a and disturbance d ; $\text{End}(a, d)$: the endpoint values visited; $\text{Succ}(a, d)$: the successor set after any endpoint reset.

- $E_{\text{typ}}(z)$, $E_w(z)$, $E_{\text{tube,phys}}(z)$, $E_{\text{end}}(z)$, $E_{\text{end,typ}}(z)$: the assessment operators of Section 3.1.
- $\mathcal{V}[E] = \{z : E(z) \neq \emptyset\}$: the accepted-state set of operator E ; in particular $\mathcal{V}_{\text{typ}} = \mathcal{V}[E_{\text{typ}}]$, $\mathcal{V}_w = \mathcal{V}[E_w]$, $\mathcal{V}_{\text{weak}} = \bigcap_{w \in W_+} \mathcal{V}_w$, $\mathcal{V}_{\text{phys}} = \mathcal{V}[E_{\text{tube,phys}}]$, $\mathcal{V}_{\text{end}} = \mathcal{V}[E_{\text{end}}]$.
- Witness datum (Section 4.5): phase state (q, x, s_1, s_2) ; gain vector $e = (1/4, 1/4)$; rescue cost $c = 1$; depth of worst-case dip 2; menu $\{\text{NO-SWITCH, FAST, SLOW, STAGED}\}$.
- Discrepancy sets (Section 4.6): $\mathcal{Q}$ aggregate-versus-direct-floor discrepancy region; R rescue set; I impossibility region; $\text{FP}_{\text{agg}} = \mathcal{V}_{\text{weak}} \setminus \mathcal{V}_{\text{typ}}$ the genuine acceptance gap; ρ_1, ρ_2 the per-weight licensing thresholds.
- $\mathcal{A}$: the action set (Sections 4.1 and 5.5); Aug_κ : the resource-augmentation map, with increment κ and minimal rescue threshold κ^* .

The gain vector e of the witness datum and the standard basis vector e_k in Remark 2 are distinct objects.

3. Assessment Operators

3.1. Five operators

Fix a typed exact-tube transition datum with transition-safe sets and a destination set. For a state z and action a , let $\text{Tube}(a, d)$ be the full set of states visited during the interval under disturbance d , let $\text{End}(a, d)$ be the endpoint values visited, and let $\text{Succ}(a, d)$ be the successor set after any endpoint reset. Five operators are distinguished. They share the same disturbance quantifier and differ in constraint structure, except that the endpoint operator also replaces tube evaluation by endpoint evaluation. The fifth is the typed-endpoint operator, used in Section 5.5.

The **noncompensatory typed** operator (each floor separately binding):

$$E_{\text{typ}}(z) = \{a : \forall d, \text{Tube}(a, d) \subseteq S \text{ and } \text{Succ}(a, d) \subseteq G\},$$

where $S = S^{\text{phys}} \cap \{s_i \geq 0, i = 1..n\}$ and $G = G^{\text{phys}} \cap \{s \geq 0\}$.

The **scalarized aggregate** operator at weight $w \in W_+ = \mathbb{R}_+^n \setminus \{0\}$:

$$E_w(z) = \{a : \forall d, \text{Tube}(a, d) \subseteq S^w \text{ and } \text{Succ}(a, d) \subseteq G^w\},$$

where $S^w = S^{\text{phys}} \cap \{w \cdot s \geq 0\}$ and $G^w = G^{\text{phys}} \cap \{w \cdot s \geq 0\}$.

The **exact-tube physical** operator (physical constraints only, full tube):

$$E_{\text{tube,phys}}(z) = \{a : \forall d, \text{Tube}(a, d) \subseteq S^{\text{phys}} \text{ and } \text{Succ}(a, d) \subseteq G^{\text{phys}}\}.$$

The **endpoint-only physical** operator (physical constraints at endpoints only):

$$E_{\text{end}}(z) = \{a : \forall d, \text{End}(a, d) \subseteq S^{\text{phys}} \text{ and } \text{Succ}(a, d) \subseteq G^{\text{phys}}\}.$$

The endpoint operator represents any assessment architecture that evaluates constraints only at sampled or audited endpoints, while typed floors constrain the whole trajectory: evaluation on the full tube versus evaluation on endpoints only. The endpoints are the photograph; the tube is the trajectory. Endpoint values cannot in general certify the condition of the system that produced them (cf. the productivity-illusion analysis in Abaee, 2026, *Typed Flux Ledgers and Depletion*

Arithmetic); the typed-endpoint operator below states the floor-level version of this observation exactly on the witness. Because $\text{End}(a, d) \subseteq \text{Tube}(a, d)$, the endpoint operator is the weakest, and the chain

$$E_{\text{typ}}(z) \subseteq E_w(z) \subseteq E_{\text{tube,phys}}(z) \subseteq E_{\text{end}}(z)$$

holds for every $w \in W_+$ (Proposition 3(i) below). The last inclusion is strict in general, though it collapses on the witness of Section 4.5 because on the witness the physical constraint involves only the monotone reserve stock.

Definition (typed-endpoint operator). The **typed-endpoint** operator evaluates typed floors at endpoints only:

$$E_{\text{end,typ}}(z) = \{a : \forall d, \text{End}(a, d) \subseteq S \text{ and } \text{Succ}(a, d) \subseteq G\}.$$

It satisfies $E_{\text{typ}}(z) \subseteq E_{\text{end,typ}}(z) \subseteq E_{\text{end}}(z)$, since $\text{End}(a, d) \subseteq \text{Tube}(a, d)$, $S \subseteq S^{\text{phys}}$, and $G \subseteq G^{\text{phys}}$. On the witness datum of Section 4.5, FAST’s endpoint values of s_1 equal the initial s_1 and its successor lies in G whenever $x \geq 0$ (action table in Section 4.5; proof of Theorem 5(1)), so $E_{\text{end,typ}}(z) \neq \emptyset$ at every $z \in X_0$, while $E_{\text{typ}}(z) \neq \emptyset$ holds exactly on $\{x \geq 1\} \cup \{s_1 \geq 2\} \cup \{s_2 \geq 2\}$ (Theorem 5(1)): on the stricken slice of the discrepancy region, $\mathcal{Q} \cap \{x < 1\}$, the endpoint evaluation is nonempty where the typed operator is empty — the photograph admits what the trajectory forbids.

The five operators differ along two axes — what is constrained, and where it is evaluated — and the two axes organize them (Table 2). The second axis is what endpoint-only reporting regimes collapse (the fourth implication of Section 5.6).

Table 2: Evaluation semantics: constraint structure (rows) by evaluation locus (columns). The scalarized-endpoint cell is not isolated in this paper; the endpoint chain of Section 3.1 covers the physical and typed extremes.

Evaluated on the full tube		Evaluated at endpoints only
T yped floors	E_{typ} — typed path safety	$E_{\text{end,typ}}$ — typed snapshots
S calarized floor	E_w — scalarized path safety	— (not isolated here)
P hysical con- s traints	$E_{\text{tube,phys}}$ — physical viability	E_{end} — audited physical snapshots

3.2. Accepted-state operators

For any operator E , write

$$\mathcal{V}[E] = \{z : E(z) \neq \emptyset\}$$

for the set of states at which E admits at least one action. Thus $\mathcal{V}_{\text{typ}} = \mathcal{V}[E_{\text{typ}}]$, $\mathcal{V}_w = \mathcal{V}[E_w]$, $\mathcal{V}_{\text{phys}} = \mathcal{V}[E_{\text{tube,phys}}]$, $\mathcal{V}_{\text{end}} = \mathcal{V}[E_{\text{end}}]$, and the compensatory accepted set

$$\mathcal{V}_{\text{weak}} = \bigcap_{w \in W_+} \mathcal{V}_w.$$

The central noncommutation is then:

$$\underbrace{\{z : \bigcap_w E_w(z) \neq \emptyset\}}_{\mathcal{V}_{\text{typ}} \text{ (common plan)}} \subseteq \underbrace{\bigcap_w \{z : E_w(z) \neq \emptyset\}}_{\mathcal{V}_{\text{weak}} \text{ (per weight)}}.$$

4. Results: The Separation Theorem

4.1. A general quantifier-separation remark

Remark 1. Let $\{E_\lambda(z)\}_{\lambda \in \Lambda}$ be a family of action sets. Then always

$$\left\{ z : \bigcap_\lambda E_\lambda(z) \neq \emptyset \right\} \subseteq \bigcap_\lambda \left\{ z : E_\lambda(z) \neq \emptyset \right\}.$$

Equality holds exactly when, for every z in the right-hand set, the family $\{E_\lambda(z)\}_\lambda$ admits a common selector. No topology on the action space is imposed. In finite menus the check is combinatorial: if the action set is finite, only finitely many distinct $E_\lambda(z)$ occur and common-selector existence is a finite enumeration (on the witness of Section 4.5, $|\mathcal{A}| = 4$). For infinite families, classical sufficient conditions — the finite-intersection property together with compactness/closedness in a specified topology — are available but not needed here.

Proof. See Appendix A. The separation in Section 4.5 is a strict instance of this inclusion, at states where the nonemptiness witnesses differ with the weight.

4.2. The full-cone identity

Remark 2 (full-cone pointwise equivalence). For $v \in \mathbb{R}^n$,

$$v \geq 0 \text{ componentwise} \iff w \cdot v \geq 0 \text{ for every } w \in W_+ = \mathbb{R}_+^n \setminus \{0\}.$$

Proof. See Appendix A. At a *fixed* trajectory the full-cone aggregate is therefore lossless: nonnegativity of every weighted sum is equivalent to componentwise nonnegativity. The separation below is entirely dynamic — a matter of quantifier order — not static scalarization blindness.

4.3. The assessment identity and hierarchy

Proposition 3 (assessment identity and hierarchy).

(i) *Hierarchy.* For every $w \in W_+$ and every z ,

$$E_{\text{typ}}(z) \subseteq E_w(z) \subseteq E_{\text{tube,phys}}(z) \subseteq E_{\text{end}}(z).$$

(ii) *Localization.* For every z ,

$$E_{\text{typ}}(z) = \bigcap_{w \in W_+} E_w(z).$$

Hence $\mathcal{V}_{\text{typ}} = \{z : \bigcap_w E_w(z) \neq \emptyset\}$.

Proof. See Appendix A. Part (i) is standard constraint-set monotonicity under nested safe sets (Aubin, 1991; Frankowska, 1989); part (ii) is proved above, and its strictness is witnessed in Theorem 5.

4.4. Weight-family monotonicity

Proposition 4 (weight-family monotonicity). *For a weight family $W \subseteq W_+$ define $\mathcal{V}_W = \bigcap_{w \in W} \mathcal{V}_w$. Then*

$$\mathcal{V}_{\text{typ}} \subseteq \mathcal{V}_W \subseteq \mathcal{V}_{\text{phys}},$$

and if $W_1 \subseteq W_2$ then $\mathcal{V}_{W_2} \subseteq \mathcal{V}_{W_1}$.

Proof. See Appendix A. Proposition 4 makes precise the claim that *restricting* the weight family *enlarges* the compensatory accepted set. The full cone is therefore the strictest compensatory reading, and any separation exhibited at the full cone persists for every restricted family.

Theorem 6 (finite-menu geometric characterization). *Let a finite menu $\mathcal{A}$ of plans be given, and suppose each plan's coordinate-wise admissibility is governed by a required-margin vector $d_a \in \mathbb{R}_+^n$: plan a is admissible at floor margins $s \geq 0$ exactly when $d_a \leq s$ componentwise, every plan satisfying the destination constraint. (Plans outside this margin class — reserve-financed or destination-failing plans — are handled separately, as in Section 4.5.) Write $\mathcal{D} = \{d_a : a \in \mathcal{A}\}$. Then:*

(i) *coordinate-wise acceptance:*

$$\{s : \exists a \in \mathcal{A}, d_a \leq s\} = \bigcup_{a \in \mathcal{A}} (d_a + \mathbb{R}_+^n);$$

(ii) *scalarized acceptance:*

$$\{s : \forall w \in \mathbb{R}_+^n \setminus \{0\} \exists a \in \mathcal{A}, w \cdot d_a \leq w \cdot s\} = \text{conv}(\mathcal{D}) + \mathbb{R}_+^n;$$

hence the acceptance gap is

$$[\text{conv}(\mathcal{D}) + \mathbb{R}_+^n] \setminus \bigcup_{a \in \mathcal{A}} (d_a + \mathbb{R}_+^n),$$

and replacing the menu by its convex hull — admitting every convexified plan — collapses the gap to the empty set.

Proof. See Appendix A. On the witness datum of Section 4.5 the margin class is $\{\text{FAST}, \text{SLOW}\}$ with $D = \{(2, 0), (0, 2)\}$, so $\text{conv}(\mathcal{D}) + \mathbb{R}_+^2 = \{s_1 + s_2 \geq 2, s_1, s_2 \geq 0\}$ and $\bigcup_a (d_a + \mathbb{R}_+^2) = \{s_1 \geq 2\} \cup \{s_2 \geq 2\}$: the gap is the discrepancy triangle $\mathcal{Q}$. Theorem 9 is the closure statement of this theorem applied to the blend family, and Theorem 5 its instantiation.

4.5. The witness datum

The witness datum of this paper is the explicit typed transition datum on which the separation of Theorem 5 is exhibited. Two architectures — extraction ($q = 0$) and regeneration ($q = 1$) — share one review interval $[0, 1]$, with phase state (q, x, s_1, s_2) : a physical reserve stock x and two typed floors s_1 (protected-group service surplus) and s_2 (remediation-liability coverage), both normalized to 0. The transition-safe set is $S_0 = \{x \geq 0, s_1 \geq 0, s_2 \geq 0\}$; the destination set is $G = \{(1, x, s) : x \geq 0, s \geq 0\}$, with destination maintainability witnessed by the destination hold policy. The destination reset applies the gain vector $e = (1/4, 1/4)$ componentwise to both typed floors; e is strictly positive, so successors are interior to G , and no inequality of Theorem 5 binds on its magnitude. The rescue cost is $c = 1$ (STAGED spends one unit of x).

Disturbance convention (action-indexed). The disturbance set is $\{\beta, \alpha\}$, where α triggers the worst-case dip — of fixed depth 2, per the action table below — in the *active action's characteristic coordinate* (s_1 for FAST, s_2 for SLOW), and β is the analogous disturbance applied to the

other coordinate; the two disturbances are never required to hit a coordinate the action does not move. Each action’s worst-case tube below is the tube under its own characteristic dip, which is the worst case for that action’s constraint. A disturbance label shared across the two actions still takes effect only on the coordinate each action moves, so FAST’s s_2 -tube and SLOW’s s_1 -tube are unchanged, and the distinction does not affect the separations of Theorem 5.

The four actions available from any initial state $(0, x, s)$ with $x \geq 0, s \geq 0$, written as piecewise-linear maps on the breakpoints $\{0, \frac{1}{2}, 1\}$, monotone on each piece (so every tube is the exact visited set — there is no outer approximation):

Action	Within-interval trajectory (worst case)	Successor
NO-SWITCH	state constant	$\{(0, x, s)\}$ — misses G
FAST	$s_1(t) = s_1 - 4t$ on $[0, \frac{1}{2}]$, $s_1 - 4(1 - t)$ on $[\frac{1}{2}, 1]$; s_2, x constant	$\{(1, x, s + e)\}$
SLOW	$s_2(t) = s_2 - 4t$ on $[0, \frac{1}{2}]$, $s_2 - 4(1 - t)$ on $[\frac{1}{2}, 1]$; s_1, x constant	$\{(1, x, s + e)\}$
STAGED	$x(t) = x - t$ on $[0, 1]$; floors grow to $s + e$	$\{(1, x - 1, s + e)\}$

The worst-case tube of FAST is thus $[s_1 - 2, s_1]$ in the s_1 coordinate with s_2 and x constant; the worst-case tube of SLOW is symmetric; the worst-case tube of STAGED is $[x - 1, x]$ in the x coordinate with both floors nondecreasing.

4.6. Witnessed separation

Presentation note. Theorem 6 of Section 4.4 is stated before the present theorem but carries the larger number: this paper’s statements share one counter assigned in citation order, so the numbering is non-sequential in presentation order. No content depends on the ordering; Theorem 6 is the general finite-menu form and Theorem 5 its witness instantiation.

Theorem 5 (witnessed separation). *On the witness datum of Section 4.5, over initial states $X_0 = \{(0, x, s) : x \geq 0, s \geq 0\}$, the following hold.*

- (1) $\mathcal{V}_{\text{typ}} = \{x \geq 1\} \cup \{s_1 \geq 2\} \cup \{s_2 \geq 2\}$.
- (2) $\mathcal{V}_{\text{weak}} = \bigcap_{w \in W_+} \mathcal{V}_w = \{x \geq 1\} \cup \{s_1 + s_2 \geq 2\}$.
- (3) $\mathcal{V}_{\text{phys}} = \mathcal{V}_{\text{end}} = X_0$.

Define: - $\mathcal{Q} = \{s_1 < 2, s_2 < 2, s_1 + s_2 \geq 2\}$ — the aggregate-versus-direct-floor discrepancy region, with x unrestricted. This is **not** entirely a false-positive set, because for $x \geq 1$ the STAGED action is typed-admissible. - $R = \mathcal{Q} \cap \{x \geq 1\}$ — the **rescue set**: the already-typed slice of the discrepancy region, whose states are typed-transformable at cost $c = 1$, witnessed by STAGED. - $I = \mathcal{Q} \cap \{x < 1\}$ — the **impossibility region**. - $\text{FP}_{\text{agg}} = \mathcal{V}_{\text{weak}} \setminus \mathcal{V}_{\text{typ}}$ — the genuine compensatory-versus-noncompensatory acceptance gap, relative to the typed criterion and the specified menu.

(4) $\text{FP}_{\text{agg}} = I$. *The genuine acceptance gap is exactly the impossibility region; it is nonempty, and its relative interior (in the $q = 0$ slice) is $\{0 < x < 1, 0 < s_1 < 2, 0 < s_2 < 2, s_1 + s_2 > 2\}$. The rescue set R is a subset of $\mathcal{V}_{\text{typ}}$, not of the gap: it is the part of $\mathcal{Q}$ that is not a false positive.*

(5) *Both hierarchy inclusions are strict: every point of I lies in $\mathcal{V}_{\text{weak}} \setminus \mathcal{V}_{\text{typ}}$, and the point $(x, s_1, s_2) = (\frac{1}{2}, \frac{1}{10}, \frac{1}{10})$ lies in $\mathcal{V}_{\text{phys}} \setminus \mathcal{V}_{\text{weak}}$.*

(6) Per-weight plan disagreement. On the relative interior of $\mathcal{Q}$ (where $s_2 > 0$ and $s_2 < 2$), writing $r = w_2/w_1$, the FAST-certifying weight ratios are exactly $\{r \geq \rho_1\}$ and the SLOW-certifying ratios exactly $\{r \leq \rho_2\}$, with

$$\rho_1 = \frac{2 - s_1}{s_2}, \quad \rho_2 = \frac{s_1}{2 - s_2}, \quad \rho_2 \geq \rho_1 \iff s_1 + s_2 \geq 2.$$

High- s_2 -weight assessors ($r > \rho_2$) license FAST only; low- s_2 -weight assessors ($r < \rho_1$) license SLOW only; intermediate assessors license both. On I no single action serves every weight — $\bigcap_w E_w(z) = E_{\text{typ}}(z) = \emptyset$ by Proposition 3(ii) — while on R the STAGED action is typed-admissible and hence serves every weight.

(7) The rescue split. R is typed-transformable, witnessed by STAGED: bridging at physical cost $c = 1$ keeps both floors intact and lands in G . I is aggregate-feasible for every cone weight yet admits **no** typed-admissible action, with four exhibited violations: FAST violates the s_1 floor under the adverse disturbance; SLOW violates the s_2 floor; STAGED drives x negative; NO-SWITCH misses G .

Proof. See Appendix B.

Boundary conventions. The formulas for ρ_1, ρ_2 require $s_2 > 0$ and $s_2 < 2$. On the boundary faces ($s_2 = 0$, $s_2 = 2$, $s_1 = 0$, $s_1 = 2$, $s_1 + s_2 = 2$) the classification follows directly from the closed-set identities of (1)–(2); no interior formula is applied there.

Remark. The compensatory assessment’s binding condition is the total budget $s_1 + s_2 \geq 2$. The noncompensatory assessment requires either one floor to survive its own worst-case dip ($s_i \geq 2$) or the bridge stock to cover the rescue cost ($x \geq 1$). The region between the two conditions is where the compensatory doctrine certifies — per weight, with weight-dependent plans — transitions the noncompensatory doctrine rejects.

Example (the witness point). At $(x, s_1, s_2) = (\frac{1}{2}, \frac{6}{5}, \frac{6}{5})$, the interior witness of Figure 1, Panel A, the typed verdict is rejection: STAGED’s x -tube reaches $\frac{1}{2} - 1 = -\frac{1}{2} < 0$; FAST’s worst-case s_1 dips to $\frac{6}{5} - 2 = -\frac{4}{5} < 0$; SLOW’s worst-case s_2 dips to $-\frac{4}{5} < 0$; NO-SWITCH misses G . Yet $s_1 + s_2 = \frac{12}{5} \geq 2$, so the state is in $\mathcal{V}_{\text{weak}}$, with thresholds $\rho_1 = (2 - \frac{6}{5})/(\frac{6}{5}) = \frac{2}{3}$ and $\rho_2 = (\frac{6}{5})/(2 - \frac{6}{5}) = \frac{3}{2}$. FAST is licensed exactly for $r \geq \frac{2}{3}$ and SLOW exactly for $r \leq \frac{3}{2}$; the intervals overlap, so every weight ratio $r \geq 0$ licenses at least one action, but the licensed action depends on r : $r = 2$ (heavy on s_2) licenses FAST only, $r = \frac{1}{2}$ licenses SLOW only, $r = 1$ licenses both. The aggregate thus certifies this state under every weighting while no single plan passes every weighting. By contrast, $(x, s_1, s_2) = (\frac{1}{2}, \frac{1}{10}, \frac{1}{10})$ has $s_1 + s_2 = \frac{1}{5} < 2$ and $x < 1$, so it fails even the weak test while remaining in $\mathcal{V}_{\text{phys}}$ — the point that makes the second inclusion strict. At $(x, s_1, s_2) = (1, \frac{6}{5}, \frac{6}{5})$, $x \geq 1$ with the same floors, STAGED’s x -tube is $[0, 1]$, the floors grow, and the successor $(1, 0, s + e) \in G$, so the state is typed-accepted, as R promises.

Table 3 records the two verdicts region by region.

4.7. The separation in three figures

The three figures of this section visualize the separation on the witness datum: the discrepancy triangle in the floor plane, the weight-interval view at the witness point, and the path view under the worst-case disturbance.

4.8. Propagation under hold-prefix extension

This subsection asks whether the one-interval separation survives when the horizon is lengthened by prepending hold intervals. Two results delimit the answer. Remark 7 establishes persis-




Figure 1: The discrepancy triangle $\mathcal{Q} = \{s_1 < 2, s_2 < 2, s_1 + s_2 \geq 2\}$ in the floor plane (boundary conventions as in Section 4.6). **Panel A** ($x < 1$): $\mathcal{Q}$ is the impossibility region $I = \text{FP}_{\text{agg}}$ — every weighting licenses a plan, no plan satisfies the floors — and the worst-case dip arrows show FAST ($s_1 \rightarrow s_1 - 2$) and SLOW ($s_2 \rightarrow s_2 - 2$) crossing the zero floors from the witness point $(6/5, 6/5)$. **Panel B** ($x \geq 1$): the same triangle is the rescue set R , served by STAGED. Below the diagonal $s_1 + s_2 = 2$, both assessments reject when $x < 1$ and both accept when $x \geq 1$. The reserve bars show the bridge stock x against the rescue cost $c = 1$ (Panel A: $x = 0$; Panel B: $x = 1$). The bottom strip records the endorsement zones at the witness point: SLOW-only for $r < \rho_1 = \frac{2}{3}$, both for $\rho_1 \leq r \leq \rho_2 = \frac{3}{2}$, FAST-only for $r > \rho_2$. The reserve bars mark representative slice values (Panel A: $x = 0$, the impossibility-slice representative; Panel B: $x = 1$, at the rescue threshold); the dip arrows are x -independent (the floor dynamics do not involve the reserve), and the canonical witness point of Sections 4.6 and 6.3 sits at $x = \frac{1}{2}$, between the drawn slices.

Region (witness datum)	Aggregate (per-weight) verdict	Common-plan (typed) verdict	Why
$x \geq 1$, all s	accept	accept	STAGED serves every weight
$s_1 \geq 2$ or $s_2 \geq 2$	accept	accept	FAST/SLOW survives its own dip
I : $x < 1$, $s_1 < 2$, $s_2 < 2$, $s_1 + s_2 \geq 2$	accept (weight-dependent plans)	reject	dips breach floors; reserve cannot finance phasing
R : $x \geq 1$, $s_1 < 2$, $s_2 < 2$, $s_1 + s_2 \geq 2$	accept	accept	reserve finances STAGED
$s_1 + s_2 < 2$, $x < 1$	reject	reject	even the index fails; both doctrines agree

Table 3: Region dictionary on the witness datum (Theorem 5): the per-weight aggregate verdict and the common-plan (typed) verdict on each salient region of X_0 .




Figure 2: Weight-interval view at the witness point $(x, s_1, s_2) = (\frac{1}{2}, \frac{6}{5}, \frac{6}{5})$. FAST is licensed exactly for $r \geq \rho_1 = \frac{2}{3}$ and SLOW exactly for $r \leq \rho_2 = \frac{3}{2}$. The two intervals together cover every weight ratio $r \geq 0$, but neither covers it alone: every weight licenses a plan, no plan is licensed by every weight. On the rescue set, STAGED supplies the full interval and closes the gap.




Figure 3: Path view at the witness point. Rows: FAST, SLOW, STAGED under the worst-case disturbance; solid curves are the constrained paths, the dashed line the zero floor, and the colored curves the scalarized index for $r \in \{2, 1, \frac{1}{2}\}$. Under FAST (top row) the index stays nonnegative for $r = 2$ and $r = 1$ while s_1 breaches its floor; under SLOW (middle row) the mirror holds. The only plan that protects both floors (bottom row, STAGED) consumes the reserve linearly: it is infeasible at $x = \frac{1}{2}$ and feasible at $x = 1$, which is what separates the impossibility region from the rescue set.

tence under a precisely specified hold-prefix extension; Theorem 8 exhibits a two-stage datum on which the gap is erased, showing the persistence claim is architecture-conditional.

Remark 7 (propagation). *Extend the witness datum to $m \geq 2$ review intervals by prepending hold intervals (sole action HOLD: constant tube $\{z\}$, successor $\{z\}$, safe set S_0 ; the final interval carries the witness). Assume:*

- (H1) *HOLD is available at every prefix state;*
- (H2) *the hold tube is exactly $\{z\}$;*
- (H3) *the prefix safe sets contain the witness states;*
- (H4) *terminal and successor architecture labels are unchanged;*
- (H5) *no additional reset or disturbance branches are introduced.*

Then:

(i) *Hierarchy. For every stage j ,*

$$\mathcal{V}_j^{\text{typ}} \subseteq \bigcap_{w \in W_+} \mathcal{V}_j^w \subseteq \mathcal{V}_j^{\text{phys}}.$$

This inclusion holds for every multi-interval typed exact-tube datum; it is constraint monotonicity under backward induction.

(ii) *Persistence of strictness. The stage-0 accepted regions are the witness regions pulled back through the holds, so both strictness witnesses of Theorem 5 persist under this hold-prefix extension.*

(iii) *The displayed separation is not an artifact of the one-interval framing for the constructed hold-prefixed datum. Strictness for arbitrary multi-stage systems does not follow without a separate persistence theorem: a later stage can erase the gap if every aggregate-feasible state has a common typed-safe continuation, if the weak and typed terminal sets coincide on the reachable subset, or if actions couple across stages.*

Proof. See Appendix B.

Theorem 8 (erasure witness). *The propagation claim of Remark 7 is architecture-conditional: there is a two-interval datum whose single-interval truncation exhibits the acceptance gap and whose two-stage extension erases it, through the terminal-coincidence mechanism of Remark 7(iii).*

Proof. See Appendix B.

The witness datum of Section 4.5 is not subject to this erasure. Its gap is tube-driven: on I every member of the four-action menu violates the stage-1 path constraint S_0 itself (the two dips and the bridge deficit of Theorem 5(7)), and a later stage cannot repair a violated tube — the stage-1 constraint is local to stage 1. Remark 7’s safe direction is secured by exactly the assumptions the erasure datum violates (unchanged terminal architecture, no new branches). The two results together delimit the persistence claim: hold-prefix extension preserves the separation, arbitrary multi-stage extension does not, and the erasure mechanisms of Remark 7(iii) are realizable.

4.9. Machine verification

Proposition 3, Proposition 4, Theorem 5, Theorem 6, Remark 7, Theorems 8–9, and Proposition 10 are proved in Appendices A–C, and Proposition 11 in Appendix C. An accompanying software artifact (deterministic exact-integer arithmetic; no floating point, tolerances, or randomness) checks the finite rational instance of Proposition 3, Proposition 4, Theorem 5, and Remark 7: the action

classifications, the region identities, and the accepted-set identities on a $31^3 = 29,791$ -state grid over (x, s_1, s_2) , with a finite verification set containing the critical weight ratios ρ_1, ρ_2 and their midpoint for every enumerated grid state. All 25 checks pass and re-execution reproduces them exactly; they are enumerated in the Supplementary Material (S8). This grid verifier is distinct from the software companion’s separate 24-check benchmark suite, which re-derives the Section 6.3 fishery instantiation; the two counts index different check lists, not a disagreement.

Claim layer	What is established	By what
Continuum statements	Exact identities and inequalities on the witness datum	Appendices A–C
Finite rational instance	Classifications on the enumerated grid	Deterministic computation
Empirical claims	None in this article	—

The finite grid does not by itself prove the continuum identities; it validates the symbolic classification on the enumerated instance.

A third warrant exists alongside the appendices and the grid. The appendix proofs of Theorem 5 (1)–(7), Remark 2 with Proposition 3, Theorem 9 (i) and (iii) with Proposition 10, Lemma A (the handshake identity and the engine equivalence), Lemma B, the master-equation reductions, and Theorem S2 are formalized in a mechanized layer (Lean 4; no axioms, no admitted gaps, no computer-algebra oracle), every statement being derived from the ordered-field interface alone. The layer and the appendices were produced independently and agree; neither is a substitute for the other, and the grid check is a third and separate thing again — it validates the symbolic classification on one enumerated instance, which is all a finite check can do.

4.10. Menu convexification: the blend family

The action space is the finite menu $\{\text{FAST}, \text{SLOW}, \text{STAGED}, \text{NO-SWITCH}\}$ of deterministic actions; fractional or alternating policies are not members of it. The convexification instrument is the blend family: for each $\delta \in (0, 1)$, the action BLEND_δ is the *convexified action* obtained by fractional allocation of the primitive control flows, $u_\delta = \delta u_{\text{FAST}} + (1 - \delta) u_{\text{SLOW}}$; by linearity of the flow map its worst-case tube is the pointwise convex combination of the two primitive worst-case tubes and its successor the convex combination of the two successors. This is menu convexification at the level of control inputs — an explicit model assumption, not the time-sharing of discrete actions in alternation.

A limitation to state before the theorem. The blend family is an explicitly defined convexified control map — an instrument of the abstract witness datum, not by itself a claim that fractional plans are available anywhere. Convexification must be demonstrated from the model equations (as here, by linearity of the flow map in the control) or carried explicitly as a management assumption; and a blended control requires physical and institutional implementability — fractional allocation of effort, quota, or seasonal resolution, with monitoring and enforcement — before Theorem 9 is applied to a real plan menu. The theorem is a statement about certification relative to the convexified menu; the menu’s availability is external data (the fourth design test of Section 5.5 and the instruments of Section 5.6).

Theorem 9 (blend collapse). *On the witness datum of Section 4.5, over initial states X_0 :*

(i) BLEND_δ is typed-admissible at z exactly when

$$\delta \in \left[1 - \frac{s_2}{2}, \frac{s_1}{2}\right] \cap [0, 1],$$

an interval nonempty exactly when $s_1 + s_2 \geq 2$ (a singleton at equality), independent of $x \geq 0$ and of the weight; the intersection with $[0, 1]$ is the convex-combination domain of BLEND_δ and is inactive on $\mathcal{Q}$, where $0 < 1 - s_2/2 \leq s_1/2 < 1$ automatically.

(ii) The typed acceptance set of the menu augmented by the blend family is exactly the compensatory region:

$$\mathcal{V}_{\text{typ}}^{\text{blend}} = \{x \geq 1\} \cup \{s_1 + s_2 \geq 2\} = \mathcal{V}_{\text{weak}}.$$

Convexification of the menu collapses the acceptance gap: FP_{agg} vanishes under the augmented menu, and the blend family reaches exactly the compensatory region and nothing beyond — on $\mathcal{V}_{\text{phys}} \setminus \mathcal{V}_{\text{weak}}$ no blend is typed-admissible.

(iii) The admissible window depends on the state through (s_1, s_2) alone: a single δ serves every weight $w \in W_+$. On the impossibility region the blend is therefore the common action that no member of the finite menu supplies — the acceptance gap is a property of the deterministic menu, not of the doctrines.

Proof. See Appendix B.

Remark. The collapse is exact and weight-independent: a single convexified plan serves every assessor on I , where the per-weight plans of Theorem 5(6) necessarily differ. The nonconvexity at work is menu geometry — the finiteness and determinism of the action space — not the Pareto-frontier geometry under which weighted sums fail to reach nonconvex frontier parts (Das and Dennis, 1997); the two mechanisms are distinct. The structure is precisely that of pure-versus-mixed strategies in the assessor–planner game (von Neumann, 1928; Sion, 1958): with the planner picking an action, the assessor a weight, and the payoff the worst-case aggregate, common-plan acceptance is the pure-strategy value and per-weight acceptance the mixed value, the gap being the pure-strategy duality gap that the convex (mixed) extension closes. The two-player reach/avoid formulation of such games is the discriminating kernel of Cardaliaguet (1996), with numerical treatment in Cardaliaguet, Quincampoix, and Saint-Pierre (1999). This is also the here-and-now versus wait-and-see separation of adjustable robust optimization (Ben-Tal et al., 2004).

A class-conditional caveat on the blend collapse. Theorem 9’s window is computed under the action-indexed disturbance convention of Section 4.5, and it is class-conditional: under the common-shock class of Section 5.9 the window is nonempty exactly on $\{s_1 \geq \delta, s_2 \geq \delta, s_1 + s_2 \geq 3/2 + 2\delta\}$ (at the datum’s event depth $\delta = 1/2$: $\{s_1 \geq 1/2, s_2 \geq 1/2, s_1 + s_2 \geq 5/2\}$, against the action-indexed $\{s_1 + s_2 \geq 2\}$ of Theorem 9(i)). The common-shock acceptance gap — which survives the shared shock (Theorem M1, Section 5.9) — is therefore *not* blend-closed on the fragile band $s_1 + s_2 \in [2, 5/2]$: the witness $(\frac{1}{2}, 1, 1)$ is a common-shock gap state with *no* admissible blend, while every common-shock gap state with $s_1 + s_2 \geq 5/2$ has one (Theorem M3 there; the coupled variant, by contrast, is blend-closed — the fragility is class-conditional). An agency that would otherwise invoke Theorem 9’s collapse at a shallow common-shock gap state cannot: the blended menu itself is inadmissible there.

4.11. *The converse: discrete time-sharing does not erase the gap*

Proposition 10 (sequential time-sharing does not erase the gap). *On the witness datum of Section 4.5, let a sequential time-shared policy alternate between FAST and SLOW over*

the interval, so the set of visited states is the union of the two primitive worst-case tubes. Then such a sequential policy is typed-admissible if and only if $s_1 \geq 2$ and $s_2 \geq 2$; on the impossibility region I (where $s_1 < 2$, $s_2 < 2$, $s_1 + s_2 \geq 2$) no such policy is typed-admissible, and the acceptance gap therefore survives discrete time-sharing.

Proof. See Appendix B.

Remark. Together Proposition 10 and Theorem 9 delimit the notion of mixing precisely: continuous convexification — fractional allocation of control flows — closes the acceptance gap exactly at the compensatory region $s_1 + s_2 \geq 2$, whereas discrete time-sharing — alternation in time, whose visited set is the union of the action tubes — closes it only where the two dips are independently subsumed, $s_1 \geq 2$ and $s_2 \geq 2$.

On the regime $s_1 + s_2 \geq 2$ with $\min(s_1, s_2) < 2$ (which contains the impossibility region) the gap survives sequential time-sharing and closes only under convexification, so the structural character of the gap is a property of the *convexity of the action space*, not of the assessment doctrine and not of temporal sharing as such. A terminological precision is required here. In the language of relaxed controls, the collapse theorem is a statement about the convexified (relaxed) action program. The converse is a statement about discrete, plan-granular alternation at full strength: each primitive plan is held for a whole sub-interval, so the visited set is the union of the primitive tubes. That stipulated visited set is the pointwise worst case over alternation schedules — a literal rate-preserving fine-grained alternation visits sub-ranges of it — so the impossibility established here is scoped to the worst-case alternation. It is not a statement about the chattering limit of classical relaxed-control theory (Filippov; Warga): infinitesimally fast switching recovers the convexified trajectory and is therefore subsumed by Theorem 9, not distinct from it. What fails to erase the gap is alternation at finite rate and plan granularity, not rapid switching.

5. Interpretation

5.1. The doctrinal reading

The scalarized operators $\{E_w\}_{w \in W_+}$ formalize **one** compensatory assessment doctrine. The doctrine comprises: a single aggregate index $w \cdot s$; nonnegative scalarization weights on capital forms; substitution across floors permitted at those weights; and disturbances respected. We refer to this as *the scalarized aggregate doctrine as formalized here*, not as “weak sustainability” simpliciter. The weak-sustainability literature is broader than this operator: it encompasses aggregate wealth, constant total capital, nondeclining comprehensive consumption, substitutability assumptions, discounting, shadow prices, and intertemporal investment conditions (Neumayer, 2013; World Bank, 2011; Boos, 2015; Dasgupta and Mäler, 2000; Asheim, 1994). Likewise, strong sustainability is not always equivalent to requiring every stock coordinate to remain nonnegative at every instant. Critical-natural-capital frameworks may use thresholds, safe operating spaces, irreversibility, resilience, or minimum-service conditions (Ekins et al., 2003; Rockström et al., 2009; Raworth, 2012; Doyen and Gajardo, 2020). The typed operator E_{typ} is a formal idealization of the separately-binding-floor reading, in the same sense that the scalarized operators idealize one aggregation doctrine. The typed operator is the coordinate-wise criterion — each floor is required to hold separately — and the scalarized family is the compensatory criterion; the two vocabularies name the same operators, with *typed* denoting the framework object and *coordinate-wise* its constraint structure.

Within this precise scope, Theorem 5 reads as follows. The two formalized doctrines can disagree on the same transition datum, with the same robustness standard and the same action set, in the direction compensatory-accepts / noncompensatory-rejects, on a relatively open region. The disagreement is not an artifact of one bad weight: every weight accepts, each licensing a different physical transition. The plans are genuinely different transitions (FAST and SLOW violate different floors at different times), which is the dynamic formalization of substitution across floors. By Proposition 3(ii), the precise seat of the disagreement is the noncommutativity of “choose an action” with “for all weights.”

At the static level the full-cone aggregate is lossless (Remark 2), so the compensatory doctrine’s limitation is not the existence of an aggregate index but the *policy dependence* of the aggregate-feasible transition: the index certifies a set of transitions, none of which the noncompensatory assessment accepts. The two quantifier orders also carry an information reading. $\exists a \forall w$ is a commitment made before the weight is known — the strong doctrine’s robustness demand — while $\forall w \exists a_w$ lets the action be chosen after the weight is observed. On the witness, the acceptance gap between the two orders measures the value of information about the assessment weight. Theorem 9 qualifies this reading: a single convexified blend serves every weight on the impossibility region, so the value of information reflects the finite deterministic menu rather than the weight information itself. Proposition 10 sharpens the point: the value of information is removed by convexifying the action space, not by sharing actions in time — under discrete time-sharing the gap (and hence the value of information about the weight) survives on the impossibility region.

Throughout, we use “nonnegative scalarization weights” for elements of W_+ and reserve the word “prices” for the interpretive discussion. Actual prices may be strictly positive, endogenous, dynamically determined, state-dependent, or dimensionally heterogeneous, and the theorems require none of those properties. The operators are linear scalarizations with weights fixed over the review interval. Nonlinear aggregate indices — CES-type substitutability or endogenous, state-dependent weighting of the kind that diverges near zero-stock boundaries — are different assessment operators, and the theorems claim nothing for them. Section 5.8 begins the programme of establishing what can be claimed for the CES family on the witness datum.

5.2. Positioning against established theory

The backward recursion of Section 3 is a typed instance of established robust-predecessor, reach-avoid, capture-basin, and hybrid-reachability constructions (Aubin, 1991; Aubin, Bayen, and Saint-Pierre, 2011; Saint-Pierre, 1994; Lygeros, Tomlin, and Sastry, 1999); the fixed-point and algebraic character of these objects is developed in Aubin and Catté (2002). Viability methods are likewise developed in natural-resources management, with fisheries a canonical application (De Lara and Doyen, 2008; Doyen et al., 2012). Proposition 3(i) is constraint-set monotonicity of viability kernels and reachability sets (Aubin, 1991; Frankowska, 1989). Static scalarization limitations — weighted sums that cannot reach nonconvex parts of Pareto fronts (Das and Dennis, 1997) — act through frontier geometry under a single optimization, a mechanism distinct from the action-quantifier mechanism established here. The per-literature details — the commensurability foundation (Martinez-Alier, Munda, and O’Neill, 1998), the maximin and indicator formulations of strong sustainability (Martinet, 2011; Cairns and Martinet, 2014; Doyen and Gajardo, 2020), and the compensability mapping in multi-criteria decision analysis — including the outranking approach PROMETHEE, whose compensatory properties the cited analysis dissects (Cinelli, Coles, and Kirwan, 2014; Schär, Pohl, and Geldermann, 2025) — and the acceptance-semantics reading

of the protocols as decision-aiding sorting rules (Roy, 1996; Vincke, 1992) — are collected in the Supplementary Material (S10).

For multi-criteria decision analysis, the protocols of Table 1 are acceptance semantics: they specify the rule by which a composite reading sorts a transition into the accepted or rejected category — the sorting problematic of decision aiding (Roy, 1996; Vincke, 1992). The separation is then a statement about which semantics a composite index can support: under weight-adaptive acceptance a transition is accepted when every weighting licenses some plan, possibly a different plan for each weighting, and Theorem 5 exhibits transitions that this rule accepts and the coordinate-wise rule rejects on an open region. The licensing thresholds of Theorem 5(6) give the verdict’s exact weight-sensitivity intervals — the preference ranges over which each plan is certified — the acceptance-side analogue of the sensitivity analyses applied to composite-indicator rankings (Nardo et al., 2008). Static compensability analyses — the compensatory potentials of multi-criteria methods for sustainability assessment (Cinelli, Coles, and Kirwan, 2014) and the partial compensability of the outranking approach (Schär, Pohl, and Geldermann, 2025) — classify aggregation behaviour at a fixed profile; the separation established here acts one level up, in the acceptance protocol over a transition, and arises even where every static level of the index is reported correctly (Remark 2).

Against this background, the finite-menu characterization of Section 4.4 is the general form, of which the witness separation (Theorem 5) and the blend collapse (Theorem 9) are the instances.

To our knowledge, no explicit exact-tube witness for this dynamic separation in a compensatory/noncompensatory assessment setting is reported in the literature. We do not assert priority over the elementary quantifier fact $\exists a \forall w \neq \forall w \exists a_w$, which is standard. Whether an equivalent dynamic separation appears in adjacent literatures (multi-objective robust control, viability theory, multi-criteria decision analysis) is not established either way.

5.3. A translation guide for readers from neighbouring fields

Because the paper sits at a confluence — multi-criteria decision analysis and operations research, control theory and viability, environmental modelling and governance — Table 4 translates its central objects into each neighbouring literature’s working vocabulary. The readings are exact correspondences of mathematical role, not analogies: the theorems are venue-neutral, and a reader may enter through any column.

5.4. Scope delimitations

- **No separation on every datum.** Where a single action is safe for all weights, the assessments coincide. The theorem is an existence separation with an open region, plus the always-valid hierarchy and localization.
- **No infinite-horizon, stochastic, partial-observation, or endogenous-event extension.** All statements are for finite-horizon exact-tube data with specified disturbance sets.
- **No claim that the full nonnegative cone is the only reasonable weight family.** Proposition 4 shows that restricting the family enlarges the compensatory accepted set, so the full cone is the strictest compensatory reading, and the separation persists a fortiori for restricted families.
- **No welfare claim about prices.** The weights model an assessment doctrine, not a normative endorsement.
- **No empirical transfer.** The theorems concern assessment operators on a specified datum and imply no empirical claim about any resource system.

Table 4: Terminology map: this paper’s objects read from three neighbouring vocabularies.

	MCDM / operations research	Control / viability theory	Environmental modelling / governance
Typed floors $s_i \geq 0$, enforced path-wise	coordinate-wise (vector) criteria; per-criterion lower bounds along the whole path	state constraints; safe set $\mathcal{V}$; path (tube) constraints	planetary boundaries; safe operating spaces; regulatory floors (e.g., a biomass limit reference point)
Scalarized operator E_w	weighted-sum acceptance; weight space W_+	scalarized robust constraint on an aggregate output	composite sustainability index; the assessment dashboard
Licensing thresholds ρ_1, ρ_2	weight intervals of acceptance; weight-space robustness of a recommendation	parameter-dependent feasibility margins	indicator weighting schemes; stakeholder weight elicitation
Acceptance gap FP_{agg}	scalarization error: accepted by some weighted sum, infeasible coordinate-wise	gap between the scalarized program and the viability kernel	the aggregation illusion: an index certifying what the floors reject
Rescue set R ; threshold κ^*	alternatives repairable by a budgeted side payment	reachability under an auxiliary resource input	transition finance: adjustment funds, fleet buy-backs
Impossibility region I	robust infeasibility: rejected at every weighting	empty common safe-action set over the belief; kernel exclusion	hidden collapse zone: index nonnegative, floor crossed mid-transition
Blend collapse (Thm 9) vs. time-sharing (Prop 10)	randomized (mixed) strategies vs. deterministic sequencing	relaxed (convexified) controls vs. plan-granular alternation (not the chattering limit)	blended policy portfolios vs. alternating policies
Exact-tube semantics	robust path-wise feasibility over a scenario set	viability tubes; worst-case reachability	avoiding transient breaches (e.g., extinction vortices during a transition)

- **No claim that $\mathcal{V}_{\text{weak}}$ is the genuine-savings or inclusive-wealth criterion.** The operator is the strictest compensatory reading (every weight, each licensing an action). The genuine-savings literature’s criterion is a different object, and the witness’s reserve stock is excluded from the weighted aggregate by construction (Section 4.5), not by consequence of aggregation. If the reserve entered the aggregate with positive weight, the gap could shrink or disappear; the separation isolates the case in which the aggregate measures the service floors while the reserve is a separate policy instrument.
- **No absolute impossibility.** The impossibility region is a negative certificate relative to the specified four-action menu; with a larger menu it can shrink or vanish (see the completeness requirement in Section 5.7).
- **No universal doctrinal ranking.** The inclusion $\mathcal{V}_{\text{typ}} \subseteq \mathcal{V}_{\text{weak}} \subseteq \mathcal{V}_{\text{phys}}$ is a theorem about the defined operators under common action and disturbance classes, common safe-set inclusion, common horizon, common tube semantics, and common terminal condition. It does not establish a universal ranking of endpoint, weak, and strong sustainability doctrines.
- **No separation under coupled all-floor shocks.** The witness disturbance convention is action-indexed (Section 4.5): each action’s worst case is its own characteristic dip. A datum whose disturbance class couples simultaneous dips across all floors of every action degenerates the per-weight licensing structure — the principal actions fail together and the compensatory/noncompensatory divergence collapses into universal rejection. The exact statement, computed in Section 5.9, is reading-conditional and depth-dependent: under the additive reading the rejection is universal on I at the datum’s full magnitude, with the exact depth boundary $\delta = 5/4$ below which deep- I states survive (witness $(\frac{1}{2}, \frac{19}{10}, \frac{19}{10})$ at heatwave magnitude) and with the separation relocating — not vanishing — at full magnitude (the coupled gap $\{x < 1, 2 \leq s_1 < \frac{7}{2}, 2 \leq s_2 < \frac{7}{2}, s_1 + s_2 \geq \frac{11}{2}\}$, witness $(\frac{1}{2}, 3, 3)$); under the replace reading the principal actions fail together under one weight-independent requirement and what collapses is the typed/weak distinction itself ($\mathcal{V}_{\text{weak}} = \mathcal{V}_{\text{typ}}$, Theorem CT there), with acceptance — not rejection — surviving at $\{s_1 \geq 2, s_2 \geq 2\} \cup \{x \geq 1, \min(s_1, s_2) \geq \frac{15}{8}\}$, a strictly more demanding reading (the witness $(\frac{1}{2}, 3, 1)$ is action-indexed typed-accepted yet fully rejected under it). The separation is exhibited for, and scoped to, the specified action-indexed class: the witness datum assumes separable transition risks, with each plan exposing a different floor to its own adverse disturbance. That the coupled regime is not a remote possibility is indicated by evidence that planetary-boundary processes interact and amplify one another (Lade et al., 2020).

The disturbance-class scoping, organized (Table 5): the witness’s action-indexed separable class is one of three regimes; the second’s verdict is the common-shock computation of Section 5.9 (Theorem M1: the gap survives the shared shock, reduced by the exact 1/2-cut, with licensing thresholds unshifted on the surviving region) — a verdict computed, not asserted, exactly as the regime demands; the third is the universal-rejection collapse stated above, refined there into two exact reading-conditional theorems (M4 and CT).

- **Convexification of the menu.** Fractional action policies are not members of the action space, and every stated separation is scoped to the finite deterministic menu. Theorem 9 resolves the convexification question: the blend family collapses the typed acceptance set exactly onto the compensatory region, so the separation is a property of the deterministic menu, not of the doctrines — menu geometry, distinct from the Pareto-frontier geometry of scalarization limits (Das and Dennis, 1997). The associated delimitation (Proposition 10) is

Table 5: Disturbance regimes on the witness datum. The middle row’s verdict is the exact common-shock computation of Section 5.9 (Theorem M1); the third row’s exact statement is reading-conditional (Theorems M4 and CT there).

Regime	Disturbance class	Separation status
(i) Action-indexed separable	each action’s worst case is its own characteristic dip in the coordinate it moves (Section 4.5)	proven: the witness datum exhibits the acceptance gap (Theorem 5)
(ii) Common biomass shock	one shared shock strikes the same floor coordinate of every plan	computed (Theorem M1, Section 5.9): the gap <i>survives</i> the common shock, reduced by the exact 1/2-cut — every action-indexed gap state with $\min(s_1, s_2) < 1/2$ weak-rejects — with licensing thresholds unshifted on the surviving region; the canonical datum remains a member
(iii) Coupled all-floor shocks	simultaneous dips across all floors of every action	reading-conditional (Theorems M4 and CT, Section 5.9): additive — universal rejection exact on I at the datum’s full magnitude, depth-dependent below the exact boundary $\delta = 5/4$, the gap relocating (not vanishing) at full magnitude; replace — the per-weight licensing structure degenerates to one weight-independent requirement, the actions fail together, and the typed/weak distinction itself collapses ($\mathcal{V}_{\text{weak}} = \mathcal{V}_{\text{typ}}$), with acceptance — not common rejection — surviving (this section; Lade et al., 2020, on boundary interactions)

that discrete time-sharing does not effect the same collapse: on the impossibility region the gap survives alternation in time, so the structural character of the separation is due to the convexity of the action space, not to temporal sharing.

5.5. Indicator validity: four design tests

The separation theorem implies that a composite indicator can be statically lossless yet dynamically unsound as a transition-certification device. For indicator developers the theorem converts into four design tests — the questions an indicator’s documentation must answer before the index is used to certify a transition. Each test is a reading of a proved statement; none is a heuristic.

1. **Static losslessness.** Does the aggregate preserve component violations at a fixed trajectory for the declared weight family? At the full cone, yes (Remark 2); for a restricted family, the test is to characterize the component violations the restricted family can miss (Proposition 4). An index that fails this test fails statically; the separation of Section 4 is the further, dynamic failure.
2. **Protocol quantifiers.** Does the certification rule require a common plan ($\exists a \forall w$) or permit weight-adaptive plan selection ($\forall w \exists a_w$)? If the latter — Protocol 2 of Table 1 — the indicator is exposed to the acceptance gap on exactly the region Theorem 5 exhibits.
3. **Path-wise observability.** Is the index evaluated along the full tube, or only at audited endpoints? Endpoint evaluation is the weakest operator of the chain (Section 3.1) and cannot detect transient breaches — mid-interval floor crossings that are invisible at snapshots ($E_{\text{end,typ}}$ versus E_{typ} on the witness datum).
4. **Menu completeness and convexity.** Is the plan menu exhaustive, and are blended (convexified) plans implementable? Negative feasibility verdicts are menu-relative (Section 5.4), and the gap is a property of the deterministic menu that convexification collapses exactly (Theorem 9, Proposition 10); the availability of the convexified menu is external data that the indicator’s documentation must declare.

A consolidated practitioner checklist. For an agency using a composite dashboard to certify a sustainability transition, the four tests and the theorems behind them operationalize as ten audit points:

1. **Object of certification:** state what is being certified — a static state, an endpoint, or the existence of a path-wise safe management plan.
2. **Binding floors:** list the floors that are legally binding, irreversible, or lack verified substitutability; these require noncompensatory (Protocol 4) treatment.
3. **Policy quantifier:** declare whether the approval rule assumes one robust plan or lets the plan adapt to the weighting.
4. **Weight-dependent plans:** if weight-adaptive, report the management plan associated with each relevant weight range (the licensing thresholds ρ_1, ρ_2 of Theorem 5(6) are the pattern).
5. **Path-wise evaluation:** evaluate the floors along the full transition, not only at assessment snapshots (the endpoint operators are the weakest of the chain; Section 3.1).
6. **Disturbance structure:** specify whether shocks are common to all plans, action-indexed, or correlated across floors (Table 5); the separation’s scope rides on the class.
7. **Menu completeness:** acknowledge that negative feasibility results are menu-relative unless the action space is demonstrably exhaustive (Section 5.4).

8. **Blended policies:** do not infer convexification from mathematical convenience; verify that fractional effort or quota allocation is administratively implementable (the limitation stated before Theorem 9).
9. **Rescue margins:** if a resource-controlled action exists, report the minimum increment κ^* that converts impossibility into rescue (Section 5.7).
10. **Mandatory reporting:** require that component trajectories and their minimum margins be reported alongside the aggregate index, so the typed verdict can be independently reconstructed (the minimum report of Section 5.6).

5.6. Policy implications

The separation theorem has four implications.

First. Aggregate indices alone cannot certify noncompensatory transition safety under the present assessment semantics. A scalarized assessment can certify its own criterion — aggregate feasibility; what it cannot certify is typed safety unless an additional bridge theorem is supplied. This follows from Theorem 5: the compensatory accepted set strictly contains the noncompensatory accepted set, and the excess region is where the two criteria diverge.

Second. When the typed recursion identifies an admissible rescue action whose feasibility is controlled by a resource margin, investment in that margin is a candidate remedy; reweighting alone cannot substitute for it on the witness datum (Section 5.7). This applies where a resource-controlled rescue action exists; it is not a universal policy claim.

Third. Per-floor reporting alongside aggregate reporting is necessary for detecting this class of discrepancy. It is not universally necessary for every sustainability judgement; but where the question is whether a transition respects separately-binding floors, an aggregate index alone cannot distinguish the rescue region from the impossibility region. The same separately-checked-floor logic underlies the safe-and-just-space (“doughnut”) literature, which evaluates social thresholds and biophysical ceilings independently, without cross-compensation, and reports that no country satisfies all floors within all ceilings (Raworth, 2012; O’Neill et al., 2018; Fanning et al., 2022). A minimum report for such an assessment should include: (1) typed floors and their provenance; (2) aggregate weights or the weight family; (3) policy quantifiers; (4) the disturbance set; (5) the observation and implementation model; (6) exact versus conservative tube status; (7) terminal maintainability; (8) action-exhaustion status; (9) uncertainty and approximation bounds; (10) normative premises.

Fourth. Reporting regimes that evaluate endpoints only, as in audited annual snapshots, evaluate only the weakest operator of the chain of Section 3.1. Under those semantics they license transitions that violate typed floors mid-interval without detection (on the witness datum, FAST is typed-endpoint-admissible at every state of X_0 but typed-tube-admissible only for $s_1 \geq 2$, by inspection of the action table in Section 4.5); the per-floor reporting of the third implication detects the discrepancy. No empirical claim about any particular reporting regime is made here.

The instruments that act on the acceptance gap close it in sharply different ways; Table 6 records them.

5.7. Rescue as action synthesis, and data requirements

Rescue as action synthesis. Rescue is an operation on the impossibility region. Define a resource-augmentation map

$$\text{Aug}_\kappa : (x, s, \mathcal{A}) \mapsto (x + \kappa, s, \mathcal{A} \cup \{\text{STAGED}_\kappa\})$$

Instrument	Mathematical operation	Effect on the gap
Reweighting	change the weight w	does not close the gap
Restricting the weight family	replace W_+ by $W \subset W_+$	enlarges per-weight acceptance; leaves the common-plan problem unless one plan works for all retained weights
Reserve augmentation	add STAGED when $x + \kappa \geq 1$	closes the gap for states where only the reserve constraint binds
Menu convexification	replace $\mathcal{D}$ by $\text{conv}(\mathcal{D})$ (the blend family)	closes the gap exactly on the scalarized feasible region
Discrete time-sharing	union of the primitive tubes	generally does not close the gap; closes only where both dips are separately survivable

Table 6: Policy instruments acting on the acceptance gap.

and the minimal rescue threshold

$$\kappa^* = \inf \{ \kappa : \exists a \in \mathcal{A}_\kappa, a \in E_{\text{typ}}(z) \},$$

where $\mathcal{A}_\kappa$ denotes the augmented menu. **Proposition 11 (rescue threshold).** *On the witness datum, for every $z \in X_0$,*

$$\kappa^*(z) = (1 - x)_+ \cdot \mathbf{1}[x < 1, s_1 < 2, s_2 < 2] = (1 - x) \cdot \mathbf{1}[z \notin \mathcal{V}_{\text{typ}}].$$

Proof. See Appendix C.

On the impossibility region I , $\kappa^*(z) = 1 - x > 0$; on the rescue set R , and on all of $\mathcal{V}_{\text{typ}}$, $\kappa^*(z) = 0$, matching the observation that states of R need no augmentation. Because STAGED_κ is admissible exactly when $x + \kappa \geq 1$, independently of the floors, the resource route admits every state with $x < 1$ — including states outside $\mathcal{V}_{\text{weak}}$, such as $(\frac{1}{2}, \frac{1}{10}, \frac{1}{10})$ — at cost $1 - x$. The set on which the four-action menu fails yet resource augmentation succeeds at finite cost is the full failure set $\{x < 1, s_1 < 2, s_2 < 2\}$, which strictly contains the impossibility region; I is its part on which the aggregate reading still certifies the transition ($s_1 + s_2 \geq 2$) — the genuine acceptance gap of Theorem 5(4). The increment converts any state of the failure set into a typed-transformable one; on I , it converts an aggregate-licensed impossibility into a rescue.

The class-conditional rescue trichotomy (a labelled extension). Proposition 11’s threshold is computed under the action-indexed disturbance convention of Section 4.5, under which the resource route is unconditionally finite — every failure state is rescuable at cost at most $1 - x$. The common-shock computation of Section 5.9 changes this qualitatively: under the common class (a shared environmental shock of the datum’s heatwave magnitude striking every plan) the threshold becomes a trichotomy (Theorem M2 there) — $\kappa^* = 0$ on the typed accepted set; $\kappa^* = (1 - x)_+$ on the failure states with both floors at or above $3/8$; and $\kappa^* = \infty$ on the failure states with $\min(s_1, s_2) < 3/8$, the first unrescuable-by-reserve failure states on this datum (witness $(\frac{1}{2}, \frac{1}{4}, 3)$: no reserve increment whatsoever rescues it, the shared shock’s exposure of the staged route’s growing floor exceeding what a floor below $3/8$ can carry). The canonical datum’s own shortfall is unchanged at $\kappa^* = 1/2$. The management reading: at an asymmetric gap state the prescription changes from reserve accumulation (bridge to $x = 1$) to floor rebuilding (rebuild

the weaker floor above $1/2$ first, above $3/8$ for the staged route) — different actions with different costs and timescales.

Institutional forms. The rescue margin $\kappa^*(z) = 1 - x$ maps onto concrete transition-finance instruments: license buy-backs and vessel decommissioning (the fund spends $c = 1$ to remove effort), transition assistance protecting the income floor during the staged draw-down, and quota banking held against the staged rebuild. The software companion — the SafeTransition library — realizes exactly this menu on the benchmark datum, as a staged rebuild financed by a fund buy-back with the rescue threshold truncated at zero once the fund covers the cost ($\kappa^* = \max(0, 1 - x)$). The mapping is a reading aid, not a theorem: each institution must verify its own physical and legal feasibility before the threshold is read as a financing requirement.

Remark (error-bound modulus). Along the resource route — varying the reserve x with the floors held fixed — the distance to the route’s feasible slice $\{x \geq 1\}$ equals the violation $f(z) = (1 - x)_+$ itself. The error-bound modulus (Fabian et al., 2010) of this one-parameter relation — the constant τ in $\text{dist} \leq \tau f$ — is therefore identically 1: the route is perfectly conditioned. The rescue threshold $\kappa^*(z) = 1 - x$ coincides with this route-distance, a state-dependent function; the modulus is the constant 1 that scales it. On data with nonlinear resource dynamics the distance can cease to equal the violation, and the modulus then governs only the scaling of distance-to-feasibility under unstructured perturbation, a regime outside the present scope (Section 6.2).

Data requirements for empirical application. The theorem is a result about assessment operators on a specified datum. Empirical application requires, in addition to the typed floors, disturbance set, action set, tube model, and destination maintainability witness of Section 4.5, five further specifications: the completeness status of the action set (whether $\mathcal{A}$ is exhaustive or merely a listed menu), calibration and identifiability data, policy and authority data, the threshold-uncertainty semantics, and the destination maintainability status. The itemized requirements are given in the Supplementary Material (S9).

5.8. The substitutability spectrum on the witness datum (a labelled extension)

Section 5.1 closes by conceding the nonlinear territory: CES-type substitutability aggregates are different assessment operators, and the theorems of Section 4 claim nothing for them. This subsection — a labelled extension, not a part of the separation theorem — begins the programme of establishing what *can* be claimed, on the witness datum only, with exact rational witnesses throughout. Its results are stated on the datum of Section 4.5 and inherit all of its scope delimitations (Section 5.4); nothing here is an empirical claim.

Notation. Throughout this subsection σ denotes the elasticity of substitution of a CES aggregate, $\sigma = 1/(1 - \theta)$ with θ the CES exponent. This σ is distinct from the Schaefer surplus $\sigma(B)$ of Section 6.3, from the weight ratios ρ_1, ρ_2 (which are weight quotients, not elasticities), and from the σ -notations of the companion papers — the timing certificate $\sigma^*(B_0)$ of the obstruction-calculus companion (Abaee, 2026, *An Obstruction Calculus for Viability under Incomplete Observation*) and the flow shares σ_f, σ_c of the two-land biocapacity companion (Abaee, 2026, *How Aggregation Can Conceal Composition*).

The family. Work in floor-referenced indices $\lambda_i = 1 + s_i$ — $\lambda_i = 1$ exactly at the floor, $\lambda_i \leq 0$ a full floor-unit below it — the currency in which geometric-mean abundance indices are defined. For weights normalized to sum to one, consider the CES/power means

$$M_\theta(\lambda; w) = \left(\sum_i w_i \lambda_i^\theta \right)^{1/\theta}, \quad \sigma = \frac{1}{1-\theta},$$

and evaluate Protocol 2 of Table 1 under the θ -aggregate: a state is accepted when, for every weight, some plan of the Section 4.5 menu keeps $M_\theta \geq 1$ along its worst-case tube (the physical and destination constraints are aggregator-independent; STAGED needs $x \geq 1$ and then serves every weight at every rung). Table 7 situates the members.

Table 7: The CES family on floor-referenced indices, on the witness datum. The linear member is the paper’s own engine (Lemma A); the Leontief member is the typed operator.

θ	σ	reading
1	∞	the linear aggregate — the paper’s own compensatory engine (Lemma A)
2/3	3	intermediate CES
1/2	2	intermediate CES
0	1	the geometric mean — the Living Planet Index functional structure (the identification, below)
-1	1/2	the harmonic mean
$-m$	$1/(m+1)$	the power means below the harmonic member
$-\infty$	0	the Leontief minimum — exactly the typed operator

Lemma A (the handshake). *With weights normalized, $\sum_i w_i \lambda_i \geq 1 \iff w \cdot s \geq 0$; the $\theta = 1$ member therefore reproduces the compensatory engine of Sections 3–4 exactly — the same acceptance set at every state and weight, and at the Section 6.3 datum the same per-weight plan assignment (the thresholds $\rho_1 = 2/3$, $\rho_2 = 3/2$). The proof is the identity $\lambda_i = 1 + s_i$ and the normalization. The extension is therefore genuine: it does not move the paper’s goalposts; it interpolates between the paper’s operator ($\sigma = \infty$) and the typed operator ($\sigma = 0$).*

Lemma B (collapse convention). *For every $\theta < 1$ member, any tube value $\lambda_i \leq 0$ rejects the plan: a coordinate at or below the collapse level cannot be compensated by surpluses elsewhere. The geometric mean is 0 there; CES with $\theta \leq 0$ is undefined there; for $\theta \in (0, 1)$ the convention is the family’s defining property. The $\theta = 1$ member alone compensates across collapse — in itself an exact statement of the difference between the doctrines: only the perfectly-substitutable dashboard can certify through a collapsed coordinate.*

The master equation and the exact ladder. On the gap-region diagonal $z = (x < 1, s, s)$, $s \in (1, 2)$ — the Section 6.3 datum sits on it at $s = 6/5$ — the cover condition collapses to a single comparison:

- $\theta > 0$: accept $\iff (s-1)^\theta + (s+1)^\theta \geq 2$;
- $\theta = 0$: accept $\iff s^2 \geq 2$;
- $\theta < 0$: accept $\iff (s-1)^\theta + (s+1)^\theta \leq 2$ (the inequality flips with the sign of the denominator).

Each rung of the ladder then has an exact rational closed form (Table 8); every decision is a pure rational comparison, cross-validated by a second independent code path in the deposited verifier.

The critical floors form a strictly increasing ladder $1 < 5/4 < \sqrt{2} < \varphi < \sqrt{3} < \dots \rightarrow 2$: the deeper a state sits in the discrepancy region, the lower the substitutability needed to expose the breach, and the ladder’s collapse point is exactly the typed boundary. The appearance of $\sqrt{2}$, φ ,

Table 8: The exact closed forms and critical floors of the ladder (diagonal gap states $z = (x < 1, s, s)$). The critical floors are strictly increasing in depth and converge to the typed boundary $s = 2$.

θ	accept $\iff$	critical floor $s^*(\theta)$
1	$s \geq 1$	1 (rational)
2/3	$s^2 \geq 3$, else $27(s^2 - 1)^2 \geq (3 - s^2)^3$	root of $27(s^2 - 1)^2 = (3 - s^2)^3$, in $(117/100, 59/50]$
1/2	$s \geq 5/4$	5/4 (rational)
0	$s^2 \geq 2$	$\sqrt{2}$, in $(141/100, 142/100]$
-1	$s^2 \geq s + 1$	the golden ratio $\varphi = (1 + \sqrt{5})/2$ (Fibonacci witnesses: 8/5 rejects, 13/8 accepts)
-2	$s^2 \geq 3$	$\sqrt{3}$, in $(173/100, 174/100]$
-m	$(s-1)^m + (s+1)^m \leq 2(s^2-1)^m$	increasing in m , brackets in the verifier
$-\infty$ (Leontief)	$s \geq 2$	2 — the typed boundary

and $\sqrt{3}$ as exact critical floors of the geometric, harmonic, and $\sigma = 1/3$ aggregators on this datum is, to our knowledge, a new exact witness for the substitutability debate.

The critical-elasticity landscape. For each gap state — typed-reject, linear-accept — the accepted rungs form a prefix of the ladder (Theorem S1), so there is a critical elasticity $\sigma^*(z)$: the σ -aggregator false-certifies z exactly when $\sigma \geq \sigma^*(z)$. Table 9 lists exact brackets on the diagonal $x = 1/2$.

Table 9: The critical elasticity $\sigma^*(z)$ on diagonal gap states $z = (1/2, s_1, s_2)$ (exact brackets; strict at both ends unless noted).

state $z = (x, s_1, s_2)$	$\sigma^*(z)$	note
(1/2, 1, 1)	∞ exactly	only the linear member certifies
(1/2, 3/2, 1/2)	∞ exactly	min coordinate ≤ 1 : only-linear
(1/2, 6/5, 6/5)	$\in (2, 3)$, strict	the Section 6.3 canonical datum
(1/2, 5/4, 5/4)	$= 2$ exactly	the $\theta = 1/2$ rung's equality state
(1/2, 13/10, 13/10)	$\in (1, 2)$	
(1/2, 3/2, 3/2)	$\in (1/2, 1)$, strict	the LPI-form witness
(1/2, 13/8, 13/8)	$\in (1/3, 1/2)$	the harmonic member's false-certification witness
(1/2, 9/5, 9/5)	$\in (1/5, 1/4)$	
(1/2, 39/20, 39/20)	$\in (1/17, 1/13)$	near the typed boundary: $\sigma^* \rightarrow 0$

The structure on the witness: gap states with $\min(s_1, s_2) \leq 1$ are only-linear ($\sigma^* = \infty$); states with both coordinates > 1 have finite σ^* , given on the diagonal by the master-equation root, with $\sigma^*(s) \rightarrow 0$ as $s \rightarrow 2^-$ and $\sigma^*(s) \rightarrow \infty$ as $s \rightarrow 1^+$.

Theorem S1 (nesting). *For every pair of rungs $\theta > \theta'$ and every state z , acceptance at θ' implies acceptance at θ . Consequently the accepted rungs at each state form a prefix of the ladder; $\sigma^*(z)$ is well defined; and it is strictly positive and finite on every gap state with both coordinates > 1 .*

Proof. For fixed positive λ and weights, M_θ is nondecreasing in θ (the power-mean inequality); the collapse convention preserves this (acceptance at any $\theta' < 1$ forces $\lambda > 0$ on the serving plan's

tube, where the family is defined). Per plan and weight, hence per weight, hence for the protocol. Finiteness: on the diagonal, $(s - 1)^\theta \rightarrow \infty$ as $\theta \rightarrow -\infty$ since $s - 1 < 1$, so some finite θ rejects. Positivity: $s_1 + s_2 \geq 2$ accepts at $\theta = 1$. $\square$

Theorem S2 (the two-sided Leontief statement). (i) *Pointwise: every gap state is safe under some strictly positive elasticity — full Leontief is never necessary at a fixed state.* (ii) *Uniform: for every $\theta < 1$ the rung false-certifies an exact rational interval of gap states (the critical floor $s^*(\theta) < 2$ strictly; concretely $\sigma = 3$ certifies $s = 6/5$, $\sigma = 2$ certifies $s = 13/10$, $\sigma = 1$ certifies $s = 3/2$, $\sigma = 1/2$ certifies $s = 13/8$, and $\sigma = 1/4$ certifies $s = 9/5$). Hence no positive σ is uniformly safe on the gap region; the uniform critical elasticity is $\inf_z \sigma^*(z) = 0$; and the only uniformly safe aggregator is $\sigma = 0$ — the Leontief member, which is exactly the typed operator ($\mathcal{V}^0 = \mathcal{V}_{\text{typ}}$).*

Proof. (i) is Theorem S1's prefix together with acceptance at $\theta = 1$ (every gap state is linear-accepted). (ii) Evaluate the master comparison at $s = 2$: for $\theta \in (0, 1)$ the accepting side is ≥ 2 and $(s - 1)^\theta + (s + 1)^\theta = 1 + 3^\theta > 2$ strictly; for $\theta < 0$ the accepting side is ≤ 2 and $1 + 3^\theta < 2$ strictly — both strictly accepting at $s = 2$, so by continuity some $\varepsilon > 0$ has every $s \in (2 - \varepsilon, 2)$ accepted at that rung, and these are gap states (typed acceptance needs $s \geq 2$). The Leontief member's acceptance is weight-independent ($\min_i \lambda_i \geq 1$ does not involve w), so its protocol collapses into the common-plan criterion — the typed operator. $\square$

The LPI identification, with its caveats. The $\sigma = 1$ member is the Living Planet Index's functional form — the geometric mean of abundance indices — evaluated here on floor-referenced indices ($\lambda_i = 1 + s_i$, so $\lambda = 1$ at the floor) under the paper's Protocol 2 quantifier structure. Two caveats state the exact scope: (a) the LPI's inputs are abundance indices relative to a base year, not floor margins — the identification is at the level of the functional form on the ratio scale, with the floor as reference; (b) the LPI's operational weighting (equal weights on system-level indices, chained) is a point in the weight cone, whereas Protocol 2 quantifies over the whole cone — the rung decision is therefore the adversarial version of the LPI reading, and over-, not under-states the index's exposure. No claim is made that the LPI as published implements Protocol 2.

The relevance chain, at the Section 6.3 datum. The extension passes the relevance test — a named decision, a changed indicator report, a new exact datum, and a changed management action. (1) The decision: the benchmark's closure decision at the canonical datum, under a DFO precautionary-approach limit rule. (2) The report flip: at the canonical trough $\lambda = (1/5, 11/5)$, the linear dashboard (equal weight) reports $6/5 \geq 1$ — certified — while the LPI-structured geometric index reports $\sqrt{11/25} < 1$ (exactly: $11 < 25$) — a decline signal; the same datum, opposite reports. (3) The new exact datum: $\sigma^* \in (2, 3)$ strict at the canonical state, the algebraic ladder, and the per-rung false-certification witnesses of Theorem S2(ii). (4) The action flip: at the canonical datum, every $\sigma \geq 3$ aggregator certifies — no mid-transition LRP response — while every $\sigma \leq 2$ member rejects, triggering the mandatory closure and rebuilding response; and at the witness $(x, s_1, s_2) = (1/2, 3/2, 3/2)$ the flip is sharper — the linear and the LPI-form both certify (no closure) while the harmonic form rejects, so an agency that replaced its linear dashboard with the LPI-structured form would still not trigger the response at this state; the flip requires σ below the state's critical elasticity $\sigma^* \in (1/2, 1)$, and the triggered response is concretely the reserve top-up $\kappa^* = 1 - x = 1/2$ financing the staged plan (Section 5.7). The aggregator choice — a reporting decision, not a management decision — flips a management action.

Cross-references within the programme. Two companion results bracket this extension. The two-land biocapacity study (Abae, 2026, *How Aggregation Can Conceal Composition*) es-

establishes the *composition illusion* — its continuous-time instance of the compensatory-aggregation gap, in which a weighted aggregate meets its floor while the typed floor on ecological capital is violated; Theorems S1 and S2 are the aggregator-side theorems behind that instance. The obstruction-calculus companion (Abaee, 2026, *An Obstruction Calculus for Viability under Incomplete Observation*) states the feasibility-side complement — its substitution-pathway certificates are separation certificates, not universal exchange rates — of which the present spectrum is the elasticity-side statement.

Verification and limits. Every number in this subsection is an exact rational, verified in a dedicated fail-loud exact-arithmetic verifier (57 checks; standard-library fractions only; deterministic; with a committed byte-reproducible run log) — a third check list, separate from and not pooled with the grid verifier’s 25 checks and the software companion’s 24 (the programme’s non-pooling policy). Two limits are stated. Off the diagonal, the geometric member’s cover comparison is log-transcendental, and the verifier decides it only on proven sufficient and necessary rational conditions — the headline results are diagonal, so nothing load-bearing rests on the residual undecided cells. And interior- θ values of $\sigma^*(z)$ are bracketed between tested rungs — the brackets are exact and strictness-refined; the interior value is the master root, transcendental in general.

The off-diagonal closure. Every off-diagonal cell that the rational conditions leave undecided — and with it the complete two-dimensional σ^* -landscape — is decided by a dedicated exact computation in the same programme (same datum, same exactness standard, same delimitations). The log-transcendental comparison admits no closed form (it has none); it is *decided*: every verdict below is a finite rational proof. The device is the *certified rational enclosure*: every logarithm is bracketed between two rationals carrying a checkable certificate (the all-positive-term artanh series for $\ln$, its tail dominated geometrically), every fractional power by integer k -th-root brackets, and every cover comparison below is multilinear in the enclosed quantities, so that its extrema over the enclosure box are attained at the vertices — a finite rational computation. Under the fail-loud discipline (any cell the schedule cannot decide aborts the run), zero cells remain undecided.

Theorem G (the exact off-diagonal criterion). *On the viable gap quadrant — $x < 1$, both coordinates in $(1, 2)$ — the geometric member’s cover comparison reduces, by cross-multiplication of the two plans’ weight thresholds, to*

$$\text{cover} \iff A \cdot D \leq B \cdot C,$$

with $A = \ln(s_1 - 1)$, $B = \ln(s_2 + 1)$, $C = \ln(s_1 + 1)$, $D = \ln(s_2 - 1)$; equivalently

$$r(s_1)r(s_2) \leq 1, \quad r(s) = \frac{\ln(1/(s-1))}{\ln(s+1)},$$

so the acceptance region is $\{s_2 \geq \rho(s_1)\}$, where $\rho = r^{-1} \circ (1/r)$ is a strictly decreasing involution with the unique fixed point $\rho(\sqrt{2}) = \sqrt{2}$ (exactly: $r(s) \leq 1 \iff s^2 \geq 2$) and limits $\rho(s) \rightarrow 2^-$ as $s \rightarrow 1^+$, $\rho(s) \rightarrow 1^+$ as $s \rightarrow 2^-$. The pivot has two sides: $\min(s_1, s_2) \geq \sqrt{2}$ is sufficient for the cover, and the cover forces $\max(s_1, s_2) \geq \sqrt{2}$; the residual wedge $\{\min < \sqrt{2} < \max\}$ is split by the curve ρ , which is transcendental — bracketed at rational columns with certified endpoints in the verifier. A coordinate $s_i \geq 2$ accepts outright — its own plan’s worst tube keeps both coordinates at or above 1 ($s_i - 1 \geq 1$ and $s_j + 1 > 1$), serving every weight at every rung — so the residual decision problem lives on the open quadrant. For the interior rungs $\theta \in (0, 1)$ the same cross-multiplication, applied to the four powers $A = (s_1 - 1)^\theta$, $B = (s_2 + 1)^\theta$, $C = (s_1 + 1)^\theta$, $D = (s_2 - 1)^\theta$, gives the

multilinear form $F = C(B - 1) + A(1 - D) - (B - D) \geq 0$, which reduces on the diagonal to the master equation of Table 8.

Proof sketch (machine-anchored). FAST's per-weight acceptance set is the downward interval $[0, U]$, SLOW's the upward interval $[L, 1]$, with $U = B/(B - A)$ and $L = -D/(C - D)$ at the geometric member (both denominators positive on the quadrant); the cover is $L \leq U$, and the cross-multiplication is the criterion. The function r is strictly decreasing (its numerator decreasing, its denominator increasing), so ρ is well defined and strictly decreasing, and the criterion's symmetry makes it its own inverse. The enclosure lemmas certify the four quantities — and the four powers — between rationals; the comparison forms are multilinear, so the vertex ranges decide; every rejection additionally carries an exhibited rational witness weight served by neither plan, cross-validated in both directions against the independent per-weight algebra. $\square$

The residual cells decided. The two wedge witnesses certify that the residual was genuine: $(5/4, 3/2)$ and its mirror $(3/2, 5/4)$ **reject**, with certified margins ≈ 0.218 and exhibited rational witness weights served by neither plan, while $(5/4, 7/4)$ and its mirror **accept**, with margins ≈ 0.422 — no per-coordinate rational condition separates them; the curve ρ does. Together with the six $(5/4, \geq 2)$ -family cells, accepted by the $s_i \geq 2$ boundary argument, all ten undecided cells of the 64-state grid are decided with certified strict verdicts.

The completed two-dimensional landscape. With the geometric member decidable off the diagonal and the interior rungs $\theta = 1/2$ and $2/3$ given the same treatment, all 13 rungs are decided at all 64 grid states, with zero cells undecided: the accepted rungs form a prefix at every state (Theorem S1's nesting, now machine-complete off the diagonal), the acceptance regions are up-closed in each coordinate — $\sigma^*(z)$ is nonincreasing in s_1 and in s_2 (increasing either coordinate raises both plans' worst-tube aggregates) — and the diagonal closed forms of Table 8 are reproduced, the $\theta = 1/2$ rung's $(5/4, 5/4)$ equality state included. The nine finite- σ^* states of the core $\{5/4, 3/2, 7/4\}^2$ form clean off-diagonal bands mirroring the diagonal ladder: $(5/4, 3/2)$ and its mirror in $(1, 2)$; $(5/4, 7/4)$ and its mirror in $(1/2, 1)$; $(3/2, 7/4)$ and its mirror in $(1/3, 1/2)$; $(7/4, 7/4)$ in $(1/4, 1/3)$; the diagonal $(5/4, 5/4)$ is $= 2$ exactly and $(3/2, 3/2) \in (1/2, 1)$, both already in Table 9. Off-grid specials: $(6/5, 7/5)$ and $(13/10, 3/2)$ lie in $(1, 2)$. The remaining thirteen gap states of the grid are only-linear ($\sigma^* = \infty$, the $\min(s) \leq 1$ boundary states).

Remark (the fixed-sum tolerance structure). Slice the gap quadrant by total margin $s_1 + s_2$ (the anti-diagonals). The elementary anchors: a slice's balanced point accepts exactly when $s_1 = s_2 \geq \sqrt{2}$ (the diagonal closed form), and the $s = 1$ edge accepts exactly from total 3 up (the boundary argument); below total 3 the near-edge rejects ($r \rightarrow \infty$ at the edge, against a fixed positive partner). On the scanned grids the slices assemble into four regimes:

- **R1** (totals $\leq 141/50$, the canonical $12/5$ among them): every scanned point rejects — no asymmetry saves the total;
- **R2** (an interior window, $141/50 < \text{total} < 2\sqrt{2}$; exemplar $707/250$): the balanced point rejects, moderate asymmetry accepts, the edges reject;
- **R3** (an anchored window, $2\sqrt{2} \leq \text{total} < 3$; exemplar $71/25$): the balanced point accepts, strong asymmetry rejects;
- **R4** (totals ≥ 3): the whole scanned slice accepts.

The regime boundaries $2\sqrt{2}$ and 3 are elementary; the slice-flip total — where a slice's first accepting point appears — is transcendental and machine-bracketed: every scanned point of the $141/50$ slice rejects, while the $353/125$ slice already contains acceptors, so the flip lies in

(141/50, 353/125], strictly below $2\sqrt{2}$. The ecological content of the interior window: *moderate concentration of a deficit is more certifiable than balance* — below $2\sqrt{2}$, the balanced point is not the slice's best point.

The anchor decision: distribution versus total. The benchmark's asymmetric-allocation licensing choice fixes the total margin at $s_1 + s_2 = 707/250$ (total biomass 1707/250 kt at the Section 6.3 floor $B_{\text{lim}} = 2$ kt per stock) and asks how the verdict depends on the allocation of the deficit across the stocks. Both allocations are gap states, and the LPI-structured composite rejects the balanced allocation (707/500, 707/500) by the diagonal closed form — $(707/500)^2 = 499849/250000 < 2$ exactly, by the exact excess 151/250000 — while accepting the concentrated allocation (607/500, 807/500) by the certified log-comparison (r -product ≈ 0.9847 , margin $\approx 1.17 \times 10^{-2}$). The same total margin, opposite verdicts: the false certification of the compensatory dashboard is visibly two-dimensional — blind to the *distribution* of the deficit, not only to its total; the linear dashboard licenses both allocations, reporting literally the same number, $707/500 \geq 1$, on every trough of both states. The report layer carries the same structure: at the concentrated state, which stock bears the pulse flips the equal-weight geometric report — the FAST-trough $\sqrt{139849/250000} < 1$, a decline signal, against the SLOW-trough $\sqrt{339849/250000} > 1$, certified — a plan-design decision with certification content. And the action layer: the balanced allocation triggers the mandatory closure and rebuilding response (the reserve top-up $\kappa^* = 1 - x = 1/2$ financing the staged plan, Section 5.7); the concentrated allocation is licensed with no mid-transition response. The allocation — a design-level choice — flips a management action.

Verification (a fifth check list). Every number in this closure is an exact rational, verified in a dedicated fail-loud exact-arithmetic verifier (50 checks; standard-library fractions only; no floats, tolerances, or randomness; deterministic; with a committed byte-reproducible run log) — a fifth check list, separate from and not pooled with the grid verifier's 25, the software companion's 24, the σ -extension's 57, and the common-shock extension's 77 (the programme's non-pooling policy). Limits: the curve ρ , the slice-flip total, and every interior σ^* value are transcendental — the closure provides certified rational brackets, not closed forms; the slice scans certify the scanned grid points, the continuum slice-minimum remaining open (the interval character of the accepting windows is machine-evidenced, not proved); and the LPI-identification caveats above carry over unchanged.

The continuum slice closure. The continuum slice-minimum — the open residual recorded above — is decided: a dedicated exact computation in the same programme (same datum, same exactness standard, same delimitations) settles the continuum. The device is the ϕ -reduction: a point $(s, t - s)$ of the total- t slice accepts at the geometric member iff $\phi(s) := s + \rho(s) \leq t$, so the fixed-sum programme is the sublevel-set programme of one explicit scalar function on $(1, 2)$, with $\phi(1^+) = \phi(2^-) = 3$ and $\phi(\sqrt{2}) = 2\sqrt{2}$. The tangency identity $\rho'(s) = -h(s)/h(\rho(s))$, with $h = r'/r$, makes the stationary points of ϕ the zeros of $D(s) = h(s) - h(\rho(s))$; there are exactly three — the elementary fixed point $\sqrt{2}$, and a tangent pair $a^* < \sqrt{2} < b^*$ with $\rho(a^*) = b^*$ and $\rho'(a^*) = -1$ — giving the certified four-phase shape of ϕ : descending from the limit 3 to its minimum at a^* , rising through $2\sqrt{2}$, descending to the equal minimum at b^* , climbing back to 3. The consequences, each a certified rational proof:

- **The slice-flip total.** $t^* = \min_{s \in (1,2)} \phi = a^* + b^* \in (2.82256976476, 2.82256976491)$, certified to ten places by two independent code paths (the tangency bisection on D ; the direct slice bisection) — five orders inside the scan bracket (141/50, 353/125] and strictly below $2\sqrt{2}$. The flip is attained at the tangent pair, where the curve ρ touches the anti-diagonal: $a^* \in$

(1.14796785390, 1.14796785398), $b^* \in (1.67460191085, 1.67460191093)$.

- **The interval character.** Every accepting window is a sublevel interval of one monotone phase of ϕ — no internal holes: the fixed-sum remark’s scan-level finding is thereby proved.
- **The total-3 boundary proved in the interior.** $\phi < 3$ on $(1, 2)$ — every interior point of every total-3 slice accepts — and with it the R4 regime wholesale: $\phi(s) < 3 \leq t$ for every point of every total- $t \geq 3$ slice.
- **The harmonic circle.** The $\theta = -1$ member’s off-diagonal acceptance region is exactly the exterior of $(2s_1 - 1)^2 + (2s_2 - 1)^2 = 10$ — a polynomial identity; the circle passes through $(1, 2)$, $(\sqrt{2}, \sqrt{2})$, and $(2, 1)$. So nothing at or below $\theta = -1$ covers anything below total 3, and the m-side flip total is exactly 3, attained at the corners — which every rung accepts.
- **The σ^* -depth sandwich.** Writing $\sigma_{\min}(t)$ for the slice-minimum of the critical elasticity, $\sigma_{\min}(t) \in (1/2, 1]$ on $(t^*, 3)$ (the circle’s strict interior below, the ϕ -minimum above); at the anchor total 707/250, $\sigma_{\min} \in [12/13, 1]$ — the interior window’s best point needs nearly full substitutability and lies near the tangent branch, sliding toward the corners as t climbs to 3. Beyond total 3 the corner cusps ($\delta \sim \varepsilon^m/m$) give $\inf \sigma^* = 0$: certified rational witnesses at every tested depth m .
- **The flip-total ladder.** The linear member exactly 2 (the dashboard identity: cover iff $s_1 + s_2 \geq 2$); $\theta = 2/3$ exactly 2×the cubic master root; $\theta = 1/2$ exactly 5/2 (the rung’s own equality state $(5/4, 5/4)$); $\theta = 0$ the transcendental t^* — moderate concentration beats balance from the geometric member down; every $\theta = -m$ and the Leontief member exactly 3 (the circle; the corners). Strictly increasing in depth, nesting-consistent.

Every number above is an exact rational, verified in a dedicated fail-loud exact-arithmetic verifier (44 checks; standard-library fractions only; deterministic; with a committed byte-reproducible run log) — a sixth check list, separate from and not pooled with the five above (the programme’s non-pooling policy). Limits: the tangent quantities (a^* , b^* , t^* , every interior σ^*) are transcendental and carry certified rational brackets only; the four-phase certificate’s end segments rest on value anchors and the elementary asymptotics of h at the ends of $(1, 2)$, where h diverges.

5.9. The common-shock variant of the witness (a labelled extension)

Section 5.4’s disturbance scoping (Table 5) leaves the middle cell — the common biomass shock, one shared disturbance striking the same floor coordinate of every plan — as the one undetermined verdict: its assertion demands its own exact computation, a common-shock variant of the witness. This subsection — a labelled extension in the sense of Section 5.8, not a part of the separation theorem — records that computation and sharpens the coupled regime’s row-(iii) assertion into two exact, reading-conditional theorems. Its results are stated on the datum of Section 4.5 with only the disturbance class replaced as specified below; all of Section 5.4’s scope delimitations are inherited; nothing here is an empirical claim.

The variant. Decompose the Section 4.5 datum per the Section 6.3 arithmetic: each principal plan’s announced schedule is a mid-interval tent drawdown of depth $3/2$ in its characteristic coordinate (FAST: the quota pulse on s_1 ; SLOW: the gradual burden on s_2), the other coordinate constant; each environmental event is a transient tent dip of depth δ in one floor coordinate, troughing at $t = 1/2$ and recovered by $t = 1$, successors unchanged; the Section 4.5 worst-case depth $2 = 3/2 + 1/2$ is the quota pulse plus the heatwave, so the datum’s event depth is $\delta = 1/2$. STAGED spends x (it needs $x \geq 1$) while its floors grow linearly to $s + e$, $e = (1/4, 1/4)$; an event of depth δ on a growing floor dips it to $s + 1/8 - \delta$ — by $\max(0, \delta - 1/8)$. The *common*

disturbance class is: the two coordinate events are separate disturbances, each striking every plan, so each plan's per-coordinate worst case takes its own schedule plus the event on that coordinate — ecologically, a marine heatwave striking the same biomass coordinate regardless of which harvest-rebuild plan is running. The disturbance quantifier remains innermost throughout (Section 2.3). One methodological device is used throughout: STAGED's aggregate dip $\max(0, \text{raw}(p) - 1/8)$ is kinked in the weight share $p = w_1/(w_1 + w_2)$ (with raw affine in p), and the admissibility requirement decomposes exactly into the two affine conditions of the kink's two branches — every other plan's dip is affine in p — so every cover decision below reduces to exact endpoint evaluations.

The handshake. Before any new class is evaluated, the computation anchors itself to the deposited datum: under the action-indexed class of Section 4.5 the variant's machinery reproduces Theorem 5 exactly — the same typed and weak acceptance sets, the same licensing thresholds $\rho_1 = 2/3$, $\rho_2 = 3/2$ at the canonical datum $(\frac{1}{2}, \frac{6}{5}, \frac{6}{5})$, the same strictness witness $(\frac{1}{2}, \frac{1}{10}, \frac{1}{10})$ — verified on 676 grid states plus the witness list, the same anchoring discipline as the σ -extension's verifier.

Theorem M1 (the middle-cell verdict). *On the common-shock variant at the datum's event depth $\delta = 1/2$:*

$$\mathcal{V}_{\text{typ}}^{\text{cs}} = \{s_1 \geq 2, s_2 \geq \tfrac{1}{2}\} \cup \{s_2 \geq 2, s_1 \geq \tfrac{1}{2}\} \cup \{x \geq 1, s_1 \geq \tfrac{3}{8}, s_2 \geq \tfrac{3}{8}\},$$

$$\mathcal{V}_{\text{weak}}^{\text{cs}} = \{x < 1, s_1 \geq \tfrac{1}{2}, s_2 \geq \tfrac{1}{2}, s_1 + s_2 \geq 2\} \cup \{x \geq 1, s_1 \geq \tfrac{3}{8}, s_2 \geq \tfrac{3}{8}\},$$

and the acceptance gap survives the common-shock regime:

$$\text{FP}_{\text{cs}} = \mathcal{V}_{\text{weak}}^{\text{cs}} \setminus \mathcal{V}_{\text{typ}}^{\text{cs}} = \{x < 1, \tfrac{1}{2} \leq s_1 < 2, \tfrac{1}{2} \leq s_2 < 2, s_1 + s_2 \geq 2\}.$$

The gap is strictly smaller than the action-indexed impossibility region I by the exact 1/2-cut — every action-indexed gap state with $\min(s_1, s_2) < 1/2$ is weak-rejected under the common class (witness $(\frac{1}{2}, \frac{1}{10}, \frac{19}{10})$, action-indexed weak-accepted by $s_1 + s_2 = 2$) — it is nonempty with interior, and the canonical datum $(\frac{1}{2}, \frac{6}{5}, \frac{6}{5})$ is a member: the compensatory/noncompensatory separation is not an artifact of the action-indexed convention. On the surviving region the licensing structure is unshifted: FAST's common-class licensing endpoint equals its action-indexed endpoint $1/(1 + \rho_1)$ at every gap state of the surviving region (SLOW's the mirror) — the middle cell's verdict is reduced by the 1/2-cut, not shifted.

Proof sketch (machine-anchored). Write $A(p) = p s_1 + (1 - p) s_2$ for the aggregate margin at weight share p . FAST's common-class requirement is $A(p) \geq \max(2p, p + \frac{1}{2})$ (its own tent $3/2$ in s_1 plus the s_1 -event, and the s_2 -event weighted by $1 - p$); SLOW's is the mirror; STAGED's ($x \geq 1$) is $A(p) \geq \max(p, 1 - p)/2 - 1/8$ (the kink device: the event exposure of a growing floor). The pointwise minimum of the two principal requirements is a tent with its only kink at $p = 1/2$, so an affine A covers it exactly at the three abscissae $p \in \{0, \frac{1}{2}, 1\}$; the closed forms follow. Every constant is an exact rational; the verifier checks the closed forms against independent grid computations (5,000-state grids for the common class) plus a curated witness list, both directions, at every grid point. $\square$

Theorem M2 (the rescue trichotomy). *Under the common class the rescue threshold of Proposition 11 becomes*

$$\kappa_{\text{cs}}^*(z) = \begin{cases} 0, & z \in \mathcal{V}_{\text{typ}}^{\text{cs}}, \\ (1 - x)_+, & z \notin \mathcal{V}_{\text{typ}}^{\text{cs}}, \min(s_1, s_2) \geq \tfrac{3}{8}, \\ +\infty, & z \notin \mathcal{V}_{\text{typ}}^{\text{cs}}, \min(s_1, s_2) < \tfrac{3}{8}. \end{cases}$$

The qualitative change: the resource route fails to rescue every failure state — the common shock introduces the first unrescuable-by-reserve failure states on this datum, with witness $(\frac{1}{2}, \frac{1}{4}, 3)$: no reserve increment whatsoever rescues it, the shared shock's exposure of the staged route's growing floor $(3/8)$ exceeding what a floor below $3/8$ can carry. The canonical datum's shortfall stays $\kappa^ = 1/2$, its action-indexed value. The $3/8$ boundary is proved structurally, not sampled: the event exposure of a growing floor is $\max(0, \delta - 1/8) = 3/8$ at the datum's depth, and no floor below it can carry the event.*

Proof sketch. The augmented menu adds STAGED_κ , admissible at $x + \kappa \geq 1$ with the same floor requirement $s_i \geq 3/8$; κ cannot move the floors, so on $\{\min(s) < 3/8\} \setminus \mathcal{V}_{\text{typ}}^{\text{cs}}$ no finite κ certifies. On the complement the requirement is the action-indexed one, $\kappa = (1 - x)_+$. $\square$

Theorem M3 (blend fragility). *Theorem 9's blend window is class-conditional: under the action-indexed class at event depth δ it is nonempty exactly on $\{s_1 + s_2 \geq 3/2 + \delta\}$ (Theorem 9 is the $\delta = 1/2$ instance, $\{s_1 + s_2 \geq 2\}$); under the common class it is nonempty exactly on*

$$\{s_1 \geq \delta, s_2 \geq \delta, s_1 + s_2 \geq \frac{3}{2} + 2\delta\}.$$

The common-class gap is therefore not blend-closed on the fragile band $s_1 + s_2 \in [2, \frac{5}{2})$: the witness $(\frac{1}{2}, 1, 1)$ is a common-shock gap state with no admissible blend, while every common-shock gap state with $s_1 + s_2 \geq 5/2$ has one. The coupled gap, by contrast, is blend-closed — its window threshold and weak boundary coincide — so the fragility is class-conditional.

Proof sketch. The blend's s_1 -trough is $s_1 - \delta - \delta_{\text{blend}} \cdot \frac{3}{2}$ and its s_2 -trough $s_2 - \delta - (1 - \delta_{\text{blend}}) \cdot \frac{3}{2}$ (each coordinate takes its event plus the corresponding fraction of the principal tent); the window is the nonnegativity interval of the two, intersected with $[0, 1]$. $\square$

Theorem M4 (the coupled regime, additive reading). *Under the coupled class — one disturbance carrying both coordinate events simultaneously, riding on the plans' own schedules — at event depth δ : the typed accepted set coincides with the common class's at every depth (the per-coordinate tube minima are event-combination independent, each coordinate taking its own schedule plus its own event), and*

$$\mathcal{V}_{\text{weak}}^{\text{cpl}} = \{x < 1, s_1 \geq \delta, s_2 \geq \delta, s_1 + s_2 \geq \frac{3}{2} + 2\delta\} \cup \{x \geq 1, \min(s_1, s_2) \geq \max(0, \delta - \frac{1}{8})\}.$$

At the datum's full magnitude ($\delta = 2$) every enumerated I-state is weak-rejected — the Section 5.4 claim holds on the witness's gap region — but the claim is depth-dependent with the exact boundary $\delta = 5/4$: at heatwave magnitude ($\delta = 1/2$) deep-I states survive as coupled gap states (witness $(\frac{1}{2}, \frac{19}{10}, \frac{19}{10})$), no I-state survives at $\delta \geq 5/4$ while $(\frac{1}{2}, \frac{79}{40}, \frac{79}{40})$ survives at $\delta = 5/4 - 1/40$ (every I-state has $s_1 + s_2 < 4$, the survival requirement at $\delta = 5/4$); and at full magnitude the separation relocates rather than vanishes: the coupled gap

$$\{x < 1, 2 \leq s_1 < \frac{7}{2}, 2 \leq s_2 < \frac{7}{2}, s_1 + s_2 \geq \frac{11}{2}\}$$

is nonempty — the witness $(\frac{1}{2}, 3, 3)$ is weak-accepted and typed-rejected. “Universal rejection” is exact on I, not global: the divergence re-emerges at larger floors.

Proof sketch. Under the coupled class each plan's aggregate dip is the sum of both events' weighted contributions ($3p/2 + \delta$ for FAST) against the common class's maximum of single-event contributions; the cover analysis is the tent argument of Theorem M1 with the summed requirements. The boundary: an I-state has $\min(s) < 2$, hence $s_1 + s_2 < 4$, and survives exactly when $s_1 + s_2 \geq 3/2 + 2\delta$, i.e. $\delta \leq (s_1 + s_2 - 3/2)/2 < 5/4$; the witness states realize both sides. $\square$

Theorem CT (the replace-reading collapse). *Under the coupled-total class — the replace reading, each announced-schedule plan’s worst case being the full-magnitude dip 2 in each floor, tent schedules suppressed — FAST and SLOW degenerate to the same, weight-independent requirement $A(p) \geq 2$ (they fail together at every weight: at the canonical datum both serve nowhere, at $(\frac{1}{2}, \frac{5}{2}, \frac{5}{2})$ both serve all of $[0, 1]$, at $(\frac{1}{2}, 3, 1)$ both serve exactly $[\frac{1}{2}, 1]$), and*

$$\mathcal{V}_{\text{weak}}^{\text{ct}} = \mathcal{V}_{\text{typ}}^{\text{ct}} = \{s_1 \geq 2, s_2 \geq 2\} \cup \{x \geq 1, \min(s_1, s_2) \geq \frac{15}{8}\} :$$

the typed/weak separation collapses entirely — the gap is empty on every state and witness probed — and the replace reading is strictly more demanding than the action-indexed class: the witness $(\frac{1}{2}, 3, 1)$ is action-indexed typed-accepted ($s_1 \geq 2$) yet a coupled-total full reject on both tests, as is the canonical datum itself. The mechanism: the weak test’s advantage is weight-wise plan choice; when every plan faces the same plan- and weight-independent dip, choice buys nothing, and the per-weight cover condition degenerates to the same per-coordinate condition the typed test applies. STAGED’s rescue route survives — its growth mechanics are not a “schedule” in the cost sense; the dip lands on the growing floor, threshold $15/8$, exactly its coupled full-magnitude branch.

Proof sketch. With the tents suppressed and the full dip 2 in each floor, every principal plan’s aggregate dip is the constant 2 — plan- and weight-independent — so the weak cover condition $A(p) \geq 2$ for all p is the pair of endpoint conditions $s_1 \geq 2, s_2 \geq 2$, which is the typed condition; STAGED’s branch is identical under both tests. $\square$

The two readings bracket the coupled truth; both are stated rather than choosing between them. Two further exact structures complete the scoping. The *single-event sub-case* (the literal row-(ii) reading, only the s_1 -event existing) has typed acceptance $\{s_1 \geq 2\} \cup \{s_2 \geq 3/2, s_1 \geq 1/2\} \cup \{x \geq 1, s_1 \geq 3/8\}$, weak acceptance $\{x < 1 : s_1 \geq 1/2, s_1 + s_2 \geq 2\} \cup \{x \geq 1 : s_1 \geq 3/8\}$, and gap $\{x < 1, \frac{1}{2} \leq s_1 < 2, s_2 < \frac{3}{2}, s_1 + s_2 \geq 2\}$, the canonical datum a member — the gap survives this degenerate class too. And the δ -family: all closed forms above are verified at $\delta \in \{0, \frac{1}{4}, \frac{1}{2}, 1, \frac{5}{4}, \frac{3}{2}, 2\}$ on cross-section grids; $\delta = 0$ degenerates to the benign datum on which all disturbance classes coincide — the middle-cell verdict is a genuine disturbance-structure effect, not a modeling artifact.

The relevance chain, at the Section 6.3 datum. The extension passes the same four-part test as the σ -spectrum. (1) The decision: whether a composite biomass index may certify a stock-rebuilding transition (after the Northern cod case, the Section 6.3 benchmark) when the assessment’s disturbance class is a shared environmental shock — a marine heatwave striking the same biomass coordinate regardless of which harvest-rebuild plan is running — rather than plan-specific risk; Table 5 row (ii) at the canonical datum, with the checklist’s sixth audit point demanding the class be specified. (2) The report flips, on two exact surfaces: the $1/2$ -cut (any action-indexed gap state with $\min(s) < 1/2$ flips its compensatory verdict from accept to reject — an index that certified “aggregate above 2, transition licensed” now reports rejection, witness $(\frac{1}{2}, \frac{1}{10}, \frac{19}{10})$); and the rescue margin (on $\{\min(s) < 3/8\}$ the reported κ^* changes from a finite reserve requirement to infinity — “rescuable with reserve κ^* ” becomes “not rescuable by any reserve”, Theorem M2). (3) The new exact datum: FP_{cs} and its $1/2$ -cut, the $3/8$ rescue boundary, the licensing-invariance identity, the $5/4$ coupled boundary with its $79/40$ witness, the $11/2$ relocation gap, and Theorem CT’s collapse — all exact rationals, machine-anchored. (4) The action flips: at an asymmetric gap state the action-indexed prescription — accumulate bridge reserve to $x = 1$ ($\kappa^* = (1-x)_+$) — is replaced under a shared-shock assessment by a floor-targeted action, rebuilding the weaker floor above $1/2$ first (above $3/8$ for the staged route), because no reserve increment alone

certifies ($\kappa^* = \infty$); reserve accumulation and floor rebuilding are different management actions with different costs and timescales. Second, on the fragile band $s_1 + s_2 \in [2, \frac{5}{2})$, blended quota policies are not admissible under a shared shock (Theorem M3) — an agency that would otherwise blend FAST and SLOW shares at such a state must not.

Verification and limits. Every number in this subsection is an exact rational, verified in a dedicated fail-loud exact-arithmetic verifier (77 checks; standard-library fractions only; no floats, tolerances, or randomness; deterministic; with a committed byte-reproducible run log) — a fourth check list, separate from and not pooled with the grid verifier’s 25 checks, the software companion’s 24, and the σ -extension’s 57 (the programme’s non-pooling policy). Limits: everything is on the witness datum’s decomposition (tent events, exact tubes, the specific STAGED mechanics, the Section 6.3 arithmetic); no claim about other data beyond the δ -family verified. The two regime-(iii) readings (additive/replace) bracket the coupled truth; Theorem CT is reading-conditional, and both are stated rather than choosing. Grid verification is fail-loud but finite; the closed forms carry proofs (affine interval arithmetic; the kink device), and the two were cross-checked against each other at every grid point — the standard the σ -extension set. κ^* under the common class is adjudicated on a reserve grid of step $1/40$ plus closed-form confirmation on the survivors; the infinity verdict on $\{\min(s) < 3/8\}$ is proved structurally, not sampled. No stochastic, partial-observation, or infinite-horizon claim is made; the disturbance quantifier remains innermost throughout (Section 2.3).

6. Related work, and what is actually new here

Three literatures run close to this paper. The separation of Section 4 is new only in a specific and bounded sense, and the purpose of this section is to state that sense exactly rather than to claim a wide berth.

6.1. The scalarization gap in multi-objective optimization

It is a textbook fact that weighted-sum scalarization recovers only the *supported* non-dominated outcomes — those lying on the boundary of the convex hull of the non-dominated set — and that no choice of weights recovers unsupported points on non-convex portions of the front (Koski, 1985; Stadler and Dauer, 1992; Stadler, 1995; Athan and Papalambros, 1996; Chen, Hansen, Jaumard and Tuy, 1999; Das and Dennis, 1998; Messac and Mattson, 2002; Huang, Li and Liu, 2007; Miettinen, 1999; Ehrgott, 2005). The remedy literature — adaptive weighted-sum schemes, Tchebycheff and weighted-maximization scalarizations, ϵ -constraint methods — exists precisely because of it.

Theorem 4.2(ii) and Theorem 6(ii) reproduce this geometry exactly: the scalarized acceptance set is $\text{conv}(\mathcal{D}) + \mathbb{R}_+^n$, the convex hull, and the acceptance gap is its excess over the union of the plans’ orthants. **The reader should therefore not read the convex-hull characterization as a new fact about scalarization.** It is the scalarization gap, transported. What is transported, and what is not, is the subject of Section 6.4.

6.2. Adjustable and randomized robust optimization

Robust optimization distinguishes the robust counterpart, in which all decisions are here-and-now, from the adjustable robust counterpart, in which a second stage may respond to realized

uncertainty (Ben-Tal, Goryashko, Guslitzer and Nemirovski, 2004, *Mathematical Programming* 99(2), 351–376; Ben-Tal, El Ghaoui and Nemirovski, 2009). Ben-Tal et al. note that under constraint-wise uncertainty the two coincide for linear problems, and Bertsimas and Goyal (2012) bound how far a static solution may lie from an adjustable one. The question of when the two *agree* has since been settled sharply: Wei and Zhang (2024) give a necessary and sufficient condition, in the form of a theorem-of-the-alternatives, for adjustability to be zero on a polyhedral uncertainty set. Their gap is measured in optimal objective *value*; the gap studied here is one of *feasibility* — whether any common implementable plan exists at all — and their polyhedrality hypothesis has no counterpart in the present setting. The two results describe the same quantifier inversion from opposite faces: they characterise when it closes, whereas this paper exhibits when it does not, and does so on a region of nonempty interior. The analogue of the present quantifier inversion — a single plan for all weights against a plan per weight — is therefore structurally familiar in that literature.

Separately, it is well established that *randomizing* over decisions can strictly improve worst-case performance against a deterministic choice: this is the content of the mixed-strategy and complete-class results of statistical decision theory, and it has been revisited in robust optimization, where problems are classified as randomization-receptive or randomization-proof according to whether a randomized decision strictly beats the deterministic optimum (Krause, Singla and Golovin, 2011; Vorobeychik and Li, 2014; Sinha, Fang, An and Kiekintveld, 2018; Sessa, Bogunovic, Kamgarpour and Krause, 2020; Kobayashi and Takazawa, 2023). Theorem 9 — that admitting fractional blends closes the aggregate-certified gap exactly — sits in this line.

A reader from game theory will reach for Kuhn’s theorem at this point, and the point should be met directly. In an extensive-form game with perfect recall, *mixed* strategies (a distribution over plans, drawn once) and *behavioural* strategies (an independent randomisation at each decision point) are realization-equivalent: they induce the same distribution over terminal nodes (Kuhn, 1953). Theorem 9 asserts that blending over plans closes the aggregate-certified gap while alternating between plans over time does not, which on its face contradicts that equivalence. It does not, and for a reason that is already the paper’s own scope restriction. Kuhn’s equivalence is an equality of *distributions over paths*, evaluated under an *expected* criterion. The certified constraint here is *path-wise*: it must hold on every disturbance trajectory in the uncertainty set. Under a worst-case criterion the two notions of randomisation are not interchangeable, because convexification must occur *within* a time step — a blend applied at each step yields the convex combination of required margins at that step, whereas alternation varies *across* steps and therefore yields only a vertex at any given step. Since the constraint binds at every step, across-step variation never convexifies. The distinction is reinforced from the other side: Main and Randour (2024) show that Kuhn’s equivalence fails outright once strategies are restricted to *finite memory*, which is the setting here, so the classical theorem does not apply even before the worst-case criterion is imposed.

6.3. Weak and strong sustainability

The distinction this paper formalizes is the one environmental economics draws between weak sustainability, under which natural and manufactured capital are substitutes and only the aggregate need be maintained, and strong sustainability, under which critical natural capital is non-substitutable and must be maintained separately (Pearce and Atkinson, 1995; Daly, 1995; Beckerman, 1994; Ayres, 1996; Neumayer, 2003; the critical-natural-capital literature surveyed in Dietz and Neumayer, 2007). That debate is largely conducted in verbal and policy terms. It has not, to our knowledge, been given a separation theorem with an exact witness, and the practical

maxim advanced here — that where floors bind separately, reporting them individually is a detection requirement rather than a presentation preference — is a consequence of the formal result rather than a restatement of the policy position.

6.4. What is claimed as new

Given the above, the contribution is deliberately narrow. We claim four things and nothing more.

1. **The setting.** The separation is established for *robust transition safety over a finite horizon with path-wise constraints*, not for the recovery of Pareto points by a scalarized optimization. The objects compared are acceptance *sets of states* under a viability constraint, and the argument turns on quantifier order in that setting. The multi-objective literature asks whether a weighted-sum optimizer can *reach* an unsupported point; we ask whether a *transition certified by every weighting* is implementable at all.
2. **The negative result on time-alternation, which inverts the usual ordering.** Theorem 9 shows that fractional blending closes the gap exactly, while alternating between plans over time does not. In robust optimization the ordering runs the other way: adjustability weakens conservatism, so the adjustable counterpart dominates the static one. Here convexification in *plan space* achieves what open-loop variation over *time* cannot. We have not found this comparison, or this inversion, in the robust optimization or randomization literatures surveyed above, which study worst-case *objective value* rather than path-wise *feasibility*. This is the sharpest claim in the paper and the one most likely to be contested.
3. **The detection reading.** Because the gap has nonempty interior on an explicit datum, an aggregate index cannot distinguish the rescuable from the impossible without a separate check of the floors or a verified common implementable plan. That is a statement about what a reported index can and cannot certify, and it is a consequence of (i) and (ii), not an independent assertion.
4. **The witness.** The datum is realized in exact rational arithmetic with a fishery transition under a heatwave disturbance, units anchored to a regulated biomass limit, with the gap a region of nonempty interior rather than a measure-zero boundary phenomenon.

We do *not* claim the quantifier-order observation itself (Remark 1 is elementary order theory), the convex-hull geometry (it is the known scalarization gap), or the usefulness of randomization under worst-case objectives (well established). A reader familiar with any of the three literatures above will recognize the ingredients; the claim is about the specific object they assemble here and, principally, about item 2.

7. Discussion

7.1. Negative certificates

A negative certificate — a rejection with an exhibited violated constraint for each action, exhausting the action set — is a complete verdict, not an inconclusive one. Theorem 5(7) is an instance: four actions, four exhibited violations, certifying impossibility on I together with the resource threshold of Section 5.6 and the per-weight licensing thresholds of Theorem 5(6). Related scored forecast-evaluation studies apply assessment discipline of the same kind to the Northern cod and Edwards Aquifer systems (Abaee, 2026, *Does a Surplus-Production Ladder Improve Forecasts of Northern Cod?*; Abaee, 2026, *Does a One-Pool Water-Balance Model Improve Forecasts of Edwards Aquifer Head?*); the cited studies are referenced for related work only.

7.2. Limitations

- (i) The separation theorem is an existence result with an open region; it does not imply the gap is large in any given application, and where the assessments coincide it yields nothing.
- (ii) Novelty statements reflect bounded literature search (Section 5.2).
- (iii) The operators cover finite horizons with exact tubes and specified disturbance sets; infinite horizons, partial observation, stochastic chance constraints, and endogenous event times are not treated.
- (iv) The finite-menu theorem of Section 4.4 governs plans whose admissibility is characterized by a required-margin vector; plans outside that class (reserve-financed plans, destination-failing plans) are handled separately, and the geometric form of the gap is established within the margin class.
- (v) If a convexified (blended) plan is physically and institutionally available, the gap disappears by Theorem 9; the separation is therefore a statement about certification relative to the available plan menu, not a claim that the gap is unavoidable in every institution.
- (vi) Conservatism of the benchmark’s certified tubes (Section 6.3) is regime-conditional: the exact inequality that makes the tubes conservative uses the Schaefer surplus’s monotonicity below $K/2$. Depensation in the stock–recruitment relationship, or heatwave mortality amplified nonlinearly at low stock, would reduce the recovery surplus, and the certified tubes would then no longer be conservative for the nonlinear realization; the certification statements of Section 4, which concern the datum rather than the realization, are unaffected.
- (vii) No empirical claims are made.
- (viii) The governance, intergenerational, and composition extensions are stated at partial status in the Supplementary Material and are not used in the proofs.
- (ix) Computational tractability: on a finite explicit state–action–disturbance graph with constant-time predicate evaluation, the backward recursion is polynomial in the graph size and horizon. For continuous, hybrid, or belief-state models, the graph itself may be exponentially large, infinite, or only approximately representable; exact-tube inclusion may be computationally hard or undecidable; and grid cardinality grows as $N_{\text{grid}} = \prod_{i=1}^n N_i$ — exponentially in dimension when floors are coordinates. The witness is tractable because its datum is small, finite, and rational.
- (x) The off-diagonal closure of Section 5.8 decides every cell of the 64-state grid by finite rational proof, not closed form: the curve ρ , the slice-flip total, and interior σ^* values are transcendental and carry certified rational brackets only.
- (xi) The σ -extension of Section 5.8 states two limits. Off the diagonal, the geometric member’s cover comparison is log-transcendental, and the verifier decides it only on proven sufficient and necessary rational conditions — the headline results are diagonal, so nothing load-bearing rests on the residual undecided cells; the off-diagonal closure decides every such cell by finite rational proof. And interior- θ values of $\sigma^*(z)$ are bracketed between tested rungs — the brackets are exact and strictness-refined; the interior value is the master root, transcendental in general.
- (xii) The common-shock extension of Section 5.9 states its limits: everything is on the witness datum’s decomposition (tent events, exact tubes, the specific STAGED mechanics, the Section 6.3 arithmetic); no claim about other data beyond the δ -family verified. The two regime-(iii) readings (additive/replace) bracket the coupled truth — Theorem CT is reading-conditional, and both are stated rather than choosing. Grid verification is fail-loud but finite, the closed

forms carry proofs (affine interval arithmetic; the kink device), and the two were cross-checked against each other at every grid point — the standard the σ -extension set. κ^* under the common class is adjudicated on a reserve grid of step $1/40$ plus closed-form confirmation on the survivors, and the infinity verdict on $\{\min(s) < 3/8\}$ is proved structurally, not sampled; no stochastic, partial-observation, or infinite-horizon claim is made, and the disturbance quantifier remains innermost throughout (Section 2.3).

- (xiii) The continuum slice-minimum of the fixed-sum tolerance structure — where a slice's best point lies — is closed (the continuum slice closure of Section 5.8): the slice-flip total is certified at $t^* = a^* + b^*$ — the tangent pair of the boundary involution — to ten places, by two independent code paths; the accepting windows are intervals (each a sublevel interval of a monotone phase of $\phi = s + \rho(s)$), so the interior window's ecological content (moderate concentration of a deficit is more certifiable than balance) is certified on the continuum, not only at the scanned points; the total-3 regime boundary is proved in the interior ($\phi < 3$); the harmonic member's region is exactly the circle $(2s_1 - 1)^2 + (2s_2 - 1)^2 \geq 10$, so nothing at or below $\theta = -1$ covers below total 3 and the m-side flip total is exactly 3; $\sigma_{\min}(t) \in (1/2, 1]$ on $(t^*, 3)$, with $\sigma_{\min}(707/250) \in [12/13, 1]$ and $\inf \sigma^* = 0$ beyond total 3 (corner cusps). The transcendental quantities (a^* , b^* , t^* , the interior σ^* values) carry certified rational brackets only.




Figure 4: The witness datum realized as a fishery-transition benchmark (Section 6.3). (a) Under the pulse-and-closure plan, the composite index $s_1 + s_2$ at $w = (1, 1)$ remains certified at every instant (minimum $2/5$), while the ecological margin breaches its floor mid-transition ($-4/5$ under the heatwave strike, $-3/10$ without it): the accepted-at-every-weighting reading is blind to the breach. (b) The plan menu as management schedules: pulse plus closed season; sustained-yield status quo; and the reserve-financed staged transition, which rebuilds the stock while improving both margins (the rescue set). All quota, stock, margin, and fund values are the exact rational values verified in Section 6.3.

7.3. A resource-transition benchmark

The witness datum of Section 4.5 is deliberately spare: a fund, two typed floors, four plans, and a tube semantics. This section realizes the same datum as a standard biomass–yield resource model, so that the reader can read every symbol as a quantity a resource agency actually monitors, and so that the theorems can be seen to bear on a textbook dynamical system rather than on a

discrete datum alone. The instantiation is an exact rational re-reading of the witness datum in the language of a Schaefer (1954) production model: every number below is re-derivable from the deposited benchmark script, which verifies it in exact rational arithmetic, and the certified tubes are conservative for the nonlinear realization by an exact inequality. It is not a fitted case study, and no empirical claim is made.

Read as a decision problem, the benchmark has the structure of a certification study in multi-criteria decision aiding. The decision maker is the resource authority; the decision alternatives are the menu’s management schedules (NO-SWITCH, FAST, SLOW, and the reserve-financed STAGED); the criteria are the two typed floors — the biomass margin above the limit reference point and the fleet income margin — evaluated along the whole transition under the declared disturbance; and the scalarization weights are the preference orderings of the assessment doctrine under which the composite index is read. The two criteria are conflicting objectives — each of the two principal schedules protects one at the expense of the other — and each is defended by a distinct constituency, so the weight space reads both as an assessor’s calibration and as the preference profile of the stakeholder dispute. The acceptance gap is then a robustness failure of the decision-aiding process in the precise sense of Theorem 5: at every preference weighting the index licenses some alternative — possibly a different one at different weightings — while no single alternative satisfies the criteria path-wise. The deliverables are ones a decision analyst can communicate: the licensing thresholds — the preference ranges over which each alternative is certified — and the rescue threshold — the adjustment funding that converts an unlicensable state into a staged transition. At the asymmetric allocations of the substitutability extension (Section 5.8), the same total margin draws opposite verdicts and opposite management responses, so the certification boundary is operationally decision-relevant.

The system. Let $B(t)$ denote spawning biomass (kt) in a stock-specific fishery over one annual review period $t \in [0, 1]$, with logistic surplus production $\sigma(B) = rB(1 - B/K)$ at $r = 4$, $K = 10$, and a biomass limit $B_{\text{lim}} = 2$ kt below which the spawning stock is judged unable to sustain recruitment. The ecological floor is the margin $s_1(t) = B(t) - B_{\text{lim}} \geq 0$; the income floor $s_2(t) \geq 0$ is the fleet income margin net of costs, maintained by the quota and by the transition reserve $x(t)$, an adjustment fund financing fleet draw-down. Four plans are available, standard in transition management: **NO-SWITCH** (status-quo harvest), **FAST** (an immediate large quota cut — a pulse season followed by a closure — accepting a mid-season trough on the ecological floor, the plan’s certified exposure), **SLOW** (a gradual reduction, shifting the burden onto the income floor), and **STAGED** (a below-sustained-yield quota while the reserve finances the fleet transition, after which both margins improve). Disturbances are environmental shocks: a mid-season heatwave strike ($t = 1/2$) imposing an excess mortality $\delta_0 = 1/2$ kt on the stock, present or absent. The state is observed only through the institutional dashboard, the weighted aggregate $w \cdot (s_1, s_2)$, and the reviewer checks the dashboard, not the floors. This is a deliberately adverse reporting architecture, adopted to isolate the assessment-operator issue — not a claim about current fisheries assessment practice, in which a limit-reference-point breach triggers a mandatory management response.

The datum, in resource terms. At the datum’s initial state $(x, s_1, s_2) = (1/2, 6/5, 6/5)$ — a half-financed fund, a stock $B = B_{\text{lim}} + 6/5 = 16/5$ kt ($6/5$ kt = 1.2 above the limit), and a healthy income margin — the plan menu of Section 4.5 reads: **FAST** cuts the quota hard for half a season; with the income floor protected by its quota schedule, FAST’s exposure is the ecological floor, whose worst-case (dip 2) trajectory is $s_1: 6/5 \rightarrow -4/5 \rightarrow 6/5$ — the stock troughs at $B_{\text{lim}} + (-4/5) = 6/5$ kt, i.e. $4/5$ kt below the limit ($s_1 = 6/5 - 2 = -4/5$ relative to the floor).

SLOW mirrors the exposure onto the income floor ($s_2 : 6/5 \rightarrow -4/5 \rightarrow 6/5$). **STAGED** spends the fund at the buy-back cost $c = 1$ while both margins improve by the destination gain $e = 1/4$. These are exactly the tube tables verified in Section 4.5 and in the deposited artifact; the interpretation is new, the arithmetic is not.

The benchmark schedule, exactly. The FAST plan is realized by the announced quota $H^*(t)$ that makes the benign margin line $s_1: 6/5 \rightarrow 6/5 - 3/2 \rightarrow 6/5$ — the tube-table line with the heatwave’s 1/2-kt dip removed — an exact trajectory of $\dot{B} = \sigma(B) - H^*(t)$: on $(0, 1/2)$, $H^* = 3 + \sigma(B_{\text{line}}(t))$ (a pulse declining from $\frac{1463}{125} \approx 11.7$ to $\frac{2161}{250} \approx 8.6$ kt/yr — admissible against the fleet capacity $H_{\text{max}} = 13$), and on $(1/2, 1)$ a closure, $H^* = 0$. The benign branch then tracks the tube line exactly. The heatwave strike at $t = 1/2$ drops the stock by δ_0 , from the trough 17/10 kt to 6/5 kt — the adverse tube’s value — and the closure lets the surplus rebuild the stock: on the adverse recovery leg the surplus is at least $\sigma(6/5) = 528/125 > 4$ kt/yr, while the certified recovery needs only 4, so the realized nonlinear trajectory lies on or above the certified tube throughout; the certified tubes are conservative for the Schaefer realization.

The remaining plans are read off the same model: NO-SWITCH and SLOW hold $B = 16/5$ at the sustained-yield quota $\sigma(16/5) = 1088/125 \approx 8.7$ kt/yr — itself below the model’s maximum sustainable yield (MSY) of $\sigma(K/2) = rK/4 = 10$ kt/yr at $B = K/2$; STAGED (entered at the rescue witness, fund $x = 3/2$) taxes the sustained yield by the margin improvement, rebuilding the stock $16/5 \rightarrow 69/20$ while the fund draws down $3/2 \rightarrow 1/2$ and both typed margins improve by $e = 1/4$.

The dashboard reading. Under the heatwave branch of the licensed plan, the composite index at the equal weighting, $s_1 + s_2$, has the tube values $(12/5, 2/5, 12/5)$: the dashboard never drops below 2/5, and the transition is certified at *every* weighting (Section 4.5’s per-weight thresholds $\rho_1 = 2/3$, $\rho_2 = 3/2$ place FAST inside every assessor’s acceptance set at income-heavy weights and SLOW inside every assessor’s at biomass-heavy weights, with their overlap containing $r = 1$). The floor itself, meanwhile, is breached mid-transition on both branches ($-4/5$ with the strike, $-3/10$ without), and no plan in the menu keeps both floors path-wise — the aggregate dashboard certifies a transition that violates the biomass limit on both branches — with the strike and without it — and no weighting can detect this from the index alone.

The reserve threshold is the operational number: a fund below the buy-back cost $c = 1$ makes the staged transition unfinanceable ($\kappa^*(z) = 1 - x$ at the stricken states), so every state of the failure set $\{x < 1, s_1 < 2, s_2 < 2\}$ carries the finite shortfall $\kappa^* = 1 - x$; the impossibility region is its part that the dashboard still certifies, and a modest top-up of the adjustment fund converts it into a rescue state. In management language: the dashboard is structurally unable to distinguish a transition that survives the heatwave from one that does not, and the financing gap, not the weighting, is what separates them (Figure 4).

Table 10 completes the translation guide of Table 4 into the benchmark’s own vocabulary: each of the paper’s central objects read as the fisheries quantity of this section.

A data anchor on a real stock. The benchmark’s units are not chosen in a vacuum: they are anchored to a real, regulated floor. For Northern cod (*Gadus morhua*, NAFO Divisions 2J3KL), the assessed spawning-stock series (NCAM M-shift, 1983–2015) carries a limit reference point (LRP) of 884.6 kt (DFO, 2016) — the biomass below which recruitment is judged impaired; the same series underlies the companion scored test of surplus-production forecasters on this stock (Abaee, 2026, *Does a Surplus-Production Ladder Improve Forecasts of Northern Cod?*). Reading the benchmark’s floor as this LRP fixes the dashboard arithmetic: the witness’s opening stock (1.6

Table 10: The translation guide (Table 4) extended to the fishery benchmark’s own quantities.

Object	Fisheries reading (this section)
Typed floors $s_i \geq 0$	the spawning-stock margin above the limit reference point ($B - B_{\text{lim}} \geq 0$; the LRP convention) and the fleet income margin net of costs
Scalarized operator E_w	the assessment dashboard $w \cdot (s_1, s_2)$ the reviewer checks
Licensing thresholds ρ_1, ρ_2	weight-ratio licensing intervals for the pulse (FAST) and gradual (SLOW) quota plans
Acceptance gap FP_{agg}	dashboard-certified yet LRP-breaching transitions (the index-blindness zone: trough 0.6 LRP units while the dashboard reads $2/5 > 0$)
Rescue set R ; threshold κ^*	the adjustment fund financing fleet buy-down (STAGED at cost $c = 1$; the shortfall $\kappa^* = 1 - x$)
Impossibility region I	underfunded transitions: fund $x < 1$, both margins below their plan margins, dashboard still certified
Blend collapse vs. time-sharing	fractional quota/effort allocation (convexification) vs. pulse/gradual season alternation
Exact-tube semantics	no transient LRP breach along the transition, the mid-season heatwave included

LRP units, $1.6 \times 884.6 \approx 1.42$ Mt) is the neighbourhood of the assessed mid-1980s Northern cod spawning biomass, the adverse plan’s trough ($B = 6/5$ kt, 0.6 LRP units ≈ 531 kt) lies below the LRP while remaining above the stock’s 2015 assessed level of about a third of the LRP (33.8%, ≈ 299 kt; DFO, 2016) — the zone between the two levels, in which the stock in fact resided after the collapse — and the adverse plan’s worst-case stock excursion — 2 kt, i.e. one full LRP unit (the quota’s 3/2-kt draw-down plus the heatwave’s 1/2-kt excess mortality) — is the order of the stock’s swing across the LRP itself. **The motivation, with the assessed numbers.** The anchoring statements of this subsection, made precise: the limit reference point of the assessed NCAM M-shift series is 884.6 kt — by construction the 1983–1989 mean spawning-stock biomass of that series (DFO, 2016, Table A2); the stock’s assessed 2015 level is 33.8% of that LRP (the advisory report’s “about a third”), ≈ 299 kt; the decline from the mid-1980s neighbourhood of the benchmark’s opening stock (≈ 1.42 Mt) to the 2015 level is a swing of ≈ 1.1 Mt, about 1.3 LRP units; and the benchmark’s adverse worst-case excursion — one full LRP unit — is therefore of the assessed order, with its trough (0.6 LRP) inside the assessed post-collapse band (0.34, 1) LRP units. The unit convention is explicit throughout: s_1 is a kt margin above B_{lim} , so the opening stock is $B(0) = B_{\text{lim}} + 6/5 = 16/5$ kt, the normalized stock is $b(0) = B(0)/B_{\text{lim}} = 8/5 = 1.6$ LRP units, and the margin $6/5$ kt is not itself a biomass ratio.

Three qualifications apply. The anchoring fixes units only: no parameter of the rational witness is estimated from data, and the theorems say nothing predictive about this or any other stock. And the companion scored test adds a caution of exactly the discipline this paper practises: on the same series, no structural surplus-production model in the scored ladder beat last-value persistence (one-year RMSE 98 kt versus 115–206 kt), and every model missed the collapse window — what survives institutional scrutiny on such a stock is the menu and the floors, both of which are set by management and both of which this paper’s datum takes as given, rather than a predictive fit of the dynamics.

8. Conclusions

Scalarized and coordinate-wise feasibility criteria are usually compared as doctrines — as different normative stances on substitutability, of which the weak- and strong-sustainability traditions are the canonical instances. This paper shows that the divergence between the two criteria as formalized here — the scalarized-aggregate and typed operators of Section 3.1 on a common action menu and disturbance class — survives translation into assessment mechanics, at the level of a theorem about those operators. The theorem establishes no ranking of doctrines: Section 5.1 delimits the formalizations, and by Theorem 9 the structural character of the separation is a property of the finite deterministic menu, not of either tradition (Sections 4.10–4.11). On a typed transition datum under exact-tube semantics, the compensatory reading (per-weight acceptance) and the noncompensatory reading (common-plan acceptance) are related by a quantifier commutation that can fail on an open region of state space. Where it fails, every scalarization weight certifies a transition, but no single transition is certified by all weights, and the certified set splits into states rescuable by a resource-controlled action and states that are impossible under every weight. The mechanism is not scalarization blindness — at fixed trajectories the full-cone aggregate is lossless — but the policy dependence of the aggregate-feasible transition. The finite-menu theorem of Section 4.4 states the underlying mechanism in full generality: for any finite menu, the acceptance gap is the part of the convex hull of the required margins that no single plan dominates, so the witness of Section 4.5 is an instance of a geometric fact rather than an isolated construction.

For the construction of composite sustainability indices, the theorem has a concrete consequence. An index can be sound at the level of accounting and still over-certify at the level of assessment, because certification is a quantifier statement about transitions, not a property of a number. Where separately-binding floors matter, per-floor reporting is not a presentation preference but a detection requirement: the aggregate alone cannot distinguish the rescue region from the impossibility region. The bridge a composite index needs — some single transition that serves every admissible weight, i.e. membership in $\mathcal{V}_{\text{typ}}$ — is exactly what the frameworks’ reporting conventions should be asked to exhibit. Theorem 9 qualifies this requirement: such a single transition is not, in general, a member of the specified deterministic menu but a convexified action, so the relevant question for a reporting convention is whether its action set admits the required menu convexification — and Proposition 10 shows that temporal alternation alone does not provide it.

Read from each of the three literatures this paper connects, the results take one specific form. For multi-criteria assessment, they locate the certification failure of compensatory aggregates in the acceptance protocol — not in weight choice, and not in compensability at fixed profiles. For decision making under uncertainty, the acceptance gap measures the value of information about the assessment weight, and Theorem 9 with Proposition 10 identifies menu convexification, not temporal sharing, as what removes it. For transition management, where separately-binding floors govern a transition, per-floor reporting along the path is the detection requirement the separation makes exact.

List of abbreviations

- CES — constant elasticity of substitution

- **DFO** — Fisheries and Oceans Canada
 - **LPI** — Living Planet Index
 - **LRP** — limit reference point
 - **MSY** — maximum sustainable yield
 - **NCAM** — Northern Cod Assessment Model
 - **PROMETHEE** — Preference Ranking Organisation Method for Enrichment Evaluation
-

Appendix A. Proofs of the general results

Proof of Remark 1 (Section 4.1). If $a \in \bigcap_{\lambda} E_{\lambda}(z)$, then each $E_{\lambda}(z)$ is nonempty, so z lies in the right-hand set. The inclusion fails at z precisely when each $E_{\lambda}(z)$ is nonempty but $\bigcap_{\lambda} E_{\lambda}(z) = \emptyset$, i.e. when no common selector exists; this is the stated equality condition. $\square$

Proof of Remark 2 (Section 4.2). $(\Rightarrow)$ $w \geq 0$, $w \neq 0$, $v \geq 0$ gives $w \cdot v = \sum_i w_i v_i \geq 0$. $(\Leftarrow)$ If $v_k < 0$, take $w = e_k \in W_+$: then $w \cdot v = v_k < 0$. $\square$

Proof of Proposition 3 (Section 4.3). (i) Let $a \in E_{\text{typ}}(z)$. For every disturbance d , $\text{Tube}(a, d) \subseteq S = S^{\text{phys}} \cap \{s \geq 0\}$. By Remark 2, every tube point satisfies $w \cdot s \geq 0$ for every $w \in W_+$, and likewise every successor state lies in G^w ; hence $a \in E_w(z)$. The remaining inclusions follow from $S^w \subseteq S^{\text{phys}}$, $G^w \subseteq G^{\text{phys}}$, and $\text{End}(a, d) \subseteq \text{Tube}(a, d)$. (ii) The inclusion $\subseteq$ is part (i) intersected over w . For the reverse inclusion, let $a \in \bigcap_{w \in W_+} E_w(z)$. For every d and every tube point $p \in \text{Tube}(a, d)$, every $w \in W_+$ gives $w \cdot s(p) \geq 0$, so by Remark 2 $s(p) \geq 0$ componentwise; hence $\text{Tube}(a, d) \subseteq S$. The same argument with Succ gives $\text{Succ}(a, d) \subseteq G$. Thus $a \in E_{\text{typ}}(z)$. $\square$

Proof of Proposition 4 (Section 4.4). For the second inclusion: by Proposition 3(i), $E_w(z) \subseteq E_{\text{tube,phys}}(z)$ for every w , so $\mathcal{V}_w \subseteq \mathcal{V}_{\text{phys}}$, and intersecting over $w \in W$ preserves the inclusion. For the first: if $z \in \mathcal{V}_{\text{typ}}$, then by Proposition 3(ii) there exists $a \in \bigcap_{w \in W_+} E_w(z)$; this same a lies in $\bigcap_{w \in W} E_w(z)$, so $E_w(z) \neq \emptyset$ for every $w \in W$, i.e. $z \in \mathcal{V}_W$. Monotonicity: intersecting over a larger family can only shrink the intersection. $\square$

Proof of the finite-menu theorem (Section 4.4). (i) is the margin test restated. For (ii), if $s = d + v$ with $d \in \text{conv}(\mathcal{D})$ and $v \geq 0$, then for every $w \geq 0$, $w \cdot s \geq w \cdot d \geq \min_a w \cdot d_a$, so the scalarized test holds. Conversely, suppose $s \notin \text{conv}(\mathcal{D}) + \mathbb{R}_+^n$; this set is closed and convex (a compact polytope plus a closed cone), so strict separation gives $w \neq 0$ with $w \cdot s < w \cdot d + w \cdot v$ for every $d \in \text{conv}(\mathcal{D})$ and every $v \geq 0$. Since v ranges over the whole cone, $w \in \mathbb{R}_+^n \setminus \{0\}$, and taking $v = 0$ gives $w \cdot s < \min_{d \in \text{conv}(\mathcal{D})} w \cdot d = \min_a w \cdot d_a$: the scalarized test fails. $\square$

Appendix B. Proofs for the witness datum

Proof of Theorem 5 (Section 4.6). (1) Typed admissibility requires the worst-case tube in S_0 and the successor in G . NO-SWITCH never reaches G (its successor retains $q = 0$), so it is never admissible. For FAST, the worst-case s_1 -tube is $[s_1 - 2, s_1]$, safe if and only if $s_1 \geq 2$; its successor $(1, x, s + e)$ lies in G given $x \geq 0$. Hence FAST is typed-admissible exactly when $s_1 \geq 2$. Symmetrically, SLOW is typed-admissible exactly when $s_2 \geq 2$. For STAGED, the floors grow

through the interval and the successor lies in G ; the binding constraint is the x -tube $[x-1, x]$, safe if and only if $x \geq 1$. Therefore $E_{\text{typ}}(z) \neq \emptyset$ exactly on $\{x \geq 1\} \cup \{s_1 \geq 2\} \cup \{s_2 \geq 2\}$.

(2) Fix $w = (w_1, w_2) \in W_+$. NO-SWITCH misses G^w and is never admissible. FAST is w -admissible if and only if its worst-case aggregate value stays nonnegative: $w_1(s_1 - 2) + w_2 s_2 \geq 0$, i.e. $w_1 s_1 + w_2 s_2 \geq 2w_1$ (automatic when $w_1 = 0$). Symmetrically SLOW is w -admissible if and only if $w_1 s_1 + w_2 s_2 \geq 2w_2$. STAGED is w -admissible if and only if $x \geq 1$ (the physical constraint within S^w), since $w \cdot s \geq 0$ holds automatically along its tube. Hence $z \in \mathcal{V}_w$ whenever $x \geq 1$; and if $x < 1$, then $z \in \mathcal{V}_w$ if and only if at least one of the two inequalities above holds.

Suppose first $x \geq 1$: then $z \in \mathcal{V}_w$ for every w , so $z \in \mathcal{V}_{\text{weak}}$. Suppose $x < 1$ and $s_1 \geq 2$ or $s_2 \geq 2$: say $s_1 \geq 2$; then FAST is typed-admissible by (1), hence w -admissible for every w , so $z \in \mathcal{V}_{\text{weak}}$. It remains to treat $x < 1$ with $s_1 < 2$ and $s_2 < 2$. For $w_1 > 0$, write $r = w_2/w_1$: FAST is w -admissible exactly when $s_1 + r s_2 \geq 2$, i.e. $r \geq (2 - s_1)/s_2 = \rho_1$ (using $s_2 > 0$; when $s_2 = 0$ the condition reduces to $s_1 \geq 2$, which is excluded). SLOW is w -admissible exactly when $s_1 + r s_2 \geq 2r$, i.e. $r \leq s_1/(2 - s_2) = \rho_2$ (using $s_2 < 2$). The boundary weights are the same limiting cases: at $w_1 = 0$ ($r \rightarrow \infty$) FAST is licensed provided $s_2 \geq 0$, while SLOW is strictly rejected on the region $s_2 < 2$; symmetrically, at $w_2 = 0$ ($r = 0$) SLOW is licensed provided $s_1 \geq 0$, while FAST is rejected on the region $s_1 < 2$. Therefore $z \in \mathcal{V}_w$ for every $w \in W_+$ if and only if the intervals $\{r \geq \rho_1\}$ and $\{r \leq \rho_2\}$ together cover all of $[0, \infty]$, which holds exactly when $\rho_2 \geq \rho_1$. Computing,

$$\rho_2 \geq \rho_1 \iff \frac{s_1}{2 - s_2} \geq \frac{2 - s_1}{s_2} \iff s_1 s_2 \geq (2 - s_1)(2 - s_2) \iff s_1 + s_2 \geq 2,$$

with both denominators positive in the present case. Hence, in this case, $z \in \mathcal{V}_{\text{weak}}$ if and only if $s_1 + s_2 \geq 2$. Assembling the three cases, and noting $\{s_1 \geq 2\} \cup \{s_2 \geq 2\} \subseteq \{s_1 + s_2 \geq 2\}$ under nonnegativity,

$$\mathcal{V}_{\text{weak}} = \{x \geq 1\} \cup \{s_1 + s_2 \geq 2\}.$$

(3) In the physical operators the only constraint is $x \geq 0$ (with $q = 1$ at the destination). FAST and SLOW have tubes within S^{phys} and successors in G^{phys} for every $z \in X_0$ (their dips affect only s , which is unconstrained physically); STAGED does so exactly for $x \geq 1$. Hence $E_{\text{tube,phys}}(z) \neq \emptyset$ for every $z \in X_0$, so $\mathcal{V}_{\text{phys}} = X_0$. For the endpoint operator, the endpoint values of FAST and SLOW lie in S^{phys} for every $z \in X_0$, so $\mathcal{V}_{\text{end}} = X_0$ as well.

(4) By (1)–(2),

$$\begin{aligned} \text{FP}_{\text{agg}} &= \mathcal{V}_{\text{weak}} \setminus \mathcal{V}_{\text{typ}} = [\{x \geq 1\} \cup \{s_1 + s_2 \geq 2\}] \setminus [\{x \geq 1\} \cup \{s_1 \geq 2\} \cup \{s_2 \geq 2\}] \\ &= \{x < 1, s_1 < 2, s_2 < 2, s_1 + s_2 \geq 2\} = I. \end{aligned}$$

This set is nonempty (e.g. $(x, s_1, s_2) = (\frac{1}{2}, \frac{6}{5}, \frac{6}{5})$), and its relatively open interior is the stated region.

(5) By (4), $I = \text{FP}_{\text{agg}} = \mathcal{V}_{\text{weak}} \setminus \mathcal{V}_{\text{typ}} \neq \emptyset$, so the first inclusion is strict. For the second, the point $(\frac{1}{2}, \frac{1}{10}, \frac{1}{10})$ satisfies $x < 1$ and $s_1 + s_2 = \frac{1}{5} < 2$, so it lies outside $\mathcal{V}_{\text{weak}}$ by (2), while by (3) it lies in $\mathcal{V}_{\text{phys}}$.

(6) The threshold characterization was derived in the proof of (2): with $r = w_2/w_1$, FAST is licensed exactly for $r \geq \rho_1$ and SLOW exactly for $r \leq \rho_2$, and $\rho_2 \geq \rho_1 \iff s_1 + s_2 \geq 2$. On $\mathcal{Q}$ the inequality $s_1 + s_2 \geq 2$ holds, so the two licensing intervals overlap and together cover all weights; strictly on the relative interior ($s_1 + s_2 > 2$) both thresholds are finite and the intervals overlap in a nondegenerate interval, with $r > \rho_2$ licensing FAST only and $r < \rho_1$ licensing SLOW

only. On I (where $x < 1$), STAGED is inadmissible for every w , so $E_{\text{typ}}(z) = \bigcap_w E_w(z) = \emptyset$ by Proposition 3(ii): no single action serves every weight. On R (where $x \geq 1$), STAGED lies in $E_{\text{typ}}(z) = \bigcap_w E_w(z)$ and serves every weight.

(7) On R , STAGED's tube is $[x - 1, x] \times \{s \rightarrow s + e\}$: the x -tube stays nonnegative because $x \geq 1$, the floors grow, and the successor $(1, x - 1, s + e) \in G$. On I , the four exhibited violations exhaust the action set and each violates a typed constraint: FAST's s_1 -tube reaches $s_1 - 2 < 0$; SLOW's s_2 -tube reaches $s_2 - 2 < 0$; STAGED's x -tube reaches $x - 1 < 0$; NO-SWITCH's successor misses G . $\square$

Proof of Remark 7 (Section 4.8). (i) Backward induction on stages. Base: at the final stage the terminal sets satisfy $G \subseteq G^w \subseteq G^{\text{phys}}$ by Remark 2 and the definition of the terminal sets. Step: the accepted set of a stage is the set of states from which some action keeps the tube in the safe set and the successor in the next-stage accepted set; applying Proposition 3(i) to the stage's action sets and the induction hypothesis to the next-stage accepted sets preserves the three-way inclusion at the stage level. (ii) With HOLD the unique prefix action, a prefix state is accepted exactly when it lies in the prefix safe set and in the next-stage accepted set; the stage-0 regions are therefore the witness regions of Theorem 5 pulled back through the (identity) holds, and the strictness witnesses ($I \neq \emptyset$ and the point $(\frac{1}{2}, \frac{1}{10}, \frac{1}{10})$) persist. (iii) The three listed mechanisms are each sufficient to erase the gap in a multi-stage setting, so no unconditional generalization is asserted. $\square$

Proof of Theorem 8 (Section 4.8). *Datum.* Two capital forms, floors at zero, the bridge stock absent ($x \equiv 0$); states are pairs $s = (s_1, s_2)$. Stage 2 (final interval): safe set $S_0^{(2)} = \mathbb{R}^2$ (no path constraint), terminal sets $G^{(2)} = G^{w,(2)} = G = \{s \geq 0\}$ — coincident for every weight — and a single action REPAIR with tube $\{s\}$ (safe trivially) and successor $(1, 1) \in G$. Stage 1: safe set $S_0 = \{s \geq 0\}$; two actions with constant tubes, safe at every state of S_0 : A1 with successor $(s_1 - 1, s_2 + 1)$ and A2 with successor $(s_1 + 1, s_2 - 1)$; the stage-1 terminal architecture is the stage-2 accepted sets. Initial state $z^* = (\frac{2}{5}, \frac{2}{5})$.

In the two-stage datum the gap is erased. REPAIR is typed-admissible — hence w -admissible for every weight — from every state, so $\mathcal{V}_{\text{typ}}^{(2)} = \mathcal{V}_w^{(2)} = \mathbb{R}^2$ for every w : the stage-2 accepted sets coincide. Backward induction then admits both A1 and A2 at z^* (safe tubes; successors in $\mathcal{V}_{\text{typ}}^{(2)} = \mathbb{R}^2$), so z^* is typed-accepted at stage 0, and weak acceptance follows from typed acceptance.

In the single-interval truncation — the same stage-1 datum with the final interval collapsed to its terminal sets, $G = \{s \geq 0\}$ and $G^w = \{w \cdot s \geq 0\}$ — the gap is present. Both successors leave G ($s_1 - 1 = -\frac{3}{5} < 0$ and $s_2 - 1 = -\frac{3}{5} < 0$), so z^* is typed-rejected. For the weak doctrine, writing $r = w_2/w_1$ on the interior of the weight cone: A1 is w -admissible exactly when $(s_1 - 1) + r(s_2 + 1) \geq 0$, i.e. $r \geq \frac{1-s_1}{s_2+1} = \frac{3}{7}$; A2 exactly when $(s_1 + 1) + r(s_2 - 1) \geq 0$, i.e. $r \leq \frac{s_1+1}{1-s_2} = \frac{7}{3}$; the two intervals cover $[0, \infty]$ because $\frac{3}{7} < \frac{7}{3}$, and the boundary weights are covered directly ($r = 0$: A2 gives $w_1(s_1 + 1) \geq 0$; $r \rightarrow \infty$: A1 gives $w_2(s_2 + 1) \geq 0$). Hence $z^* \in \mathcal{V}_{\text{weak}}$ in the truncation, and the gap $\mathcal{V}_{\text{weak}} \setminus \mathcal{V}_{\text{typ}}$ contains z^* . $\square$

Proof of Theorem 9 (Section 4.10). (i) FAST's worst-case tube has s_1 -coordinates $[s_1 - 2, s_1]$ and s_2 -coordinate $\{s_2\}$; SLOW's has s_1 -coordinate $\{s_1\}$ and s_2 -coordinates $[s_2 - 2, s_2]$; both keep x nonnegative throughout (their dips affect only s). The blended tube therefore has s_1 -coordinates $[s_1 - 2\delta, s_1]$, s_2 -coordinates $[s_2 - 2(1 - \delta), s_2]$, and nonnegative x -coordinates (convex combinations of nonnegative quantities). Typed admissibility requires the tube within $S_0 = \{x \geq 0, s \geq 0\}$ and the successor in G . The tube condition is exactly the displayed window. The successor is the mixture of the two successors, which the datum fixes to the same value — the action table of

Section 4.5 gives both FAST and SLOW the successor $(1, x, s + e)$ — so the blended successor is $(1, x, s + e) \in G$: $x \geq 0$ and the destination label is 1. The window is nonempty exactly when $1 - s_2/2 \leq s_1/2$, i.e. $s_1 + s_2 \geq 2$; on $\mathcal{Q}$ the raw window already lies in $(0, 1)$ ($s_1 < 2, s_2 < 2$ give $0 < 1 - s_2/2$ and $s_1/2 < 1$), so the $[0, 1]$ clip in the statement is inactive there; at equality it is the singleton $\delta = s_1/2 = 1 - s_2/2$, feasible because S_0 is closed. On the boundary faces ($s_1 = 0$ or $s_2 = 0$) the window is closed at the corresponding endpoint, and the classification follows the boundary conventions of Section 4.6.

- (ii) If $s_1 + s_2 \geq 2$, (i) supplies a typed-admissible blend whatever the value of $x \geq 0$; if $s_1 + s_2 < 2$, no blend is typed-admissible by (i), and the deterministic menu contributes exactly $\mathcal{V}_{\text{typ}} = \{x \geq 1\} \cup \{s_1 \geq 2\} \cup \{s_2 \geq 2\} \subseteq \{x \geq 1\} \cup \{s_1 + s_2 \geq 2\}$. Hence the augmented typed set equals $\{x \geq 1\} \cup \{s_1 + s_2 \geq 2\} = \mathcal{V}_{\text{weak}}$ by Theorem 5(2) — and nothing beyond, since every admissible blend satisfies $s_1 + s_2 \geq 2$ and the deterministic actions stay within $\mathcal{V}_{\text{typ}}$.
- (iii) Immediate from the window's state dependence and from typed admissibility implying w -admissibility for every weight.

□

Proof of Proposition 10 (Section 4.11). A sequential policy that spends any positive time on FAST visits every s_1 in $[s_1 - 2, s_1]$ and any positive time on SLOW visits every s_2 in $[s_2 - 2, s_2]$; the union tube stays within $S_0 = \{x \geq 0, s \geq 0\}$ exactly when both dips are nonnegative, i.e. $s_1 \geq 2$ and $s_2 \geq 2$. On I both coordinates are strictly below 2, so the union tube is never within S_0 and no sequential policy is typed-admissible; since every member of the finite deterministic menu is typed-admissible on I only in the per-weight sense of Theorem 5(6) (and never in the common sense), the noncompensatory accepted set does not extend to I under sequential time-sharing. □

Appendix C. Proof of the rescue threshold

Proof of the rescue threshold (Section 5.7). Among the augmented menu, only STAGED_κ depends on κ , and its typed-admissibility is independent of the floors: its x -tube is $[x + \kappa - 1, x + \kappa]$, the floors are nondecreasing (they grow to $s + e$ with $e = (\frac{1}{4}, \frac{1}{4}) > 0$), and the successor $(1, x + \kappa - 1, s + e)$ lies in G exactly when $x + \kappa - 1 \geq 0$. Hence $\text{STAGED}_\kappa \in E_{\text{typ}}(z) \iff x + \kappa \geq 1$, for every z . If $z \in \mathcal{V}_{\text{typ}}$, some base action of $\mathcal{A} \subseteq \mathcal{A}_\kappa$ is admissible at $\kappa = 0$, so $\kappa^*(z) = 0$. If $z \notin \mathcal{V}_{\text{typ}}$ — that is, $x < 1$ and $s_1 < 2$ and $s_2 < 2$ — then NO-SWITCH, FAST, and SLOW all fail (Theorem 5(1)), the only admissible element of $\mathcal{A}_\kappa$ is STAGED_κ , and it is admissible exactly for $\kappa \geq 1 - x$; hence $\kappa^*(z) = 1 - x$. □

References

- Abaee, A., 2026a. Governance delay and the stability of harvested stocks: mobilising and protective feedback rules, and the review interval as a design parameter. Zenodo. <https://doi.org/10.5281/zenodo.22554217>. Companion delay-dynamics study.
- Abaee, A., 2026b. Periodic review as sampled governance: sample-and-hold dynamics of assessment-driven effort control, a selected 42-stock spectral screen, and the Northern Cod case. Zenodo. <https://doi.org/10.5281/zenodo.22554297>. Companion review-screen study.
- Abaee, A., 2026c. Aggregate indices and transition safety: a quantifier-order separation between scalarized and coordinate-wise feasibility. Zenodo. <https://doi.org/10.5281/zenodo.22545740>. Companion assessment-separation study.

- Abaee, A., 2026d. Certifying a typed ledger: the predicates, the programmes, the vintages, and the reproduction bundle.
- Abaee, A., 2026e.
- Abaee, A. (2026). *Does a one-pool water-balance model improve forecasts of Edwards Aquifer head? A scored test at J-17*. Zenodo. <https://doi.org/10.5281/zenodo.22552680>.
- Abaee, A. (2026). *Does a surplus-production ladder improve forecasts of Northern cod? A scored test on NAFO 2J3KL*. Zenodo. <https://doi.org/10.5281/zenodo.22553609>.
- Abaee, A. (2026). *How Aggregation Can Conceal Composition: Aggregate Biocapacity and the Identifiability of Modelled Ecological-Capital Drawdown*. Manuscript; publication status not independently verified here.
- Abaee, A. (2026). *Typed flux ledgers and depletion arithmetic: conservation, componentwise diagnostics, and the semantics of depletion horizons*. Zenodo. <https://doi.org/10.5281/zenodo.22554177>.
- Abaee, A. (2026). *Verification code and figure pipeline for Aggregate Indices and Transition Safety: A Quantifier-Order Separation Between Scalarized and Coordinate-Wise Feasibility*. figshare. <https://doi.org/10.6084/m9.figshare.33764023>.
- World Bank. (2011). *The Changing Wealth of Nations: Measuring Sustainable Development in the New Millennium*. World Bank, Washington, D.C.
- Ahuja, R.K., Magnanti, T.L., Orlin, J.B., 1993. *Network Flows: Theory, Algorithms, and Applications*.
- Asheim, G. B. (1994). Net national product as an indicator of sustainability. *Scandinavian Journal of Economics*, 96(2), 257–265.
- Aubin, J.-P., and Catté, F. (2002). Bilateral fixed-points and algebraic properties of viability kernels and capture basins of sets. *Set-Valued Analysis*, 10(4), 379–416.
- Aubin, J.-P., Bayen, A. M., and Saint-Pierre, P. (2011). *Viability Theory: New Directions*, 2nd ed. Birkhäuser, Boston.
- Aubin, J.-P. (1991). *Viability Theory*. Birkhäuser, Boston.
- Ayres, R.U. (1996). Viewpoint: weak versus strong sustainability. Tinbergen Institute Discussion Paper 98-103/3.
- Baez, J., Li, X., Libkind, S., Osgood, N.D., Patterson, E., 2023. Compositional modeling with stock and flow diagrams. In: *Applied Category Theory 2022*. Electronic Proceedings in Theoretical Computer Science 380, 77–96. <https://doi.org/10.4204/EPTCS.380.5>
- Blomqvist, L., Brook, B.W., Ellis, E.C., Kareiva, P.M., Nordhaus, T., Shellenberger, M., 2013. Does the shoe fit? Real versus imagined ecological footprints.
- Bartlett, A.A., 2000. An analysis of U.S. and world oil production patterns using Hubbert-style curves. *Mathematical Geology* **32**(1), 1–17.
- Beckerman, W. (1994). "Sustainable development": is it a useful concept? *Environmental Values* **3**(3), 191–209.
- Becker, W., Saisana, M., Paruolo, P., and Vandecasteele, I. (2017). Weights and importance in composite indicators: Closing the gap. *Ecological Indicators*, 80, 12–22.
- Ben-Tal, A., El Ghaoui, L. and Nemirovski, A. (2009). *Robust Optimization*. Princeton Series in Applied Mathematics, vol. 28. Princeton University Press.
- Ben-Tal, A., Goryashko, A., Guslitser, E. and Nemirovski, A. (2004). Adjustable robust solutions of uncertain linear programs. *Mathematical Programming* **99**(2), 351–376.

- Bertsimas, D. and Goyal, V. (2012). On the power and limitations of affine policies in two-stage adaptive optimization. *Mathematical Programming* **134**(2), 491–531.
- Boos, A. (2015). Genuine savings as an indicator for “weak” sustainability: Critical survey and possible ways forward in measuring weak sustainability. *Sustainability*, 7(4), 4146–4163.
- Brunner, P.H. and Rechberger, H., 2016. *Handbook of Material Flow Analysis: For Environmental, Resource, and Waste Engineers*, 2nd edn.
- Cairns, R. D., and Martinet, V. (2014). An environmental-economic measure of sustainable development. *European Economic Review*, 69, 4–17.
- Cardaliaguet, P. (1996). A differential game with two players and one target. *SIAM Journal on Control and Optimization*, 34(4), 1441–1460.
- Cardaliaguet, P., Quincampoix, M., and Saint-Pierre, P. (1999). Set-valued numerical analysis for optimal control and differential games. In: *Stochastic and Differential Games: Theory and Numerical Methods*, Annals of the International Society of Dynamic Games, vol. 4, Birkhäuser, Boston, 177–247.
- Chen, M., Hansen, P., Jaumard, B. and Tuy, H. (1999). Solution of the Weber problem with attraction and repulsion.
- Chhikara, R.S., Folks, J.L., 1989. The Inverse Gaussian Distribution: Theory, Methodology, and Applications.
- Cinelli, M., Coles, S. R., and Kirwan, K. (2014). Analysis of the potentials of multi criteria decision analysis methods to conduct sustainability assessment. *Ecological Indicators*, 46, 138–148.
- Clark, C.W., 1990. *Mathematical Bioeconomics: The Optimal Management of Renewable Resources*, 2nd ed.
- Daly, H.E. (1995). On Wilfred Beckerman’s critique of strong sustainability.
- Daly, H. E. (1990). Toward some operational principles of sustainable development. *Ecological Economics*, 2(1), 1–6.
- Dasgupta, P., and Mäler, K.-G. (2000). Net national product, wealth, and social well-being. *Environment and Development Economics*, 5(1), 69–93.
- Das, I., and Dennis, J. E. (1997). A closer look at drawbacks of minimizing weighted sums of objectives for Pareto set generation in multicriteria optimization problems. *Structural Optimization*, 14, 63–69.
- Das, I. and Dennis, J.E. (1998). Normal-boundary intersection: a new method for generating the Pareto surface in nonlinear multicriteria optimization problems. *SIAM Journal on Optimization* **8**(3), 631–657.
- De Lara, M., and Doyen, L. (2008). *Sustainable Management of Natural Resources: Mathematical Models and Methods*.
- Dietz, S. and Neumayer, E. (2007). Weak and strong sustainability in the SEEA: concepts and measurement. *Ecological Economics* **61**(4), 617–626.
- Doyen, L., and Gajardo, P. (2020). Sustainability standards, multicriteria maximin, and viability. *Natural Resource Modeling*, 33(3), e12250.
- Doyen, L., Thébaud, O., Béné, C., Martinet, V., Gourguet, S., Bertignac, M., Fifas, S., and Blanchard, F. (2012). A stochastic viability approach to ecosystem-based fisheries management. *Ecological Economics*, 75, 32–42.
- Ehrgott, M. (2005). *Multicriteria Optimization*. 2nd ed. Springer.
- Ekins, P., Simon, S., Deutsch, L., Folke, C., and De Groot, R. (2003). A framework for the practical application of the concepts of critical natural capital and strong sustainability. *Ecological*

Economics, 44(2), 165–185.

Athan, T.W. and Papalambros, P.Y. (1996). A note on weighted criteria methods for compromise solutions in multi-objective optimization.

Fabian, M. J., Henrion, R., Kruger, A. Y., and Outrata, J. V. (2010). Error bounds: necessary and sufficient conditions. *Set-Valued Analysis*, 18(2), 121–149.

Fanning, A. L., O'Neill, D. W., Hickel, J., and Roux, N. (2022). The social shortfall and ecological overshoot of nations. *Nature Sustainability*, 5(1), 26–36.

Feinberg, M., 2019. Foundations of Chemical Reaction Network Theory.

Filippov, A. F. (1988). *Differential Equations with Discontinuous Righthand Sides*. Mathematics and Its Applications (Soviet Series), vol. 18. Kluwer Academic Publishers, Dordrecht.

Fischer-Kowalski, M., Krausmann, F., Giljum, S., Lutter, S., Mayer, A., Bringezu, S., Moriguchi, Y., Schütz, H., Schandl, H., Weisz, H., 2011. Methodology and indicators of economy-wide material flow accounting: state of the art and reliability across sources. *Journal of Industrial Ecology* 15, 855–876.

Fischer, S. M., Joy, M. K., Abrahamse, W., Milfont, T. L., and Petherick, L. M. (2022). The use and misuse of composite environmental indices. *bioRxiv*, 2022.03.15.484501. <https://doi.org/10.1101/2022.03.15.484501>.

Frankowska, H. (1989). Optimal trajectories associated with a solution of contingent Hamilton–Jacobi equations. *Applied Mathematics and Optimization*, 19, 291–311.

Gale, D., 1957. A theorem on flows in networks. *Pacific Journal of Mathematics* 7, 1073–1082.

Gao, P., Wang, Y., Wang, H., Song, C., Ye, S., and Wang, X. (2023). A Pareto front-based approach for constructing composite index of sustainability without weights: A comparative study of implementations. *Ecological Indicators*, 155, 110919. <https://doi.org/10.1016/j.ecolind.2023.110919>.

Illakwahhi, D.T., Vegi, M.R., Srivastava, B.B.L., 2024.

Hanley, N., Moffatt, I., Faichney, R., and Wilson, M. (1999). Measuring sustainability: A time series of alternative indicators for Scotland. *Ecological Economics*, 28(1), 55–73.

Hickel, J. (2020). The sustainable development index: Measuring the ecological efficiency of human development in the Anthropocene. *Ecological Economics*, 167, 106331.

Huang, L., Li, R. and Liu, J. (2007). A survey of multi-objective optimization methods.

Hubbert, M.K., 1956. Nuclear energy and the fossil fuels. *Drilling and Production Practice* (1956), American Petroleum Institute, pp. 1–57.

Keeney, R. L., and Raiffa, H. (1976). *Decisions with Multiple Objectives: Preferences and Value Tradeoffs*. Wiley, New York.

Kobayashi, K. and Takazawa, K. (2023). Randomized strategies for robust combinatorial optimization with approximate separation. *Algorithmica*. <https://doi.org/10.1007/s00453-023-01175-3>.

Koski, J. (1985). Defectiveness of weighting method for multicriterion optimization.

Krause, A., Singla, A. and Golovin, D. (2011). Submodular maximization under noise.

Kuhn, H.W. (1953). Extensive games and the problem of information. In: *Contributions to the Theory of Games (AM-28), Volume II*, 193–216. Princeton University Press, Princeton, NJ.

Lade, S. J., Steffen, W., de Vries, W., Carpenter, S. R., Donges, J. F., Gerten, D., Hoff, H., Newbold, T., Richardson, K., and Rockström, J. (2020). Human impacts on planetary boundaries amplified by Earth system interactions. *Nature Sustainability*, 3(2), 119–128.

Lin, D., Hanscom, L., Murthy, A., Galli, A., Evans, M., Neill, E., Mancini, M.S., Martindill, J.,

- Medouar, F.-Z., Huang, S., Wackernagel, M., 2018. Ecological footprint accounting for countries: Updates and results of the National Footprint Accounts, 2012–2018.
- Lygeros, J., Tomlin, C., and Sastry, S. (1999). Controllers for reachability specifications for hybrid systems. *Automatica*, 35(3), 349–370.
- Main, J.C.A. and Randour, M. (2024). Different strokes in randomised strategies: Revisiting Kuhn’s theorem under finite-memory assumptions. *Information and Computation*, 301, 105229.
- Martinet, V. (2011). A characterization of sustainability with indicators. *Journal of Environmental Economics and Management*, 61(2), 183–197.
- Martinez-Alier, J., Munda, G., and O’Neill, J. (1998). Weak comparability of values as a foundation for ecological economics. *Ecological Economics*, 26(3), 277–286.
- Meadows, D.H., Meadows, D.L., Randers, J., Behrens III, W.W., 1972. The Limits to Growth.
- Messac, A. and Mattson, C.A. (2002). Generating well-distributed sets of Pareto points for engineering design using physical programming. *Optimization and Engineering* 3, 431–450.
- Miettinen, K. (1999). *Nonlinear Multiobjective Optimization*. Kluwer.
- Munda, G. (2005). Multiple criteria decision analysis and sustainable development. In: Figueira, J., Greco, S., and Ehrgott, M. (eds.), *Multiple Criteria Decision Analysis: State of the Art Surveys*. Springer, New York, 953–986.
- Munda, G., Nardo, M., 2009. Noncompensatory/nonlinear composite indicators for ranking countries: a defensible setting.
- Nardo, M., Saisana, M., Saltelli, A., Tarantola, S., Hoffmann, A., and Giovannini, E. (2008). *Handbook on Constructing Composite Indicators: Methodology and User Guide*. OECD Publishing, Paris.
- Neumayer, E. (2003). *Weak versus Strong Sustainability: Exploring the Limits of Two Opposing Paradigms*. 2nd ed. Edward Elgar.
- Neumayer, E. (2013). *Weak versus Strong Sustainability: Exploring the Limits of Two Opposing Paradigms*, 4th ed. Edward Elgar, Cheltenham.
- Augeraud-Véron, E. and Pereau, J.-C., 2022. Sustainable groundwater management in a two-cell aquifer model. *Environmental Modeling and Assessment* 27, 1–19.
- O’Neill, D. W., Fanning, A. L., Lamb, W. F., and Steinberger, J. K. (2018). A good life for all within planetary boundaries. *Nature Sustainability*, 1(2), 88–95.
- Pearce, D. and Atkinson, G. (1995). Measuring sustainable development.
- Jacquez, J.A., Simon, C.P., 1993. Qualitative theory of compartmental systems.
- Brunner, P.H., Rechberger, H., 2004. Practical Handbook of Material Flow Analysis.
- Prajna, S., Jadbabaie, A., Pappas, G.J., 2007. A framework for worst-case and stochastic safety verification using barrier certificates. *IEEE Transactions on Automatic Control* 52, 1415–1428.
- Prajna, S., Jadbabaie, A., 2004. Safety verification of hybrid systems using barrier certificates. In: *Hybrid Systems: Computation and Control VII. Lecture Notes in Computer Science* 2993, 477–492.
- Raworth, K. (2012). A safe and just space for humanity: Can we live within the doughnut? *Oxfam Policy and Practice: Climate Change and Resilience*, 8(1), 1–26.
- Redner, S., 2001. A Guide to First-Passage Processes.
- Ricard, D., Minto, C., Jensen, O.P., Baum, J.K., 2012. Examining the knowledge base and status of commercially exploited marine species with the RAM Legacy Stock Assessment Database.
- Rockström, J., Steffen, W., Noone, K., et al. (2009). A safe operating space for humanity. *Nature*, 461, 472–475.

- Roy, B. (1996). *Multicriteria Methodology for Decision Aiding*. Kluwer Academic Publishers, Dordrecht.
- Saint-Pierre, P. (1994). Approximation of the viability kernel. *Applied Mathematics and Optimization*, 29, 187–209.
- Schaefer, M. B. (1954). Some aspects of the dynamics of populations important to the management of commercial marine fisheries. *Inter-American Tropical Tuna Commission Bulletin*, 1(2), 23–56.
- Schär, S., Pohl, E., and Geldermann, J. (2025). Analysing the compensatory properties of the outranking approach PROMETHEE. *Journal of Multi-Criteria Decision Analysis*, 32, e70013.
- Sessa, P.G., Bogunovic, I., Kamgarpour, M. and Krause, A. (2020). Mixed strategies for robust optimization of unknown objectives. *AISTATS 2020*, PMLR **108**.
- Sinha, A., Fang, F., An, B. and Kiekintveld, C. (2018). Stackelberg security games: looking beyond a decade of success.
- Sion, M. (1958). On general minimax theorems. *Pacific Journal of Mathematics*, 8(1), 171–176.
- Smith, H.L., 1995. Monotone Dynamical Systems: An Introduction to the Theory of Competitive and Cooperative Systems. Mathematical Surveys and Monographs 41.
- Solow, R. M. (1974). Intergenerational equity and exhaustible resources. *Review of Economic Studies*, 41, 29–45.
- DFO. (2016). *Stock Assessment of Northern Cod (NAFO Divs. 2J3KL) in 2016*. Canadian Science Advisory Secretariat, Science Advisory Report 2016/026. Fisheries and Oceans Canada, Ottawa.
- Stadler, W. and Dauer, J.P. (1992). Multicriteria optimization in engineering: a tutorial and survey.
- Stadler, W. (1995). Fundamentals of multicriteria optimization.
- Sterner, R.W. and Elser, J.J., 2002. *Ecological Stoichiometry: The Biology of Elements from Molecules to the Biosphere*.
- Tapley, B.D., Bettadpur, S., Ries, J.C., Thompson, P.F., Watkins, M.M., 2004. GRACE measurements of mass variability in the Earth system.
- Tilton, J.E., Lagos, G., 2007. Assessing the long-run availability of copper.
- Tilton, J.E., 2003. On Borrowed Time? Assessing the Threat of Mineral Depletion.
- Usubiaga-Liaño, A. (2025). Strong sustainability in the SEEA and the wider indicator debate. *One Ecosystem*, 10, e141086.
- Vincke, P. (1992). *Multicriteria Decision-Aid*.
- Vorobeychik, Y. and Li, B. (2014). Optimal randomized classification in adversarial settings.
- Abaee, A. (2026). *An Obstruction Calculus for Viability under Incomplete Observation*. Prior public preprint version (v12), Zenodo, <https://doi.org/10.5281/zenodo.22552616>.
- Wackernagel, M., Beyers, B., 2019. Ecological Footprint: Managing our Biocapacity Budget.
- Warga, J. (1972). *Optimal Control of Differential and Functional Equations*. Academic Press, New York.
- von Neumann, J. (1928). Zur Theorie der Gesellschaftsspiele. *Mathematische Annalen*, 100, 295–320.
- Wei, N. and Zhang, P. (2024). Adjustability in robust linear optimization. *Mathematical Programming*, 208(1–2), 581–628.

Supplementary Material

Framework extensions (governance constructors, the implementability ladder, the commons obstruction, intergenerational structures, the nested-impossibility theorem, composition interfaces), the planetary-boundaries application note, the full framework definitions, and the conjectures are given in the accompanying file `paper06_assessment_separation_v68_supplementary.md`, which also enumerates the machine artifact’s 25 checks (S8) and the itemized data requirements for empirical application (S9), and the per-literature positioning notes (S10).

Declarations

Ethics approval and consent to participate

Not applicable.

Consent for publication

Not applicable.

Availability of data and materials

All datasets and code generated and analysed during the current study are available in the figshare repository. The deposit serves two explicit roles: (i) this manuscript’s verification artifact — the exact-integer grid verifier for the finite rational instance (all 25 checks), the figure-generation code, the manuscript figure files, and the preserved novelty-search strings; and (ii) the software companion’s archive — the SafeTransition library (with its own separate 24-check benchmark suite), whose master deposit this shared item also is. The two check lists are distinct and are not pooled. <https://doi.org/10.6084/m9.figshare.33764023>

Competing interests

The authors declare that they have no competing interests.

Funding

Not applicable.

Authors’ contributions

A.A. conceptualized the entire work, wrote, reviewed and edited the manuscript.

Acknowledgements

Not applicable.

AI declaration

GLM (Z.ai), Qwen (Alibaba Cloud) and DeepSeek AI assisted with drafting and iterative review.