

Discrete Regimes, Carefully: Mode Uncertainty in the Obstruction Calculus

Edition 2, revised after a joint audit

Amin Abaee

Independent Researcher

ORCID: 0000-0002-0019-1842

amin_abaee@ut.ac.ir

September 24, 2026

Abstract

Discrete regimes are where hybrid formalisms tempt and where decorative extensions hide. This paper, the D11 item of the generalization catalogue, takes the catalogue’s own warning seriously — hybrid structure only where the mode is load-bearing — and exhibits the smallest setting in which it is. The working system gains an exogenous regime $m \in \{1, 2\}$ chosen adversarially each step: each patch regenerates, and the pressured patch takes damage unless protected. Under either *frozen* regime every safe pair is viable (36 of 36; each regime is separately sustainable); with the regime *unread*, the mode-blind kernel holds exactly 28 of 36 — eight beliefs, precisely those containing the thin-stock corner $(1, 1)$, are lost purely to unread regimes, every loss generated by the mode-blind death of the corner state itself (no merged-pair loss exists on this grid); with the regime *read* per (state, mode) cell, protecting the pressured patch restores all 36. The chain $\text{Viab}_{\text{blind}} \subset \text{Viab}_{\text{observed}} = \text{Viab}_{\text{frozen-1}} \cap \text{Viab}_{\text{frozen-2}}$ is exact on the audited pair family (in general only the inclusion holds; equality is the audited instance’s computed fact), and the recursive augmented-belief computation is cross-validated against brute-force game enumeration. Scope is recorded, not claimed: finite modes, total transitions, no Zeno, no continuous-mode metrics (6/6 check families).

Introduction and relation to the catalogue’s warning

The catalogue proposes hybrid generalization “carefully”: only if the mode structure is central, otherwise decorative. The paper honors this by construction: the mode is load-bearing exactly when it is adversarial and unread — the eight lost beliefs of the verified record below vanish under either frozen regime and under full reading, so the obstruction belongs to mode *uncertainty* alone. Positioning against the programme: the calculus’s Example 1 is the constant unread hidden parameter (a common-action obstruction) and its Section 8 the delayed hidden regime (Abaee, 2026, An obstruction calculus); the present note’s increment is per-step adversarial switching of the mode,

and the catalogue slot is taken accordingly — unread discrete modes, one step short of the timed hybrid layer the catalogue reserves. The paper’s model class is finite discrete modes with total transitions; the catalogue’s warning against decorative hybrid machinery (Zeno, timed-automata metrics, continuous mode dynamics) is adopted verbatim as the scope statement.

Mode uncertainty

Branches carry (x, m) ; the per-cell read-then-act semantics applies on augmented cells; after the regulator’s per-cell move the adversary re-picks the mode. The timing is causal by construction: the mode realized at step t is the one the adversary chose at the end of step $t - 1$; the observed class conditions on that realized mode, and the re-pick after the move sets the next mode. Dynamics (the audited two-patch data, Abaee, 2026, An obstruction calculus, Section 8): both patches regenerate +1 per step (cap 3); the regime names the pressured patch, which takes damage 2 unless protected ($u = m$); the unpressured patch is never damaged; instruments $u \in \{1, 2\}$; full transition data in the script. Three policy classes: frozen mode (no switching), mode-blind (cells keyed by state), mode-observed (cells keyed by (x, m)).

Proposition 1 (Frozen regimes are separately sustainable). With the mode frozen at either value, protecting the pressured patch lets both patches grow: every safe pair is viable under each frozen regime (36 of 36).

Proposition 2 (The mode-uncertainty obstruction and its exact boundary). Under an unread adversarial regime the mode-blind kernel holds exactly 28 of 36 safe pairs; the eight lost beliefs are exactly those containing $(1, 1)$; under mode observation, protecting the pressured patch restores all 36. Hence $\text{Viab}_{\text{blind}} \subset \text{Viab}_{\text{observed}} = \text{Viab}_{\text{frozen-1}} \cap \text{Viab}_{\text{frozen-2}}$ setwise — on the audited pair family; in general only the inclusion $\text{Viab}_{\text{observed}} \subseteq \text{Viab}_{\text{frozen-1}} \cap \text{Viab}_{\text{frozen-2}}$ holds (switching can deplete margins that each frozen regime preserves), and the equality is the audited instance’s computed fact.

Proof. The counts, the lost-family characterization (all containing $(1, 1)$), and the setwise chain are computed exactly by the verification script; the restoration strategy is

$u = m$ per observed cell. The boundary reading: thin-buffer beliefs cannot juggle an unread adversary, buffered ones can. The pair-level loss is generated entirely at the singleton level: the corner $(1, 1)$ is the only mode-blind-dead state (every other state admits a state-keyed law serving both regimes), every lost pair contains it, and no pair of mode-blind-viable states is merged into nonviability (all verified, E7). The mode-blind count 28 thus coincides, as a set on this grid, with the base system’s full-information kernel — the pairs avoiding the corner — while arising under a different damage routing; the coincidence is recorded, not claimed as an identity across models (E8). $\square$

Verified record (8/8 exact check families)

E1 frozen-mode kernels $36 = 36$ (the corner $(1, 1)$ viable under each frozen regime); E2 the mode-blind kernel exactly 28 of 36 with the eight lost beliefs enumerated (all contain $(1, 1)$); E3 mode-observed restoration 36 of 36; E4 cross-validation of the recursive augmented computation against brute-force game-tree enumeration (three probe beliefs $\times$ horizons 1–4 $\times$ frozen/adversarial switching); E5 the exact policy-class chain ($28 \subset 36 = 36$); E6 the scope assertions (finite modes, total transitions, load-bearing instance exhibited); E7 singleton generation: exactly $\{(1, 1)\}$ is mode-blind-dead among the nine singletons, all eight others viable as singletons, and every pair of mode-blind-viable states is mode-blind viable (no merging loss); E8 the setwise coincidence of the mode-blind kernel with the base system’s full-information kernel (both the 28 pairs avoiding $(1, 1)$) recorded against the different damage routings.

What is not claimed

No general hybrid-systems theory: no Zeno analysis, no timed or metric mode semantics, no stochastic mode processes, no continuous-mode dynamics. The regime is adversarial and exogenous; endogenous regime change (institutional responses to stock levels) is cited as the governance extension and not developed. The damage–regeneration asymmetry is the audited instance’s; no classification over damage profiles is claimed. The paper’s honesty obligation is the catalogue’s warning itself: where the mode is not load-bearing (frozen or observed), the hybrid layer disappears — and the paper says so.

Declarations

Funding. None. **Competing interests.** None. **Data availability.** No external data were used. **Code availability.** The verification script `hybrid_mode_viability_v2_verify.py` is publicly available at <https://github.com/MIKEAA2020/general-sustainability> (folder `arena agent 1/paper rewrites/latex`); it reproduces all eight check families verbatim. **AI declaration.** GLM (Z.ai), Qwen (Alibaba Cloud) and DeepSeek AI assisted with drafting and iterative review of the first edition; the joint audit

of the first edition and this revision were prepared with further AI assistance, under the programme’s verification discipline.

References

Abaee, A.: An obstruction calculus for viability under incomplete observation. Submitted. <https://github.com/MIKEAA2020/general-sustainability> (2026)