

# Computational Viability Certification under Incomplete Observation: A Continuous-to-Finite Bridge with Exact Certificates, a Solver-Certified Campaign, and a Reference Library

Amin Abaee

Independent Researcher

ORCID: 0000-0002-0019-1842

amin\_abaee@ut.ac.ir

September 24, 2026

## Abstract

Nonviability certificates are useful computationally only if they are finite, attainable by standard solvers, and independently checkable. This paper develops such a pipeline for linear systems under exogenous partial observation of a hidden mode. An outer moment approximation of all measurable information-adapted controls, combined with an inner adversarial approximation, yields finite linear programs whose optimal values sandwich the continuous-time safety value between explicit error estimates ( $\rho_{\mathcal{H},G} \leq J_{\mathcal{I}}(T) \leq \rho_{\mathcal{H},G} + R_U V_U h_u + \delta_\beta + L h_t$ ); every dual solution is a finite obstruction certificate covering every measurable policy; nonviability is detected under every sufficiently fine refinement, and positive certificates never exceed  $r + 1$  safety labels at rank  $r$  — the relevant dimension being information–time rank, not instantaneous input dimension, and a matching construction with scalar input produces obstructions requiring arbitrarily many witnesses. The certificates account for post-observation recourse: on the three-branch delayed-observation instance every pair of compatible states is viable and a common control keeps every branch safe throughout the blind window, yet nonviability is certified, with the exact delay threshold  $\tau^* = 7/50$ . The finite program is solved by a standard solver (HiGHS) across a five-point mesh refinement; optimality is certified in exact rational arithmetic by independent primal and dual witnesses, the solver agreeing with the exact values to at most  $4.2 \times 10^{-17}$ , and the refinement obeys the closed law  $\rho(h) = \max\{0, 3/50 - Th/4\}$ . A reference library, **viacert**, packages the exact finite-horizon recursion and the fibre-partition criteria behind a pluggable finite-system interface, with twelve audited identities in its self-test. All exact computations run in integer and rational arithmetic.

**Keywords:** viability; incomplete observation; obstruction certificates; linear programming; semi-infinite programming; exact arithmetic; certified computation

## 1. Introduction

Viability theory under full observation has a mature numerical apparatus: kernel approximations, set-valued meshes, and convergence results (Aubin, 1991; Cardaliaguet, Quincampoix, and Saint-Pierre, 1999). Un-

der incomplete observation the sufficiency side is carried by output-feedback conditions and the estimation-tube programme (Veliov, 1993), while the necessity side — certificates whose firing *proves* nonviability of an information state (Abaee, 2026b) — has lacked a computational bridge. The difficulty is the quantifier order: the adversary chooses the mode, the initial state inside its cell, and a measurable disturbance with knowledge of the announced policy, so naive discretization of controls proves at most the infeasibility of a held-control subclass, not of the policy class itself.

This paper builds the bridge and the tooling around it. The main result (Theorem 4) approximates the admissible controls from outside by their block moments — a necessary relaxation of continuous-time safety that restricts nothing — and the adversary from inside by explicitly stored scenarios, yielding a finite linear program whose optimal value sandwiches the exact minimax safety violation. Dual solutions of the program are finite, policy-independent obstruction certificates; positive certificates survive mesh refinement; and their size is governed by the rank of the safety-row matrix in the information–time product, which no bound on the instantaneous input dimension controls (Proposition 5).

The pipeline is demonstrated end to end on a three-branch delayed-observation instance: delay threshold, Farkas certificate, and a five-point mesh campaign in which the HiGHS solver (Huangfu and Hall, 2018) attains the exact rational optimum at every mesh, with independent primal–dual witnesses closing the duality gap exactly. The finite-side machinery — the exact backward recursion and the fibre-partition criteria — is packaged as the reference library **viacert** with a pluggable finite-system interface and a self-test of twelve audited identities.

The contributions are: the continuous-to-finite sandwich theorem with sound, refinement-complete, sparsity-bounded certificates (Section ); the recourse-aware obstruction reading and the exact three-branch instance with its certified campaign (Section ); the information–time rank separation (Section ); the refinement-erosion identity and the oracle-recourse extension (Section ); the complexity and residual-bounded verification framework (Section ); and the reference library (Section ).

## Model, information, and quantifiers

*Dynamics.* Let the hidden mode be  $\theta \in \Theta := \{1, \dots, M\}$ . Conditional on  $\theta$ , the continuous state satisfies

$$\dot{x}(t) = A_\theta(t)x(t) + g_\theta(t) + B_\theta(t)u(t) + E_\theta(t)d(t), \quad t \in [0, T],$$

under the standing assumptions:  $A_\theta, g_\theta, B_\theta, E_\theta$  bounded and piecewise continuous, with bounded-variation bounds on  $B_\theta, E_\theta$ ; the initial set in mode  $\theta$  a nonempty compact polytope  $P_\theta \subset \mathbb{R}^n$ ; the input and disturbance sets nonempty compact polytopes  $U = \{v : Fv \leq f\} \subset \mathbb{R}^m$  and  $D_\theta \subset \mathbb{R}^p$ ; safety the closed polyhedron  $\mathcal{V} = \{x : c_i^\top x \leq b_i, i = 1, \dots, s\}$ ; and the initial information state  $B_0 = \bigcup_\theta \{\theta\} \times P_\theta$  initially safe (a time-zero certificate suffices otherwise). The dynamics admit unique Carathéodory solutions for every admissible measurable signal pair.

*Information.* At known times  $0 = \tau_0 < \tau_1 < \dots < \tau_R = T$  the controller receives exogenous information about  $\theta$ : refining history partitions  $\mathcal{P}_0 \preceq \mathcal{P}_1 \preceq \dots \preceq \mathcal{P}_{R-1}$  of  $\Theta$ , the observation history during  $I_r = [\tau_r, \tau_{r+1})$  identifying the cell  $C_r(\theta) \in \mathcal{P}_r$ . These are history partitions, not necessarily partitions generated by the latest observation alone. Observations do not depend on  $x, u, d$ ; active sensing and endogenous observation branching are excluded from the exact theorem (Section gives a sound extension).

*Policies and quantifiers.* Policies are deterministic, measurable, nonanticipative, and adapted to the observation history; every such policy is represented exactly by signals  $u_{r,C} \in L^\infty(I_r; U)$ ,  $C \in \mathcal{P}_r$ . The robust quantifier order is

$$\exists \pi \in \Pi_{\mathcal{I}} \forall \theta \forall x_0 \in P_\theta \forall d \in L^\infty([0, T]; D_\theta) \forall t \in [0, T],$$

the adversary selecting mode, initial state, and measurable disturbance with knowledge of the announced policy; randomized policies are not included. The information-constrained viability kernel is  $\mathcal{K}_{\mathcal{I}}^T = \{B_0 : \exists \pi \forall \theta, x_0, d, t \ x_{\theta, x_0, d}^\pi(t) \in \mathcal{V}\}$ . Writing  $\text{Lift}(K_{\text{full}}^T) = \{B : \emptyset \neq B \subseteq K_{\text{full}}^T\}$  for the belief lifting of the full-information kernel,  $\mathcal{K}_{\mathcal{I}}^T \subseteq \text{Lift}(K_{\text{full}}^T)$ , and the inclusion is strict in general: the instance of Section has every two-state subbelief viable while the three-state belief is not.

## Adjoint safety rows and the safety value

Let  $X_\theta(t, s)$  denote the state-transition matrix. For a label  $a = (\theta, i, t) \in \mathcal{S} := \Theta \times \{1, \dots, s\} \times [0, T]$  define the input kernel and margin

$$k_a(\sigma) = \mathbf{1}_{\{\sigma \leq t\}} B_\theta(\sigma)^\top X_\theta(t, \sigma)^\top c_i,$$

$$\begin{aligned} \beta_a &= b_i - h_{P_\theta}(X_\theta(t, 0)^\top c_i) - \int_0^t c_i^\top X_\theta(t, \sigma) g_\theta(\sigma) d\sigma \\ &\quad - \int_0^t h_{D_\theta}(E_\theta(\sigma)^\top X_\theta(t, \sigma)^\top c_i) d\sigma, \end{aligned}$$

where  $h_Q(v) = \sup_{q \in Q} v^\top q$ . For a policy  $\pi$  set  $\mathcal{F}_a(\pi) = \int_0^T k_a(\sigma)^\top u_\theta^\pi(\sigma) d\sigma - \beta_a$ .

**Proposition 1** (robust rows).  *$\pi$  is robustly safe on  $[0, T]$  if and only if  $\mathcal{F}_a(\pi) \leq 0$  for every  $a \in \mathcal{S}$ . This follows from the linear solution formula; initial-state and disturbance maximizations separate from the control term*

*because the control signal in a mode does not depend on  $x_0$  or  $d$ . For polytopic  $D_\theta$ , a maximizing disturbance can be chosen measurably by selecting a maximizing vertex at each time.*

Define the minimax safety violation  $J_{\mathcal{I}}(T) = \min_{\pi \in \Pi_{\mathcal{I}}} \max_{a \in \mathcal{S}} \mathcal{F}_a(\pi)$ .

**Proposition 2** (value and nonviability). *The minimum exists: the policy-signal space is weak-\* compact and each  $\mathcal{F}_a$  is weak-\* continuous. Consequently  $B_0 \notin \mathcal{K}_{\mathcal{I}}^T$  if and only if  $J_{\mathcal{I}}(T) > 0$ . The strict positive margin is a consequence of compactness and closed safety constraints, not an extra nonviability assumption.*

## The finite obstruction certificate

Define  $\phi_U(g) := \min_{v \in U} g^\top v = -h_U(-g)$ . Choose finitely many labels  $a_\ell = (\theta_\ell, i_\ell, t_\ell)$ ,  $\ell = 1, \dots, q$ , and weights  $\lambda_\ell \geq 0$ ,  $\sum_\ell \lambda_\ell = 1$ . For each active information cell  $C \in \mathcal{P}_r$ , pool the adjoint kernels  $g_{r,C}(\sigma) = \sum_{\ell: \theta_\ell \in C} \lambda_\ell k_{a_\ell}(\sigma)$ . The certificate margin is

$$\Gamma_{\mathcal{I}}(\lambda, a) = \sum_{r=0}^{R-1} \int_{I_r} \sum_{C \in \mathcal{P}_r} \phi_U(g_{r,C}(\sigma)) d\sigma - \sum_{\ell=1}^q \lambda_\ell \beta_{a_\ell}. \quad (1)$$

*Remark 3* (semantics). If  $\Gamma_{\mathcal{I}}(\lambda, a) > 0$ , then no admissible information-adapted policy is safe: for every policy,  $\sum_\ell \lambda_\ell \mathcal{F}_{a_\ell}(\pi) \geq \Gamma_{\mathcal{I}}(\lambda, a)$ , so at least one stored label satisfies  $\mathcal{F}_{a_\ell}(\pi) \geq \Gamma_{\mathcal{I}}(\lambda, a) > 0$  — an admissible realization violates facet  $i_\ell$  at time  $t_\ell$ . The violating label may depend on the policy; the finite list and its weights do not. The certificate is not an adversarial tree: it is a finite, policy-independent set of adversarial tests combined by an adjoint identity.

## The continuous-to-finite theorem

The exact dual expression (1) alone would be close to standard infinite-dimensional convex duality (Anderson and Nash, 1987; Hettich and Kortanek, 1993). The consequential result is a sound finite construction covering unrestricted measurable controls.

*Outer control moments.* Choose a control-time mesh  $\mathcal{H} = \{J_1, \dots, J_H\}$  containing every sensing time. Introduce one vector  $v_{j,C} \in U$  for every cell  $C$  active during  $J_j$ , and let  $Q = \sum_j |\mathcal{P}_{r(j)}|$  be the number of information-time blocks. For a label  $a$  define the finite row  $A_{a,j,C} = \mathbf{1}_{\{\theta \in C\}} \int_{J_j} k_a(\sigma) d\sigma$ , the block average  $\bar{k}_{a,\mathcal{H}}(\sigma) = \frac{1}{|J_j|} \int_{J_j} k_a(s) ds$  on  $J_j$ , and the moment error

$$e_{a,\mathcal{H}} = \int_0^T h_U(\bar{k}_{a,\mathcal{H}}(\sigma) - k_a(\sigma)) d\sigma, \quad (2)$$

for which any verified upper bound  $\bar{e}_a \geq e_{a,\mathcal{H}}$  may be used. For every measurable policy, its block averages belong to  $U$ , and

$$A_a v - \int_0^T k_a(\sigma)^\top u_\theta^\pi(\sigma) d\sigma \leq \bar{e}_a. \quad (3)$$

Inequality (3) is the crucial orientation: the finite moments yield a necessary relaxation of continuous-time safety. They do not restrict the real controller to held inputs. For the sampled-data subclass whose policies are

held constant on the blocks, the same polyhedron is an exact reduction rather than a relaxation.

*Inner adversarial approximation.* One may use a right-hand side  $\beta_a^+ \geq \beta_a$  obtained from an explicit admissible initial state and disturbance signal rather than exact worst-case optimization. If  $0 \leq \beta_a^+ - \beta_a \leq \delta_\beta$  uniformly over the labels used, the approximation is safe for nonviability: it weakens the adversary and enlarges the feasible controller set. For example, with  $w_a(\sigma) = \mathbf{1}_{\{\sigma \leq t\}} E_\theta(\sigma)^\top X_\theta(t, \sigma)^\top c_i$ , on a disturbance mesh choose  $d_{a,J} \in \arg \max_{d \in D_\theta} (\int_J w_a)^\top d$ , yielding an explicit piecewise-constant admissible disturbance; if  $D_\theta$  has radius  $R_D$  about a fixed center and  $\text{TV}(w_a) \leq V_D$ , then

$$0 \leq \beta_a^+ - \beta_a \leq R_D V_D h_d, \quad (4)$$

apart from any separately accounted initial-state or numerical error.

*The finite program.* Let  $G \subset [0, T]$  be a finite time grid including every mode and safety facet at every grid time. Solve

$$\begin{aligned} \rho_{\mathcal{H},G} = \min_{v,\rho} \quad & \rho \quad \text{s.t.} \quad A_a v - \beta_a^+ - \bar{e}_a \leq \rho \\ & (a \in \Theta \times \{1, \dots, s\} \times G), \quad v_{j,C} \in U. \end{aligned} \quad (5)$$

**Theorem 4** (continuous-to-finite). *Assume: uniformly over admissible policies, modes, and facets,  $|\mathcal{F}_{\theta,i,t}(\pi) - \mathcal{F}_{\theta,i,t'}(\pi)| \leq L|t - t'|$  (6); the input kernels satisfy  $\text{TV}(k_a) \leq V_U$  (7);  $U$  has radius  $R_U$  about some fixed center; and the time grid has covering radius at most  $h_t$ , the control mesh maximum length  $h_u$ , and the adversarial approximation satisfies the bound above. Then:*

(i) *Certified value sandwich.*

$$\rho_{\mathcal{H},G} \leq J_{\mathcal{I}}(T) \leq \rho_{\mathcal{H},G} + R_U V_U h_u + \delta_\beta + L h_t. \quad (8)$$

With (4),  $J_{\mathcal{I}}(T) - \rho_{\mathcal{H},G} \leq R_U V_U h_u + R_D V_D h_d + L h_t$ . Verified numerical row widths are added to the right side.

(ii) *Sound finite certificates.* If  $\rho_{\mathcal{H},G} > 0$ , the program's dual supplies a finite certificate of nonviability of the original continuous-time system, covering every measurable information-adapted policy.

(iii) *Completeness under refinement.* If  $B_0 \notin \mathcal{K}_{\mathcal{I}}^T$ , then every sufficiently fine sequence of the above meshes, with numerical errors tending to zero, eventually produces a positive finite certificate. Nonviability is thus semidecidable in this class when the coefficients and error bounds are effectively computable.

(iv) *Exact atomic dual representation.*

$$J_{\mathcal{I}}(T) = \sup_{q < \infty, a_1, \dots, a_q, \lambda \in \Delta_q} \Gamma_{\mathcal{I}}(\lambda, a). \quad (10)$$

In particular, nonviability is equivalent to the existence of a positive finite atomic certificate.

(v) *Sparse witnesses.* Let  $r = \text{rank } A$  for the finite safety-row matrix in (5). Whenever  $\rho_{\mathcal{H},G} > 0$ , a positive certificate exists using at most  $r + 1 \leq mQ + 1$  safety labels, and it can retain any prescribed margin strictly below  $\rho_{\mathcal{H},G}$ . If the lifted continuous-time kernels span a finite-dimensional space of dimension  $r_{\mathcal{I}}$ , a continuous-time obstruction has at most  $r_{\mathcal{I}} + 1$  labels. The relevant dimension is information–time rank, not  $m$ .

The complete proof is in Supplementary S1. The asymmetry at its core is worth stating plainly: averaging arbitrary controls proves the lower bound; implementing piecewise-constant controls proves the upper bound. Confusing these two directions is what makes naïve infeasibility-by-discretization arguments unsound.

## A three-branch delayed-observation instance carried exactly

*Plant and exact delay threshold.* Consider the planar double integrator  $\dot{p} = v$ ,  $\dot{v} = u$ , with unit vectors  $n_1 = (1, 0)$ ,  $n_2 = (-\frac{3}{5}, \frac{4}{5})$ ,  $n_3 = (-\frac{3}{5}, -\frac{4}{5})$ , the hexagonal input set  $U = \text{conv}\{\pm n_1, \pm n_2, \pm n_3\}$ , equivalently  $\{v : |u_2| \leq \frac{4}{5}, |2u_1 + u_2| \leq 2, |2u_1 - u_2| \leq 2\}$ , safety  $n_i^\top p \leq 2$  ( $i = 1, 2, 3$ ) and  $|v_1|, |v_2| \leq \frac{6}{5}$ , three possible initial states  $p_j(0) = \frac{34}{25} n_j$ ,  $v_j(0) = n_j$ , no disturbance, and the mode unobserved until time  $\tau$ , when it is revealed exactly.

If  $j$  is known,  $u = -n_j$  for one second and then  $u = 0$  keeps the maximum critical position at  $\frac{34}{25} + \frac{1}{2} = \frac{93}{50} = 1.86 < 2$ : all three physical states are indefinitely viable under full information. Take three witness labels, each using its own position facet at  $t_* = \tau + 1$ , with kernels  $k_j(s) = \mathbf{1}_{\{s \leq t_*\}}(t_* - s)n_j$  and  $\beta_j = \frac{16}{25} - t_*$ , weighted by  $\lambda_1 = \frac{3}{8}$ ,  $\lambda_2 = \lambda_3 = \frac{5}{16}$ , so that  $\sum_j \lambda_j n_j = 0$ . Before observation the pooled kernel is identically zero; after observation each mode has its own control and  $\phi_U(\lambda_j(t_* - s)n_j) = -\lambda_j(t_* - s)$ . Hence

$$\Gamma(\tau) = - \int_{\tau}^{\tau+1} (\tau+1-s) ds - \left( \frac{16}{25} - \tau - 1 \right) = \tau - \frac{7}{50}, \quad (16)$$

so  $\tau > \frac{7}{50}$  implies  $B_0 \notin \mathcal{K}_{\mathcal{I}}$ . Conversely,  $u = 0$  before observation and  $u = -n_j$  for one second afterward keeps the maximum critical position at  $\frac{34}{25} + \tau + \frac{1}{2}$ , safe exactly when  $\tau \leq 0.14$ . Thus  $\tau_{\max} = \frac{7}{50} = 0.14$  is exact, and for horizons containing the stopping time  $J_{\mathcal{I}}(\tau+1) = \tau - 0.14$ .

*Pairs are viable; the blind-window exit test misses.* At  $\tau = \frac{1}{5}$ , the common control  $u = 0$  keeps all three branches safe throughout the blind window:  $n_j^\top p_j(0.2) = \frac{39}{25} = 1.56 < 2$ . No exit occurs before information arrives; nevertheless  $\Gamma(0.2) = 0.06 > 0$  — the best post-observation recourse cannot recover all branches. The obstruction is not that information arrives after unavoidable physical exit; it is that information arrives after some indistinguishable branch has necessarily accumulated more braking liability than the remaining control authority can discharge.

*The finite program and its exact witnesses.* Set  $\tau = \frac{1}{5}$ ,  $T = \frac{6}{5}$ ,  $h = \frac{1}{10}$ : two common pre-observation control blocks and ten post-observation blocks for each of three modes, so  $Q = 32$ ,  $mQ = 64$ . For the endpoint position rows the exact moment error is  $e = 12 \cdot h^2/4 = \frac{3}{100}$  (the kernel is affine on each of the  $T/h = 12$  mesh intervals and  $h_U(\alpha n_j) = |\alpha|$ ); the endpoint right-hand side is  $\beta = \frac{16}{25} - \frac{6}{5} = -\frac{14}{25}$ , relaxed to  $\beta + e = -\frac{53}{100}$ . The exact continuous certificate has margin  $\frac{3}{50} = 0.06$ ; the finite moment certificate has margin  $\frac{3}{100} = 0.03$ .

The dual witness is exact. Order the input facets by  $F = [(0, 1); (0, -1); (2, 1); (2, -1); (-2, 1); (-2, -1)]$ ,  $f = (\frac{4}{5}, \frac{4}{5}, 2, 2, 2, 2)$ , and define  $\eta_1 = (0, 0, 0, 0, \frac{1}{4}, \frac{1}{4})$ ,  $\eta_2 = (0, \frac{1}{2}, 0, \frac{3}{10}, 0, 0)$ ,  $\eta_3 = (\frac{1}{2}, 0, \frac{3}{10}, 0, 0, 0)$ , satisfying  $F^\top \eta_j = -n_j$  and  $f^\top \eta_j = 1$ . For a post-observation interval  $J_k$  put

$a_k = \int_{J_k} (T-s) ds$  and  $\mu_{j,k} = \lambda_j a_k \eta_j$ . Before observation, stationarity holds because  $\sum_j \lambda_j n_j = 0$ ; after observation, because  $\lambda_j a_k n_j + F^\top \mu_{j,k} = 0$ . Since  $\sum_k a_k = \frac{1}{2}$ , the Farkas contradiction is  $\lambda^\top (\beta + e) + \mu^\top f_{\text{blk}} = -\frac{53}{100} + \frac{1}{2} = -\frac{3}{100} < 0$  — an exact rational certificate. The primal witness is equally exact: hold the pre-observation blocks at 0 and the post-observation blocks of mode  $j$  at  $-n_j$ ; the endpoint row then reads  $-\frac{1}{2} + \frac{14}{25} - \frac{3}{100} = \frac{3}{100}$ , and every other row lies at or below the same value.

*Solver-certified mesh study.* For each mesh the program (5) is built at its claimed dimensions and solved by HiGHS; optimality is then certified in exact rational arithmetic by the witness pair above (primal feasibility of every row; dual stationarity and dual objective  $-\rho(h)$ ), so the duality gap closes exactly and the solver’s returned value is a computational confirmation of a proved optimum, not an observation. Table 1 reports the campaign: exact values  $\rho(h) = \max\{0, \frac{3}{50} - Th/4\}$ , solver values, iteration counts, and wall-clock times as measured on the verification machine.

The values obey the exact law  $\rho(h) = \frac{3}{50} - Th/4$  for  $h < \frac{1}{5}$  and vanish at  $h = \frac{1}{5}$ : the whole refinement gap is the control-moment error  $Th/4$  of the endpoint rows. All dimension counts ( $mQ + 1$  variables;  $Ms(N_t+1) + p_U Q$  inequalities), all witness inequalities, and the closed-form law are machine-verified in exact rational arithmetic by the verification scripts of Section , which regenerate Table 1 verbatim and re-run the solver campaign.

## Computation, complexity, and independent verification

With  $Q$  information-time blocks and  $N_t + 1$  time labels, the program has  $mQ + 1$  variables and at most  $Ms(N_t+1) + p_U Q$  inequalities ( $p_U$  input-polytope facets); it is polynomial-time in these dimensions and the rational data size, conditional on the validated coefficient bounds. Initial-state dependence is affine, so no vertex enumeration of  $P_\theta$  is required: at each label solve  $\max_{\xi \in P_\theta} c_i^\top X_\theta(t, 0)\xi$  and store a maximizing or near-maximizing point — the verifier needs only membership of the stored point in  $P_\theta$ , not global optimality; near-optimality bears on the convergence guarantee, not on a positive certificate’s validity. The same holds for stored disturbance signals.

For rational outer rows  $\hat{A}v \leq \hat{b}$ ,  $F_{\text{blk}}v \leq f_{\text{blk}}$ , a residual-bounded certificate test suffices: with  $\lambda, \mu \geq 0$ ,  $r = \hat{A}^\top \lambda + F_{\text{blk}}^\top \mu$ ,  $c = \lambda^\top \hat{b} + \mu^\top f_{\text{blk}}$ , the condition

$$c + \sum_j h_U(-r_j) < 0 \quad (15)$$

is sufficient, since any feasible  $v$  satisfies  $-\sum_j h_U(-r_j) \leq r^\top v \leq c$ . All integration, rounding, and model-enclosure errors must be included before applying the test, and uncertain row data must be relaxed toward a larger feasible controller set: overestimating the worst disturbance contribution can create a false obstruction. A deletion-based sparsification uses at most one feasibility solve per safety row; exact rank compression requires verified rank, and floating-point singular-value truncation must be charged to the error budget. If  $J_{\mathcal{I}}(T) = \gamma > 0$ , a mesh with  $R_U V_U h_u + R_D V_D h_d + L h_t < \gamma$  guarantees detection — an explicit, though not claimed optimal, margin-dependent certificate-size bound.

## Information-time rank: input dimension does not bound the witnesses

**Proposition 5** (arbitrarily large minimal obstructions with one input). *For every integer  $q \geq 1$  there is a continuous-time system with scalar input  $u(t) \in [0, 1]$  and  $q + 1$  indistinguishable modes such that every proper sub-belief is viable, the full  $(q+1)$ -mode belief is not, and every obstruction certificate must involve all  $q + 1$  modes.*

*Proof.* Let  $T = q$ ,  $I_j = [j - 1, j]$ . For mode  $j \leq q$ :  $\dot{x}_j = \mathbf{1}_{I_j}(t)(u - 1)$ ,  $x_j(0) = 0$ ; for mode  $q + 1$ :  $\dot{x}_{q+1} = -u$ ,  $x_{q+1}(0) = q - \varepsilon$ ,  $0 < \varepsilon < 1$ . Require  $0 \leq x \leq q$ ; there are no observations. Writing  $w_j = \int_{I_j} u$ , safety requires  $-w_j \leq -1$  for  $j \leq q$  and  $\sum_j w_j \leq q - \varepsilon$ ; summing gives  $0 \leq -\varepsilon$ , a contradiction. If mode  $q + 1$  is removed,  $u \equiv 1$  is safe; if mode  $j \leq q$  is removed,  $u = 0$  on  $I_j$  and  $u = 1$  elsewhere is safe. Equal weights give certificate margin  $\varepsilon/(q + 1)$ ; the kernel rank is  $q$ , and  $q + 1$  witnesses are necessary.  $\square$

A statement such as “at most  $m + 1$  compatible states suffice to witness nonviability” is valid for convex common-action sets in  $\mathbb{R}^m$  but false for a common blind control function: the correct dimension counts independent temporal and informational control decisions. This distinction is consequential for review intervals, delayed sensing, and certificate-size claims; its discrete counterpart is a scoping caveat of the finite audits of Abaee (2026a).

## Refinement erosion and oracle recourse

Suppose a coarse information cell  $C$  is split into finer cells  $C_1, \dots, C_k$ . Because  $\phi_U(g_1 + \dots + g_k) \geq \sum_j \phi_U(g_j)$ , the same witness satisfies  $\Gamma_{\text{coarse}} \geq \Gamma_{\text{fine}}$ , with the exact difference

$$\int \left[ \phi_U \left( \sum_j g_{C_j} \right) - \sum_j \phi_U(g_{C_j}) \right] dt. \quad (13)$$

For  $U = [-\bar{u}, \bar{u}]^m$  this is  $\bar{u} \int [\sum_j \|g_{C_j}(t)\|_1 - \|\sum_j g_{C_j}(t)\|_1] dt$ : the certificate identifies the precise time intervals on which indistinguishable branches demand conflicting control effort. On rational instances the identity is exactly checkable; the verification record confirms  $\phi_U(g_1 + g_2) \geq \phi_U(g_1) + \phi_U(g_2)$  on 3,240 deterministic rational pairs for the hexagonal  $U$  of Section .

The exact theorem assumes exogenous mode observations. A useful sound extension does not: if finitely many admissible initial-state/disturbance scenarios are provably indistinguishable until  $\tau$ , give the controller after  $\tau$  separate arbitrary control signals for every stored scenario — an oracle relaxation of the true policy class. If

$$\int_0^\tau \phi_U \left( \sum_\ell \lambda_\ell k_\ell \right) dt + \sum_\ell \int_\tau^T \phi_U(\lambda_\ell k_\ell) dt - \sum_\ell \lambda_\ell \beta_\ell > 0, \quad (14)$$

with  $\beta_\ell$  computed from the explicitly stored scenario, then no actual output-feedback policy can be safe. This covers arbitrary sensing after a known blackout; it is sound but generally incomplete.

## The reference library

The finite side of the pipeline is packaged as `viacert`, a standard-library Python package in exact integer and ra-

| $h$  | blocks $Q$ | variables $mQ+1$ | rows $Ms(N_t+1) + p_U Q$ | exact $\rho(h)$           | solver $\rho$ | iterations | time (s) |
|------|------------|------------------|--------------------------|---------------------------|---------------|------------|----------|
| 0.20 | 16         | 33               | 243                      | 0                         | 0.000000000   | 36         | 0.024    |
| 0.10 | 32         | 65               | 465                      | $\frac{3}{190} = 0.030$   | 0.030000000   | 78         | 0.014    |
| 0.05 | 64         | 129              | 909                      | $\frac{200}{200} = 0.045$ | 0.045000000   | 141        | 0.043    |
| 0.02 | 160        | 321              | 2241                     | $\frac{27}{500} = 0.054$  | 0.054000000   | 501        | 0.252    |
| 0.01 | 320        | 641              | 4461                     | $\frac{57}{1000} = 0.057$ | 0.057000000   | 1062       | 1.015    |

**Table 1:** Solver-certified mesh study on the three-branch instance ( $\tau = \frac{1}{5}$ ,  $T = \frac{6}{5}$ ): program dimensions, exact optimal values, HiGHS-attained values (agreement to at most  $4.2 \times 10^{-17}$ ), simplex iterations, and wall-clock time. Optimality of each exact value is certified by the exact primal–dual witness pair, independent of the solver. The exact continuous value is 0.06; three safety labels support each positive certificate.

tional arithmetic. Its design follows three principles. *Exactness*: no floating point on any audited path; values, witnesses, and bounds are rationals. *Fidelity to the verified recursion*: the finite-horizon selector is the calculus’s backward recursion (Abaee, 2026b) with the no-repeat protocol threaded as a parameter, matching the audited implementations rather than a reformulation. *Pluggability*: new systems enter through a `FiniteSystem` interface (states, actions, safe set, transition, null action); the audited systems — the shared two-coordinate audit system, the continuous benchmark caps, the certainty-equivalence law pair, and the review grid — ship as default instances.

The public surface comprises the recursion module (viability verdicts, kernels, and step-one witness extraction), the monitoring module (safe and common safe-action sets, the partition-form adequacy converse, maximal common-action sets, the partition census, the review-timing identity with its threshold convention, and the certainty-equivalence drift audit), and a command-line self-test of twelve audited identities — the four-cell kernel counts 24/26/25/28 with the meet identity, the census 15 total and 7 adequate partitions with maximal sets overlapping in (2, 2), the benchmark  $Y^* = 27/5$ , the drift bounds 21/100 and 0, the 48-cell delay identity with its boundary cell, and witness extraction — each a pointer to the published audits (Abaee, 2026a, 2026b). In the certification workflow the library plays the enumerate-and-check role: finite instances are enumerated exactly, candidate designs are tested against the fibre criteria, and every claim is archived with the script that regenerates it. The library reimplements no theorem of this paper; the bridge’s programs live in the campaign scripts of Section .

## Scope delimitations

The exact theorem does not cover: control-dependent observations or active sensing; arbitrary continuous-state measurements; nonlinear dynamics without additional validated enclosures; nonconvex actuator sets without a relaxation gap; rate-constrained actuators without modifying the approximation; or unrestricted infinite-horizon completeness. For general sensing after a blackout, the oracle-recourse certificate is sound but incomplete. For nonlinear systems, a uniform validated flow-error bound can be added to the row budgets, yielding a sound extension, not automatic completeness. Against the alternatives: full-information viability decides none of the information-constrained question; initial admissibility and tangency tests see no obstruction here; safety-before-first-observation tests are silent when a common safe blind control exists; exact finite safety games await a justified

continuous-to-finite bridge, which is what Theorem 4 provides; and naïve held-control infeasibility excludes only the held-control class. No general-purpose computational-calculus claim is made: the theorem is a certificate generator for the stated class, with completeness only under refinement and semidecidability only with effectively computable bounds.

## Conclusion

Nonviability certification under incomplete observation is computationally well-posed exactly when three requirements hold: the discretization must relax the controller and never the adversary’s necessity direction; optimality must be certifiable independently of the solver; and certificate size must be understood in the information–time product rather than the input dimension. The bridge theorem delivers all three, the three-branch instance shows the resulting certificates detecting failure that full-information viability, blind-window exit tests, and held-control infeasibility all miss, and the campaign shows standard solvers attaining proved optima to numerical precision while exact witnesses close every gap. The reference library carries the exact finite machinery to new systems. What remains open is completeness beyond refinement: an inner approximation of the policy class that is computationally effective for general sensing is not known, and the oracle-recourse extension is the current sound surrogate.

## Verification methods

Every number in this paper is regenerated by deterministic scripts, exact rational arithmetic unless stated as solver output, standard library plus the solver interface only. Three scripts cover the claims. The exact verification record re-derives the instance’s certificates — delay threshold, margins, witnesses, erosion identity, rank construction — in 51 exact checks. The campaign script rebuilds the program (5) at each of the five meshes, runs HiGHS, and certifies optimality by the exact primal–dual witnesses, closing the duality gap; it re-derives the mesh law and all dimension identities (nine checks). The library self-test runs the twelve audited identities. A master script chains the three, then recomputes independently: the mesh values from the closed law  $\max\{0, 3/50 - Th/4\}$ , the Farkas contradiction  $-\frac{53}{100} + \frac{1}{2} = -\frac{3}{100}$ , the primal row value  $\frac{3}{100}$ , the margins  $\frac{3}{50}$  and  $\frac{3}{100}$ , the moment error  $e = \frac{3}{100}$ , the relaxed right-hand side  $-\frac{53}{100}$ , the delay threshold  $\tau^* = \frac{7}{50}$ , the blind-window safety value  $\frac{39}{25}$ , and every headline number in the text.

## References

Abae, A. An obstruction calculus for viability under incomplete observation (submitted).

Abae, A., 2026a. Robust viability of the 2J3KL limit reference point under a surplus-production map: policy scoring, expansion, and when catch cannot help. Zenodo. <https://doi.org/10.5281/zenodo.22552060>

Anderson, E.J., Nash, P., 1987. Linear Programming in Infinite-Dimensional Spaces: Theory and Applications. Wiley, Chichester.

Aubin, J.-P., 1991. Viability Theory. Birkhäuser, Boston.

Cardaliaguet, P., Quincampoix, M., Saint-Pierre, P., 1999. Set-valued numerical methods for optimal control and differential games, in: Stochastic and Differential Games: Theory and Numerical Methods, pp. 177–247. Birkhäuser, Boston.

Hettich, R., Kortanek, K.O., 1993. Semi-infinite programming: theory, methods, and applications. SIAM Rev. 35(3), 380–429.

Huangfu, Q., Hall, J.A.J., 2018. Parallelizing the dual revised simplex method. Math. Program. Comput. 10(1), 119–142.

Veliov, V.M., 1993. Sufficient conditions for viability under imperfect measurement. Set-Valued Anal. 1, 305–317.

## Declarations

### Funding

None.

### Competing interests

None.

### Data availability

No external data were used.

### Code availability

The verification scripts `paper2_computational_certification_v1_verification.py`, `paper2_nonstandard_verification.py`, `paper2c_lp_campaign.py`, and the `viacert` package are publicly available at <https://github.com/MIKEAA2020/general-sustainability> (folder `arena agent 1/paper rewrites/latex`); they reproduce all exact values, witnesses, bounds, and the solver campaign verbatim.

### AI declaration

GLM (Z.ai), Qwen (Alibaba Cloud) and DeepSeek AI assisted with drafting and iterative review. The author reviewed and edited outputs and takes responsibility for the final work.