

Supplementary Material: An Obstruction Calculus for Viability under Incomplete Observation

Amin Abaee

September 13, 2026

This supplement collects the material condensed from the main text: the complete proofs (S1), the expanded sufficiency review (S2), the bounded constructions and scope remarks (S3), and three additional figures (S4). Notation and numbering follow the main text.

S1. Complete proofs

Theorem 1 (finite-time exit certificate). Let $q : X \rightarrow \mathbb{R}$ be a constraint function of class C^1 — the class assumed throughout the exit theorems — with $\mathcal{V} \subseteq \{q \geq 0\}$, and suppose the following hypotheses hold.

(H1.1) There exist constants $a > 0$ and $\varepsilon > 0$ such that on the strip $\mathcal{S}_a = \{x : 0 \leq q(x) \leq a\}$,

$$\sup_{u \in U(x)} \inf_{d \in D(x)} D^+q(x; f(x, u, d)) \leq -\varepsilon \quad \forall x \in \mathcal{S}_a, \quad (1)$$

where $D^+q(x; v) = \limsup_{h \downarrow 0} [q(x + hv) - q(x)]/h$ is the upper right Dini derivative of q in direction v (for C^1 q , $D^+q(x; v) = \nabla q(x) \cdot v$ — the form the proof uses).

(H1.2) **Closed-loop existence of the adverse realization.** The disturbance correspondence is lower semicontinuous, or constant, on the strip, so that the state-dependent adverse selection constructed in the proof is realized by an admissible control–disturbance pair for which the Carathéodory existence theorem applies. (Alternatively, without (H1.2), the certificate is read against the convexified (relaxed) inclusion $\dot{x} \in \text{co } f(x, u(t), D_\varepsilon(x, u(t)))$, which admits solutions under Marchaud-type conditions and preserves the drift inequality because the constraint is convex in the velocity; closure and relaxation are the standard tools for this step.)

Then, under hypotheses (H1.1)–(H1.2), for every admissible control and every initial state $x_0 \in \mathcal{S}_a$ there exists an admissible disturbance realization under which the trajectory attains $q(x(t)) < 0$ — hence leaves $\{q \geq 0\} \supseteq \mathcal{V}$ — by every time strictly larger than $q(x_0)/\varepsilon$: $\inf\{t : q(x(t)) < 0\} \leq q(x_0)/\varepsilon$, in particular within time at most a/ε (since $q(x_0) \leq a$). (Convention, used in both exit theorems: *violation* is the strict inequality $q < 0$; at $q = 0$ the state may still lie on $\partial\mathcal{V} \subseteq \mathcal{V}$.)

Proof. Trajectories of $\dot{x} = f(x, u(t), d(t))$ exist in the Carathéodory sense under the standing hypotheses — D compact-valued, f measurable in t and continuous in (x, u, d) , linearly bounded — for every measurable control–disturbance pair. The adverse realization selected below is *state-dependent*, so its closed-loop realization is an additional requirement — the content of (H1.2) (D lower semicontinuous or constant, or the convexified reading of its parenthetical); convexity of D was not needed for the closure step — closure and relaxation are the standard tools for it. Fix an admissible control $u(\cdot)$. For each $x \in \mathcal{S}_a$, condition (1) gives $\inf_{d \in D(x)} D^+q(x; f(x, u(t), d)) \leq -\varepsilon$. Define the adverse-selection correspondence

$$D_\varepsilon(x, u) = \{d \in D(x) : D^+q(x; f(x, u, d)) \leq -\varepsilon\},$$

which has nonempty values on $\mathcal{S}_a$ by (1). For q of class C^1 the Dini derivative reduces to $\nabla q(x) \cdot f(x, u, d)$, continuous in (x, u, d) ; with the closed graph and compact values of D (Section 2.1), D_ε then has closed graph and compact values, and the measurable-selection theorem (Aubin and Frankowska, 1990, Thm. 8.1.3) supplies a measurable selection $d(\cdot)$ with $d(t) \in D_\varepsilon(x(t), u(t))$ along the trajectory. (The proof is carried out for q of class C^1 , as assumed; the locally Lipschitz case requires a separate comparison argument for almost-everywhere-differentiable arcs and is not treated.) Along the realization $(u(\cdot), d(\cdot))$ from x_0 , the fundamental theorem of calculus along the absolutely continuous trajectory integrates the inequality to

$$q(x(t)) \leq q(x_0) - \varepsilon t$$

as long as $x(t)$ remains in $\mathcal{S}_a$. Hence $q \leq 0$ at a time not exceeding $q(x_0)/\varepsilon \leq a/\varepsilon$. If the trajectory has not already left the strip by time $t^* + \delta$ for some $\delta > 0$, the comparison continues to apply and gives $q(x(t^* + \delta)) \leq q(x_0) - \varepsilon(t^* + \delta) \leq -\varepsilon\delta < 0$ with $t^* = q(x_0)/\varepsilon$; in either case the trajectory violates $\{q \geq 0\} \supseteq \mathcal{V}$ before every time strictly larger than $q(x_0)/\varepsilon$. The selection $d(t) \in D_\varepsilon(x(t), u(t))$ is itself nonanticipative feedback of the pair (x, u) — a disturbance strategy in the Isaacs order, responding to the current state and control without anticipating future controls — which is what the certificate requires; no Elliott–Kalton upgrade to a closed-loop strategy is required for the stated claim. $\square$

Proposition 1 (epistemic emptiness by admissibility — a minimal construction). There exists a system with $\text{Viab}(\mathcal{V}; U, \pi_{\text{perf}}) = \mathcal{V} \neq \emptyset$ such that, for a non-injective observation map O , every admissible initial information state — the observation fibres $B_0 = O^{-1}(O(S_0))$, $S_0 \in \mathcal{V}$ — lies outside $\text{ERViab}_{\mathcal{I}}(\mathcal{V})$: the epistemic kernel over the fibre-induced initial beliefs is empty. The mechanism is *admissibility* rather than a hidden mode: the state-dependent control set of the construction leaves a single common action at the merged belief — one that is outward-drifting at the boundary state — so the emptiness is the common-action obstruction (Theorem 2’s safety case) at minimal size, produced by the control-set primitive rather than by information; the hidden-mode instance is Example 1.

Proof (by construction). Let $X = [1, 2]$, $\dot{S} = u - r(S)$ with $r : [1, 2] \rightarrow \mathbb{R}_{++}$ continuous and **injective on** $[1, 2]$ (e.g. $r(S) = S - 1 + c$, $c > 0$, strictly increasing), $U(S) = \{0, r(S)\}$, $\mathcal{V} = [1, 2]$, and the constant observation $O(S) \equiv 0$. Under perfect information, the state-feedback control $u = r(S)$ gives $\dot{S} = 0$, so every $S_0 \in \mathcal{V}$ is viable: $\text{Viab}(\mathcal{V}; U, \pi_{\text{perf}}) = \mathcal{V}$. Under the constant observation, the initial information state is the fibre $B_0 = [1, 2]$, and the set-membership semantics propagates it: at time t the belief is $B_t = \Phi_t(B_0) \cap \mathcal{V}$, the image of the initial set under the flow generated by the applied control — under a constant observation with no exit observation, this flow-image coincides with the record-compatibility definition of Section 2.3. The common control set at the initial belief is

$$U^B(B_0) = \bigcap_{S \in [1, 2]} \{0, r(S)\} = \{0\},$$

because $r(S) > 0$ for every S and, by injectivity on $[1, 2]$, no single control equals $r(S)$ for all S . An observation-based policy may select time-varying actions, but robust admissibility requires the selected action to lie in $U^B(B_t)$ at every t ; the propagated belief B_t is nonempty and, while the state remains in $[1, 2]$, $U^B(B_t) = \{0\}$ by the same computation (the flow of $\dot{S} = -r(S)$ with $r > 0$ moves B_t monotonically downward, preserving the injectivity of r on the surviving interval). Hence the policy must hold $u(t) = 0$ almost everywhere, giving $\dot{S} = -r(S) \leq -\min_{[1, 2]} r < 0$; every compatible trajectory exits $\mathcal{V}$ downward in finite time. The same argument applies verbatim to any initial fibre, since O is constant. Hence $\text{ERViab}_{\mathcal{I}}(\mathcal{V})$ contains no fibre-induced initial belief, while beliefs that are singletons (exact prior knowledge) are not admissible initial states of this observation structure. The system has no disturbance and no estimation error: the failure is caused entirely by the non-injectivity of O . $\square$

Theorem 2 (common-action obstruction). Let $B \subseteq \mathcal{V}$ be an information set of the specified observation structure (the safe-control correspondence $\mathcal{R}_{\mathcal{V}}$ is defined on $\mathcal{V}$, so the belief is taken inside it). Suppose the following hypotheses hold.

(H3.1) The selected action is held until the next observation (sampled review: no informative observation arrives before the next action must be selected, and the policy holds the selected action over the review interval — the institutional reading; without a positive dwell time the instantaneous statement below must be replaced by the tube-safety statement below).

(H3.2) Either the **common admissible action set** is empty,

$$U^B(B) = \bigcap_{x \in B} U(x) = \emptyset, \quad (\text{admissibility obstruction})$$

or the **common safe-action set** is empty while common admissibility holds,

$$\mathcal{R}_{\mathcal{V}}^B(B) := \bigcap_{x \in B} \mathcal{R}_{\mathcal{V}}(x) = \emptyset \quad \text{with} \quad U^B(B) \neq \emptyset, \quad (\text{safety obstruction}) \quad (2)$$

(H3.3) **Closed-loop existence of the local adverse selection (safety case).** The disturbance correspondence is lower semicontinuous, or locally constant near the boundary point at which the obstruction fires, so that the local adverse selection of the proof is realized along the compatible trajectory; alternatively, the certificate is read against the convexified (relaxed) inclusion, which preserves the signed drift (the constraint is convex in the velocity).

Then, under hypotheses (H3.1)–(H3.3), $B \notin \text{ERViab}_{\mathcal{I}}(\mathcal{V})$: every compatible state may be individually robustly viable, and the belief is nevertheless nonviable, because the state-specific safe actions are incompatible.

Proof. Let π be any observation-based policy, and let $a = \pi(B)$ be the action it selects at the information set B .

Admissibility case. If $a \notin U^B(B)$, then for some compatible state $x \in B$ the action is inadmissible ($a \notin U(x)$): at that state the selected action is not an admissible control, so the policy fails outright on the compatible branch from x — an obstruction independent of the safe set.

Safety case. Otherwise $a \in U^B(B)$. Because $\bigcap_{x \in B} \mathcal{R}_{\mathcal{V}}(x) = \emptyset$ while $\mathcal{R}_{\mathcal{V}}(x) = U(x)$ at interior points of $\mathcal{V}$ (no constraint is active there) and $a \in U(x)$ for every $x \in B$, the emptiness of the intersection must be caused by the boundary: there is a state $\bar{x} \in B \cap \partial\mathcal{V}$ with $a \notin \mathcal{R}_{\mathcal{V}}(\bar{x})$. By the defining inequality of the safe-control correspondence, there are an active constraint j (with $q_j(\bar{x}) = 0$) and a disturbance $\bar{d} \in D(\bar{x})$ such that

$$\nabla q_j(\bar{x}) \cdot f(\bar{x}, a, \bar{d}) < 0.$$

With $\eta = -\nabla q_j(\bar{x}) \cdot f(\bar{x}, a, \bar{d}) > 0$, define the adverse-selection correspondence $D_{\eta}(x) = \{d \in D(x) : \nabla q_j(x) \cdot f(x, a, d) \leq -\eta/2\}$ on a neighbourhood of $\bar{x}$; by continuity of the right-hand side in (x, d) and the closed graph of D (Section 2.1), it has nonempty compact values near $\bar{x}$ and admits a measurable local selection $d(x)$ with $d(\bar{x}) = \bar{d}$. Let the disturbance realize this selection on a small initial interval (its closed-loop realization is hypothesis (H3.3)’s content — D lower semicontinuous or locally constant near $\bar{x}$, or the convexified reading); then along the compatible trajectory from $\bar{x}$ (compatible, since $\bar{x} \in B$), q_j strictly decreases from $q_j(\bar{x}) = 0$ at rate at least $\eta/2$, so the trajectory leaves $\mathcal{V}$ within the first review period — before any informative observation can arrive, whichever action the policy takes. Since π was arbitrary, no observation-based policy keeps every compatible trajectory in $\mathcal{V}$: $B \notin \text{ERViab}_{\mathcal{I}}(\mathcal{V})$. $\square$

Proposition 2 (uniform margin gives a tube obstruction at every review length). Let $B \subseteq \mathcal{V}$ be a compact information set with $U^B(B) \neq \emptyset$, f and ∇q_j continuous, and suppose there is

$\eta > 0$ such that for every $a \in U^B(B)$ there exist a boundary state $x_a \in B \cap \partial\mathcal{V}$, an active constraint j with $q_j(x_a) = 0$, and a disturbance $d_a \in D(x_a)$ satisfying

$$\nabla q_j(x_a) \cdot f(x_a, a, d_a) \leq -\eta.$$

Under the closed-loop existence convention of Theorem 2, the tube-safe action set is then empty at every review length,

$$\mathcal{A}_{\text{tube}}(B, \Delta) = \emptyset \quad \forall \Delta > 0,$$

and hence $B \notin \text{ERViab}_{\mathcal{I}}(\mathcal{V})$ under sampled review of any length.

Proof. Fix $a \in U^B(B)$ and the triple (x_a, j, d_a) . By continuity of $(x, u, d) \mapsto \nabla q_j(x) \cdot f(x, u, d)$ and the closed graph of D , the drift inequality holds with margin $-\eta/2$ on a neighbourhood of (x_a, a, d_a) , and a measurable local selection $d(x)$ realizes it. Along the trajectory from x_a under the held action a and this adverse selection, $q_j(x(t)) \leq q_j(x_a) - (\eta/2)t = -(\eta/2)t < 0$ for every small $t > 0$, so the trajectory leaves $\mathcal{V}$ in arbitrarily small positive time; hence $\text{Reach}_{[0, \Delta]}(B, a) \not\subseteq \mathcal{V}$ for every $\Delta > 0$, i.e. $a \notin \mathcal{A}_{\text{tube}}(B, \Delta)$. Since $a \in U^B(B)$ was arbitrary, $\mathcal{A}_{\text{tube}}(B, \Delta) = \emptyset$ for every $\Delta > 0$, and the tube-safety form certifies $B \notin \text{ERViab}_{\mathcal{I}}(\mathcal{V})$. $\square$

Proposition 3 (obstruction ladder). Under the standing regularity assumptions and the closed-loop existence convention of Theorem 2, the common response sets of Section 2.4 are nested:

$$\mathcal{A}_{\text{tube}}(B, \Delta) \subseteq \mathcal{R}_{\mathcal{V}}^B(B) \subseteq U^B(B) \quad \forall \Delta > 0,$$

and emptiness descends the ladder:

$$U^B(B) = \emptyset \implies \mathcal{R}_{\mathcal{V}}^B(B) = \emptyset \implies \mathcal{A}_{\text{tube}}(B, \Delta) = \emptyset \implies B \notin \text{ERViab}_{\mathcal{I}}(\mathcal{V}).$$

Proof. The second inclusion is by definition, since $\mathcal{R}_{\mathcal{V}}(x) \subseteq U(x)$ for every x . For the first, fix $a \in \mathcal{A}_{\text{tube}}(B, \Delta)$, $x \in B$ with an active constraint j ($q_j(x) = 0$), and $d \in D(x)$. If $\nabla q_j(x) \cdot f(x, a, d) < 0$, then along the trajectory from x under the held action a and the constant disturbance d , $q_j(x(t)) = \nabla q_j(x) \cdot f(x, a, d)t + o(t) < 0$ for all small $t > 0$, so the trajectory leaves $\mathcal{V}$ in arbitrarily small time, contradicting $a \in \mathcal{A}_{\text{tube}}(B, \Delta)$. Hence $\nabla q_j(x) \cdot f(x, a, d) \geq 0$ for every $d \in D(x)$ and every active j , i.e. $a \in \mathcal{R}_{\mathcal{V}}^B(B)$. $\square$

Theorem 3 (delayed-information obstruction). Let B_0 be an initial information set. Suppose the following hypotheses hold.

(H4.1) No informative observation arrives before time $T_{\text{obs}} > 0$ — formally, there are branches that are **observation-equivalent on** $[0, T_{\text{obs}})$: they produce identical observation records throughout that interval, so the policy cannot distinguish them before T_{obs} ; between observations the information set evolves by the set-membership semantics of Section 2.3.

(H4.2) There exist a constraint function q of class C^1 (the class of Theorem 1's proof) with $\mathcal{V} \subseteq \{q \geq 0\}$ and a constant $\varepsilon > 0$ such that, **for every open-loop control on the blind window** — every measurable $u(\cdot) : [0, T_{\text{obs}}] \rightarrow \bigcup_{x \in B_0} U(x)$, the class of controls that any observation-based policy realizes before the first informative observation, where its actions are a fixed function of time — there are a compatible initial state $x^* \in B_0$ attaining $q(x^*) = \inf_{x \in B_0} q(x)$ (with B_0 compact and q lower semicontinuous; otherwise replace the infimum by $\inf q + \delta$ for an arbitrary δ -minimizer, let $\delta \downarrow 0$, and read the timing bound with $\inf q + \delta$) and an admissible disturbance realization such that, along the realized trajectory from x^* under that open-loop control, the record remains compatible and

$$D^+q(x(t); f(x(t), u(t), d(t))) \leq -\varepsilon \quad \text{while } q(x(t)) > 0, \quad (3)$$

(H4.3) The timing bound holds:

$$T_{\text{obs}} > \frac{\inf_{x \in B_0} q(x)}{\varepsilon}. \quad (4)$$

Then, under hypotheses (H4.1)–(H4.3), $B_0 \notin \text{ERViab}_{\mathcal{I}}(\mathcal{V})$: information may be accurate but arrive too late.

Proof. Fix any observation-based policy. By (H4.1) its actions on $[0, T_{\text{obs}})$ are a fixed open-loop function of time — the record carries no informative observation before T_{obs} , so every observation-equivalent branch induces the same actions — and (H4.2), applied to that open-loop control, supplies a compatible initial state x^* with $q(x^*) = \inf_{x \in B_0} q(x)$ and an admissible disturbance realization satisfying (3) along the trajectory from x^* under those actions, observation-equivalent on $[0, T_{\text{obs}})$ to the other compatible branches. For q of class C^1 — the class assumed in (H4.2) — the fundamental theorem of calculus along the absolutely continuous trajectory integrates (3) to

$$q(x(t)) \leq q(x^*) - \varepsilon t = \inf_{x \in B_0} q(x) - \varepsilon t$$

while $q(x(t)) > 0$, so the trajectory attains the violation convention of Theorem 1 — the strict inequality $q(x(t)) < 0$ — by every time strictly larger than $t^* = \inf_{x \in B_0} q(x)/\varepsilon$ (in particular $\inf\{t : q(x(t)) < 0\} \leq t^*$; the same argument with the δ -minimizer bounds the violation time by $(\inf q + \delta)/\varepsilon$), and t^* by (4) is strictly less than T_{obs} — the violation is strict and precedes the first informative observation. Until t^* , the record contains no informative observation: the violating branch is observation-equivalent to every other compatible branch before T_{obs} , so the policy cannot have altered its action in response to the violation before it occurs. Since the policy was arbitrary and the violating realization is compatible with the observation record, no observation-based policy is robustly viable: $B_0 \notin \text{ERViab}_{\mathcal{I}}(\mathcal{V})$. $\square$

Theorem 4 (finite-horizon soundness and completeness). An initial belief B_0 admits an observation-based policy that keeps every compatible trajectory in $\mathcal{V}$ for N steps if and only if $B_0 \in \mathcal{W}_N$.

Proof. ($\Rightarrow$) Induction on N . The case $N = 0$ is $B_0 \subseteq \mathcal{V}$, i.e. $B_0 \in \mathcal{W}_0$. If a policy is safe for N steps, its first action a is admissible at every compatible state, so $a \in U^B(B_0)$; for every possible observation y , the successor belief $\text{Post}(B_0, a, y)$ is safe for $N - 1$ steps under the same policy, so by induction $\text{Post}(B_0, a, y) \in \mathcal{W}_{N-1}$; hence $B_0 \in \text{Pre}(\mathcal{W}_{N-1}) = \mathcal{W}_N$.

($\Leftarrow$) Backward induction. If $B_k \in \mathcal{W}_{N-k} = \text{Pre}(\mathcal{W}_{N-k-1})$, choose the witness action $a_k \in U^B(B_k)$; after the observation y_k the successor belief $B_{k+1} = \text{Post}(B_k, a_k, y_k) \in \mathcal{W}_{N-k-1}$, and the recursion continues. By construction every reachable state over N steps lies in $\mathcal{V}$. $\square$

Proposition 4 (observation-fibre criterion). An exact certifier based on O exists if and only if membership in K is constant on every observation fibre:

$$O(z_1) = O(z_2) \implies [z_1 \in K \iff z_2 \in K] \quad \forall z_1, z_2 \in Z, \quad (5)$$

equivalently, $K = O^{-1}(O(K))$ on the admissible domain.

Proof. ($\Rightarrow$) If C exists and $O(z_1) = O(z_2)$, then $\mathbf{1}_K(z_1) = C(O(z_1)) = C(O(z_2)) = \mathbf{1}_K(z_2)$, so membership is fibre-constant.

($\Leftarrow$) If (5) holds, define $C(y) = \mathbf{1}_K(z)$ for any $z \in Z$ with $O(z) = y$; (5) makes C well defined on $O(Z)$ and exact by construction; extend C arbitrarily on $Y \setminus O(Z)$. A *measurable* exact certifier exists under the same condition on standard Borel spaces when $O(K)$ is measurable, taking $C = \mathbf{1}_{O(K)}$; measurability is the only additional requirement. $\square$

Corollary 1 (safety-crossing fibres and the certainly-safe set). If two admissible states share an observation and lie on opposite sides of the *same* component safety constraint — with the safe set of that certification problem read as the single floor $K = \{q_j \geq 0\}$; or, in the general form of the hypothesis, one of the two states lies in K and the other outside K — no exact observation-only certificate exists. (Crossing one floor of an intersection $K = \bigcap_j \{q_j \geq 0\}$ does not by itself put the pair on opposite sides of the intersection — the state satisfying floor j may violate floor i — hence the single-floor scoping.) The largest set of observations that can soundly be labeled safe without further information is the certainly-safe set

$$\mathcal{Y}_{\text{safe}} = \{y \in O(Z) : O^{-1}(y) \subseteq K\}.$$

Proof. In the general form (one state in K , one outside), the fibre violates (5) for that K , so Proposition 4 denies the certifier; in the single-floor form it violates (5) for the floor’s own certification problem ($K = \{q_j \geq 0\}$). On $\mathcal{Y}_{\text{safe}}$ every compatible state is safe, so the constant verdict “safe” is sound; outside $\mathcal{Y}_{\text{safe}}$ some compatible state is unsafe, so no sound “safe” verdict exists there. $\square$

Proposition 5 (monotonicity of the epistemic kernel). The robust epistemic kernel is monotone in its constituents: enlarging the action sets, reducing the disturbance sets, refining the information structure, or enlarging the policy class never shrinks the kernel. Writing $\text{ERViab}^{\cdot}$ for the kernel with the indicated constituent, $U_1 \subseteq U_2$, $D_1 \subseteq D_2$, $\Pi_1 \subseteq \Pi_2$, and “ $\mathcal{I}_1$ refines $\mathcal{I}_2$ ” imply respectively

$$\begin{aligned} \text{ERViab}^{U_1} &\subseteq \text{ERViab}^{U_2}, & \text{ERViab}^{D_2} &\subseteq \text{ERViab}^{D_1}, \\ \text{ERViab}^{\Pi_1} &\subseteq \text{ERViab}^{\Pi_2}, & \text{ERViab}^{\mathcal{I}_2} &\subseteq \text{ERViab}^{\mathcal{I}_1}. \end{aligned}$$

Proof. Each inclusion is by restriction: a policy that is viable under the smaller action sets, the larger disturbance sets, the smaller policy class, or the coarser information structure remains admissible and nonanticipative under the corresponding enlargement or refinement, while the adversary’s options are unchanged or reduced, so the same policy witnesses viability. $\square$

S2. The sufficiency landscape (full)

5. The Sufficiency Landscape

The sufficiency results against which the obstruction calculus is defined are standard; they are stated here for contrast, and proofs appear in the cited sources.

(a) Veliov’s output-feedback condition. Veliov (1993) considers the same problem as Section 3 (a tube in the state space under incomplete and inexact measurement) and gives a sufficient condition for the existence of an *output-feedback regulation map*: a set-valued feedback of the measurement under which all trajectories starting from the graph of the tube remain in it. Under perfect measurement his condition reduces to the classical viability condition (Haddad, 1981). The complementarity with this paper is exact: Veliov’s theorem certifies existence; the obstruction calculus certifies nonexistence; a problem that satisfies neither requires a stronger observer or a finer observation structure.

(b) The estimation-tube programme. Cardaliaguet, Quincampoix, and Saint-Pierre (2007) pass from imperfect information in the measurement space to perfect information in an estimation space of *sets* of compatible states: the value functions of the two problems coincide, and the estimation-space problem admits a Dini-derivative characterization. The reduction is exact for value functions; what it does not supply is a *common prescription*. The exact relationship is stated in Section 6.3: carried out under the set-membership semantics, the estimation-space solution propagates the belief under a single control signal, so its controls *are* common controls and joint admissibility is built into the estimation-space problem’s own admissibility. Pointwise selections taken without that joint-admissibility check need not be jointly admissible; the common-action obstruction (Theorem 2) is the certificate that no jointly admissible selection exists at the information state in question.

(c) Observer-and-buffer transfer. When a full-information feedback exists with a strict inward margin, and an observer supplies exponentially convergent estimates, output feedback of the estimate preserves safety on a buffered subset: if K_ζ is a compact controlled-invariant subset of the interior of the kernel, the state feedback $u = k(x)$ is Lipschitz with constant L_k , the observer satisfies $\|\hat{x}(t) - x(t)\| \leq Me^{-\lambda t}\|\hat{x}(0) - x(0)\|$, and the initial estimation error is small enough that $L_k M\|\hat{x}(0) - x(0)\|$ is absorbed by the margin, then K_ζ is viable under output feedback (observer-to-viability transfer with safety buffer; the standard separation argument

of observer-based control (Sontag, 1998)). Similarly, eroded kernels: if K is robustly invariant under full-state feedback with a strict inward margin on ∂K , and estimation and implementation errors are bounded by ζ , then an eroded set $K^{-c\zeta}$ is invariant under output feedback for a sensitivity constant $c > 0$ (erosion absorbs the error; see the buffer constructions in the literature cited in Section 1.3). These results bound the sufficiency side: under margins and convergence, output feedback *can* work; the obstruction calculus of Sections 3–4 states the conditions under which it *cannot*.

(d) The linear substitution alternative. For the finite linear model in which a resource-typed system must meet a demand vector through specified substitution pathways, exactly one of the following holds (Farkas, 1902; Gale, 1960): (i) there exists a non-negative pathway vector $a \geq 0$ satisfying the linear substitution constraints; or (ii) there exist multipliers $\alpha, \beta, \gamma \geq 0$ such that

$$\alpha^\top R + \beta^\top E - \gamma^\top Q \geq 0 \quad \text{componentwise}, \quad \gamma^\top s^{\text{req}} > \alpha^\top x + \beta^\top e.$$

The second statement is a certificate that the specified substitution pathways cannot meet demand within the typed resource and capacity bounds; writing all constraints as $Aa \leq \rho$, the pair is the Farkas lemma alternative. The multipliers are a separation certificate, not universal exchange rates: nonlinear, nonconvex, path-dependent, spatial, or irreversible technologies require their own feasibility analysis, and an elasticity fitted near one operating point cannot establish global substitutability. This is the feasibility-side complement of the observation-side obstructions of Sections 3–4: where the present paper’s certificates bound what *information* can achieve, the substitution alternative bounds what *material pathways* can achieve, and the same infeasibility logic — exhibit the separating multiplier rather than assert impossibility — applies to both.

S3. Bounded constructions and scope remarks

Appendix A: Bounded Constructions and Scope Remarks

The bounded constructions below complement the obstruction theorems of Sections 3–4. They are stated in full, together with their scope remarks.

A.1 Example (coupling creates viability absent in a factor). Take the two-patch logistic system $\dot{S}_i = g_i(S_i) + \kappa(S_j - S_i) - h_i$, $i, j \in \{1, 2\}$, $i \neq j$, with $g_i(s) = s(1 - s)$, $\kappa = 0.2$, constraints $S_i \in [0, 1]$ and $h_i \geq h_{\min, i}$ (harvest floors), and admissible controls the constant harvest vectors $h = (h_1, h_2)$ with $h_i \geq 0$. Choose $(S_1^*, S_2^*) = (0.5, 0.8)$. Define harvest floors by the equilibrium equations: $h_{\min, 1} = g_1(0.5) + 0.2(0.8 - 0.5) = 0.31$; $h_{\min, 2} = g_2(0.8) + 0.2(0.5 - 0.8) = 0.10$. Patch 1 in isolation: $\max_s g_1(s) = 0.25 < h_{\min, 1} = 0.31$, so patch 1 is not viable in isolation. Yet the coupled system has the equilibrium $(0.5, 0.8)$ with $h = (0.31, 0.10)$ admissible and the stock constraints satisfied, so its kernel is nonempty. The example is the constructive counterpart of the emptiness results of Section 3: where the obstruction theorems certify that no policy works for reasons of information, patch coupling exhibits the opposite direction — a jointly viable operating point that no factor sustains alone.

A.2 Counterexample (emptiness despite factorwise viability at MSY). Take $\kappa > 0$, $C_1 \neq C_2$, $h_{\min, i} = r_i C_i / 4$ (MSY level), in the same coupled system with $\phi_i(S_i) := g_i(S_i) - h_{\min, i}$ and growth laws $g_i(S) = r_i S(1 - S/C_i)$, so that

$$\phi_i(S_i) = -\frac{r_i}{C_i} \left(S_i - \frac{C_i}{2} \right)^2 \leq 0$$

with equality only at $S_i = C_i/2$, and the constraint set is the box $[C_1/2, \bar{S}_1] \times [C_2/2, \bar{S}_2]$. Each isolated system has kernel $[C_i/2, \bar{S}_i]$ (from $S_i > C_i/2$ the stock declines toward $C_i/2$

asymptotically, never below). For the coupled system, define the Lyapunov function $W = S_1 + S_2$; then

$$\dot{W} = \phi_1(S_1) + \phi_2(S_2) \leq 0,$$

with equality exactly at $p^* = (C_1/2, C_2/2)$, where the flow is nevertheless nonzero: $(\dot{S}_1, \dot{S}_2)|_{p^*} = (\frac{\kappa}{2}(C_2 - C_1), \frac{\kappa}{2}(C_1 - C_2)) \neq 0$. Suppose a trajectory remained in the constraint box for all time. Then W is non-increasing and bounded below by $(C_1 + C_2)/2$, so $W(t) \downarrow W_\infty \geq (C_1 + C_2)/2$; the ω -limit set of the trajectory is nonempty, compact, and invariant, and $\dot{W} \equiv 0$ on it (otherwise W would keep decreasing along the set). Hence the ω -limit set is contained in $\{p^*\}$, so it equals $\{p^*\}$ and the trajectory converges to p^* . But convergence to p^* is impossible: along a trajectory with $x(t) \rightarrow p^*$ the velocity satisfies $\dot{x}(t) = f(x(t)) \rightarrow f(p^*) \neq 0$, so $x(t) - p^* \sim (t - t_0)f(p^*)$, contradicting convergence. Therefore no trajectory remains admissible for all time: the kernel is empty, and since the box is compact and the vector field smooth, every trajectory violates the constraint in finite time. Note the logical order: the emptiness conclusion here rests on the Lyapunov- ω -limit argument, *not* on the absence of an equilibrium — absence of equilibria alone does not imply an empty kernel (periodic and nonstationary viable trajectories are possible in general), and the two are kept separate. The counterexample is specific to the MSY parameter choice; for $H_{\min,i} < r_i C_i/4$, equilibria may exist. Together with A.1 it delimits what coupling can and cannot repair: factorwise nonviability is repairable (A.1), but factorwise viability at an incompatible operating point is not (A.2).

S4. Additional figures




Figure 1: The obstruction ladder. The common response sets are nested, $U^B(B) \supseteq \mathcal{R}_V^B(B) \supseteq \mathcal{A}_{\text{tube}}(B, \Delta) \supseteq \mathcal{A}_N(B)$, and emptiness descends (Proposition 3); the finite-time exit certificate sits below the ladder, defeating every control under full information (Theorem 1).

Auxiliary definitions and remarks (for reference)

Definition 1 (robust epistemic kernel). $B_0 \in \text{ERViab}_{\mathcal{I}}(\mathcal{V})$ if there exists a nonanticipative observation-based policy π such that, for every compatible initial state $x_0 \in B_0$ and every admissible disturbance realization $d(\cdot)$, the trajectory of $\dot{x} = f(x, u(t), d(t))$ generated from x_0




Figure 2: The one-step obstruction tree of Example 1, read as a finite system: for each admissible action the adversary has a compatible branch terminating in violation, so $B_0 \notin \mathcal{W}_1$ (Theorem 4).




Figure 3: The certainty-equivalence trap. Under the perfect-information law $u = g(S)$ the stock is stationary inside $\mathcal{V}$. Applying the same law to the biased observation $u = g(\hat{S})$ with $\hat{S} = S + b$ gives $\dot{S} = g(S + b) - g(S) > 0$, so S exits above S^* in finite time (Remark 3); inverting the bias restores $g(S)$.

by π 's actions under $d(\cdot)$ — the record being the one that $(\pi, d(\cdot))$ generates — remains in $\mathcal{V}$ for all time. (This is the standard exists-strategy-for-all-realizations form. The quantifier order is essential: the record is generated by the realization, so conditioning the admissible realizations on the record would be circular.) Admissible initial beliefs are the compact information sets $B_0 \subseteq X$ generated by the observation structure — the observation fibres $O^{-1}(y)$, $y \in O(X)$, in the constructions of Sections 3.2–3.4. The existential (non-robust) counterpart $\text{ERViab}_{\mathcal{I}}(\mathcal{V})$ is the non-robust contrast class.

Proposition 6 (selector principle). If $B \in \text{ERViab}_{\mathcal{I}}(\mathcal{V})$ and π is a policy witnessing viability on B , then at every review time the action $a = \pi(B)$ selected by π is a response that is (i) *admissible*, $a \in U^B(B)$; (ii) *safe*, keeping every compatible trajectory in $\mathcal{V}$ until the next observation — $\text{Reach}_{[0, \Delta]}(B, a) \subseteq \mathcal{V}$ at review length Δ ; and (iii) *recursively viable*, in that every post-observation belief reached at the next review is again in $\text{ERViab}_{\mathcal{I}}(\mathcal{V})$ under the continuation of π . A belief is therefore viable only if each of the response correspondences of this section — $U^B(B)$, $\mathcal{R}_{\mathcal{V}}^B(B)$, $\mathcal{A}_{\text{tube}}(B, \Delta)$, and the recursive predecessor of Section 3.5 — is nonempty along every reachable branch; hence a certificate that any one of them is empty at a reachable belief certifies nonviability. In finite systems the converse is exact (Theorem 4); in continuous time the converse is the backward-recursion completion noted in Section 6.5.

Remark 1 (comparison-function form). Hypothesis (H1.1) uses a constant drift margin ε . The certificate extends to state-dependent decline: if there is a continuous $\alpha : (0, a] \rightarrow (0, \infty)$ with

$$\sup_{u \in U(x)} \inf_{d \in D(x)} D^+ q(x; f(x, u, d)) \leq -\alpha(q(x)) \quad \forall x \in \mathcal{S}_a,$$

then the exit-time bound sharpens to

$$t_{\text{exit}} \leq \int_0^{q(x_0)} \frac{ds}{\alpha(s)},$$

whenever the integral is finite. The proof replaces the linear comparison $q(x(t)) \leq q(x_0) - \varepsilon t$ with the solution r of $\dot{r} = -\alpha(r)$, $r(0) = q(x_0)$, by the standard comparison argument. If $\alpha(0) = 0$ and the integral diverges, the conclusion weakens to asymptotic approach rather than finite-time exit; the constant-margin case $\alpha \equiv \varepsilon$ is the sharp finite-time instance.

Example 1 (hidden-mode conflict). Let an unobserved parameter satisfy $\theta \in \{-1, +1\}$, with $\dot{z} = \theta u$, $u \in \{-1, +1\}$, and constraint $z \geq 0$. At $z = 0$: if $\theta = +1$, only $u = +1$ is safe; if $\theta = -1$, only $u = -1$ is safe. Both compatible states are individually robustly viable under full information (each admits its safe control), but the information set $B = \{(0, +1), (0, -1)\}$ satisfies $\mathcal{R}_{\mathcal{V}}^B(B) = \{+1\} \cap \{-1\} = \emptyset$, so $B \notin \text{ERViab}_{\mathcal{I}}(\mathcal{V})$. This is a purely informational failure: no stochasticity and no estimation-quality argument is involved. The example is the two-action skeleton of every “the stock is either recovering or collapsing, and the safe policy differs between the two” situation.

Remark 2 (threshold form of the timing obstruction). The delayed-information obstruction has a sharp threshold. Let $\sigma^*(B_0)$ be the guaranteed blind-window survival time: the largest duration such that some pre-observation policy keeps every observation-equivalent branch in $\mathcal{V}$ until then,

$$\sigma^*(B_0) = \sup_{u(\cdot)} \inf_{x_0 \in B_0, d(\cdot)} \tau(x_0, u, d),$$

with the infimum over branches that remain observation-equivalent on $[0, T_{\text{obs}})$ and τ the first violation time under the convention of Theorem 3. Then

$$\sigma^*(B_0) < T_{\text{obs}} \implies B_0 \notin \text{ERViab}_{\mathcal{I}}(\mathcal{V}),$$

because before T_{obs} a policy cannot condition its action on which branch realizes. Condition (4) certifies this threshold: under (H4.1)–(H4.2) every branch exits by time $\inf_{x \in B_0} q(x)/\varepsilon$, hence $\sigma^*(B_0) \leq \inf_{x \in B_0} q(x)/\varepsilon$, and (4) states that this upper bound lies below T_{obs} ; the threshold itself may be strictly smaller, in which case nonviability holds even when (4) fails. The threshold is the blind-window response correspondence of the timing layer: its emptiness is the timing certificate, the delayed-information analogue of the tube and common-action correspondences.

Definition 2 (exact certifier). For an admissible domain $Z \subseteq X$ and a safe set $K \subseteq Z$, an **exact certifier** based on the observation map $O : Z \rightarrow Y$ is a function $C : Y \rightarrow \{0, 1\}$ such that

$$C(O(z)) = 1 \iff z \in K \quad \text{for every } z \in Z.$$

Remark 3 (certainty-equivalence obstruction). Consider $\dot{S} = u - g(S)$ with $u \in [0, \bar{u}]$, g strictly increasing on $[S_{\min}, S^* + b]$, $g(0) = 0$, and $\mathcal{V} = [S_{\min}, S^*]$. Assume the range condition $g(S + b) \leq \bar{u}$ for every $S \in [S_{\min}, S^*]$, so that the certainty-equivalence command below is admissible rather than inadmissible. Under perfect information the feedback $u(t) = g(S(t))$ gives $\dot{S} = 0$, so every $S_0 \in \mathcal{V}$ is viable and $\text{Viab}(\mathcal{V}; U, \pi_{\text{perf}}) = \mathcal{V} \neq \emptyset$. Now take the injective, biased observation $\hat{S} = S + b$ with $b > 0$, and restrict the policy class to the **certainty-equivalence class** Π_{CE} , defined as the class of causal controllers obtained by taking a fixed *perfect-information regulation law* k of the plant and applying it directly to the uncorrected observation, $u = k(\hat{S})$ — no use of the observation map’s structure (here, the bias b) enters the law. The trap below is the member $k = g$: $u = g(\hat{S})$. Then

$$\dot{S} = g(S + b) - g(S) > 0 \quad \forall S,$$

and since g is strictly increasing on the compact interval $[S_{\min}, S^* + b]$, $\dot{S}$ is bounded below by a positive constant there; hence S strictly increases and exits above S^* in finite time from every $S_0 \in \mathcal{V}$:

$$\text{Viab}(\mathcal{V}; U, \Pi_{\text{CE}}) = \emptyset.$$

The mechanism is a *restriction of the policy class*, not a loss of information: for a known invertible observation the output-feedback class is as powerful as the state-feedback class, $\Pi_{\text{CE}} \subsetneq \Pi_{\text{output}} \cong \Pi_{\text{state}}$, so $\text{Viab}(\mathcal{V}; U, \Pi_{\text{CE}}) \subsetneq \text{Viab}(\mathcal{V}; U, \Pi_{\text{output}}) = \text{Viab}(\mathcal{V}; U, \Pi_{\text{state}})$ — the kernel empties exactly when the controller does not apply the correction. An observer who inverts the bias, $u = g(\hat{S} - b) = g(S)$, recovers the perfect-information kernel — and lies outside Π_{CE} , since the composite law $g(\cdot - b)$ is not a perfect-information law of the plant applied uncorrected. The remark is the mechanism behind the monitoring-design consequence of Section 6.4: monitoring bounds observation error so that a desired state-feedback law is implementable, but the bound is useful only if the controller uses the correction; an uncorrected biased indicator empties the kernel that a corrected one preserves. The phenomenon is the viability-theoretic analogue of Witsenhausen’s counterexample in stochastic control (Witsenhausen, 1968): the information pattern — not the dynamics and not the constraint — defeats the restricted controller class.

References

References

- Abae, A.: The limits of compensatory aggregation: a formal separation of weak and strong sustainability assessment. Zenodo. <https://doi.org/10.5281/zenodo.22545740> (2026).
- Aubin, J.-P.: Viability Theory. Birkhäuser, Boston (1991)
- Aubin, J.-P.: Viability kernels and capture basins of sets under differential inclusions. SIAM J. Control Optim. **40**(3), 853–871 (2001)
- Aubin, J.-P., Bayen, A.M., Saint-Pierre, P.: Viability Theory: New Directions, 2nd edn. Birkhäuser, Boston (2011)

- Aubin, J.-P., Catté, F.: Bilateral fixed-points and algebraic properties of viability kernels and capture basins of sets. *Set-Valued Anal.* **10**, 379–416 (2002)
- Aubin, J.-P., Frankowska, H.: *Set-Valued Analysis*. Birkhäuser, Boston (1990)
- Béné, C., Doyen, L., Gabay, D.: A viability analysis for a bio-economic model. *Ecol. Econ.* **36**, 385–396 (2001)
- Cardaliaguet, P., Quincampoix, M., Saint-Pierre, P.: Differential games through viability theory: old and recent results. In: Jørgensen, S., Quincampoix, M., Vincent, T.L. (eds.) *Advances in Dynamic Game Theory*. *Annals of the International Society of Dynamic Games*, vol. 9, pp. 3–35. Birkhäuser, Boston (2007)
- De Lara, M., Doyen, L.: *Sustainable Management of Natural Resources: Mathematical Models and Methods*. Springer, Berlin (2008)
- Doyen, L., Gajardo, P.: Sustainability standards, multicriteria maximin, and viability. *Nat. Resour. Model.* **33**(3), e12250 (2020)
- Doyen, L., Thébaud, O., Béné, C., Martinet, V., Gourguet, S., Bertignac, M., Fifas, S., Blanchard, F.: A stochastic viability approach to ecosystem-based fisheries management. *Ecol. Econ.* **75**, 32–42 (2012)
- Farkas, J.: Theorie der einfachen Ungleichungen. *J. Reine Angew. Math.* **124**, 1–27 (1902)
- Frankowska, H.: Optimal trajectories associated with a solution of contingent Hamilton–Jacobi equations. *Appl. Math. Optim.* **19**, 291–311 (1989)
- Gale, D.: *The Theory of Linear Economic Models*. McGraw-Hill, New York (1960)
- Haddad, G.: Monotone viable trajectories for functional-differential inclusions. *J. Differ. Equ.* **42**, 1–24 (1981)
- Hale, J.K., Verduyn Lunel, S.M.: *Introduction to Functional Differential Equations*. Springer, New York (1993)
- Kurzhanski, A.B., Vályi, I.: *Ellipsoidal Calculus for Estimation and Control*. Birkhäuser, Boston (1997)
- Maghenem, M., Sanfelice, R.G.: Characterizations of safety in hybrid inclusions via barrier functions. In: *Proceedings of the 22nd ACM International Conference on Hybrid Systems: Computation and Control (HSCC)*, pp. 109–118 (2019)
- Martinez-Alier, J., Munda, G., O’Neill, J.: Weak comparability of values as a foundation for ecological economics. *Ecol. Econ.* **26**(3), 277–286 (1998)
- Nagumo, M.: Über die Lage der Integralkurven gewöhnlicher Differentialgleichungen. *Proc. Phys.-Math. Soc. Japan* **24**, 551–559 (1942)
- Prajna, S., Jadbabaie, A.: Safety verification of hybrid systems using barrier certificates. In: Alur, R., Pappas, G.J. (eds.) *Hybrid Systems: Computation and Control (HSCC 2004)*. LNCS, vol. 2993, pp. 477–492. Springer, Berlin (2004)
- Prajna, S., Jadbabaie, A., Pappas, G.J.: A framework for worst-case and stochastic safety verification using barrier certificates. *IEEE Trans. Autom. Control* **52**(8), 1415–1428 (2007)
- Prajna, S., Rantzer, A.: On the necessity of barrier certificates. *IFAC Proc.* Vol. **38**(1), 526–531 (2005)
- Quincampoix, M., Veliov, V.M.: Viability with a target: theory and applications. In: *Control Theory, Multivariate Analysis and Applications*, pp. 47–63. Springer (1994)
- Saint-Pierre, P.: Approximation of the viability kernel. *Appl. Math. Optim.* **29**, 187–209 (1994)
- Smallwood, R.D., Sondik, E.J.: The optimal control of partially observable Markov processes over a finite horizon. *Oper. Res.* **21**(5), 1071–1088 (1973)
- Sontag, E.D.: *Mathematical Control Theory: Deterministic Finite Dimensional Systems*, 2nd edn. Springer, New York (1998)
- Veliov, V.M.: Sufficient conditions for viability under imperfect measurement. *Set-Valued Anal.* **1**, 305–317 (1993)

Witsenhausen, H.S.: A counterexample in stochastic optimum control. SIAM J. Control **6**(1), 131–147 (1968)